

Relative cost of matrix operations

In [1]:

```
#keep
import numpy as np
import scipy.linalg as spla
import scipy as sp

import matplotlib.pyplot as plt

from time import time
%matplotlib inline

np.alterdot()
```

In [2]:

```
#keep
n_values = (10**np.linspace(1, 3.75, 15)).astype(np.int32)
n_values
```

Out[2]:

```
array([ 10,  15,  24,  38,  61,  95, 150, 237, 372, 585,
        921,
        1447, 2275, 3577, 5623], dtype=int32)
```

In [3]:

```
#keep
```

```
def mat_mul(A):  
    return A.dot(A)  
  
for name, f in [  
    ("mat_mul", mat_mul),  
    ("lu", spla.lu_factor),  
    ]:  
  
    times = []  
    print("----->", name)  
  
    for n in n_values:  
        print(n)  
  
        A = np.random.randn(n, n)  
  
        start_time = time()  
        f(A)  
        times.append(time() - start_time)  
  
    pt.plot(n_values, times, label=name)  
  
pt.grid()  
pt.legend(loc="best")  
pt.xlabel("Matrix size $n$")  
pt.ylabel("Wall time [s]")
```

```
-----> mat_mul
```

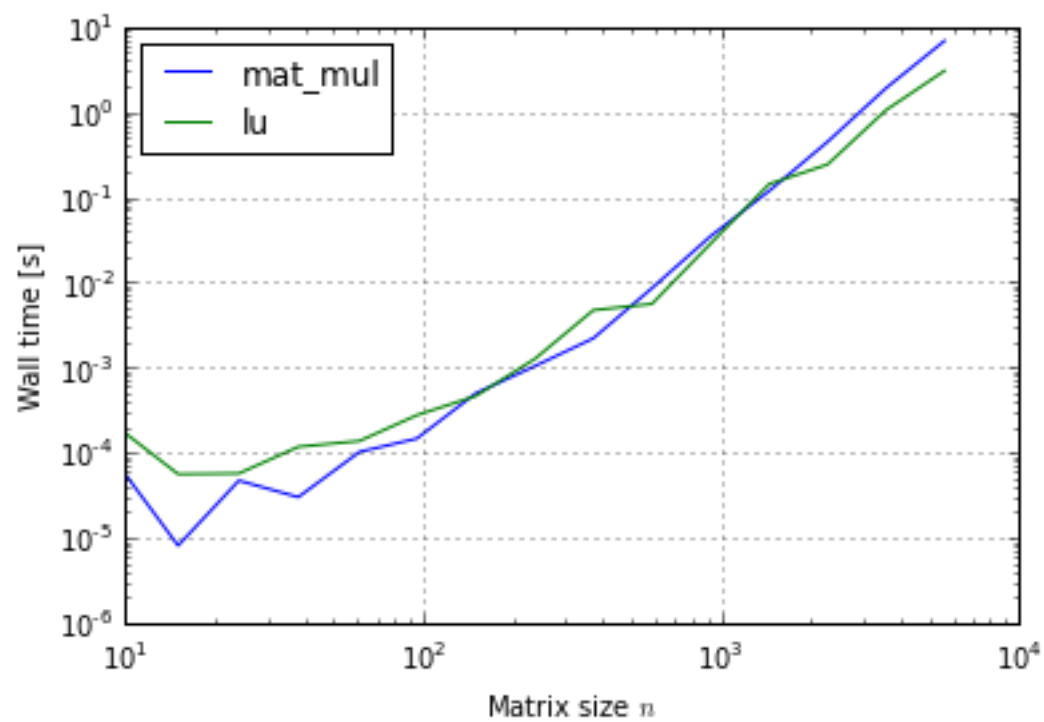
```
10
15
24
38
61
95
150
237
372
585
921
1447
2275
3577
5623
```

```
-----> lu
```

```
10
15
24
38
61
95
150
237
372
585
921
1447
2275
3577
5623
```

```
Out[3]:
```

```
<matplotlib.text.Text at 0x10dc5acc0>
```



- The faster algorithms make the slower ones look bad. But... it's all relative.
- Is there a better way of plotting this?
- Can we see the asymptotic cost ($O(n^3)$) of these algorithms from the plot?

In []:

```
pt.show( )
```

In []: