

Picking apart a floating point number

In [1]:

```
#keep

# Never mind the details of this function...

def pretty_print_fp(x):
    print("-----")
    print("Floating point structure for %r" % x)
    print("-----")
    import struct
    s = struct.pack("d", x)

    def get_bit(i):
        byte_nr, bit_nr = divmod(i, 8)
        return int(bool(
            s[byte_nr] & (1 << bit_nr)
        ))

    def get_bits(lsb, count):
        return sum(get_bit(i+lsb)*2**i for i in range(count))

    # https://en.wikipedia.org/wiki/Double_precision_floating-point_format

    print("                    1             2             3             4
5")
    print("indices                : 012345678901234567890123456789012345
67890123456789")
    print("Sign bit (1:negative):", get_bit(63))
    exponent = get_bits(52, 11)
    print("Exponent      (binary):", bin(exponent)[2:])
    print("Exponent      (shifted): %d" % (exponent - 1023))
    fraction = get_bits(0, 52)
    significand = fraction + 2**52
    print("Significand   (binary):", bin(significand)[2:])
    print("Significand   (shifted):", repr(significand / (2**52)))
```

In [8]:

```
#keep
pretty_print_fp(2**1024)
```

```
-----
Floating point structure for 179769313486231590772930519078902473361
79769789423065727343008115773267580550096313270847732240753602112011
38798713933576587897688144166224928474306394741243777678934248654852
76302219601246094119453082952085005768838150682342462881473913110540
827237163350510684586298239947245938479716304835356329624224137216
-----
```

```
-----
error                                Traceback (most recent call last)
<ipython-input-8-46758d327020> in <module>()
      1 #keep
----> 2 pretty_print_fp(2**1024)
```

```
<ipython-input-1-0c0d10989b38> in pretty_print_fp(x)
      8     print("-----")
      9     import struct
----> 10     s = struct.pack("d", x)
      11
      12     def get_bit(i):
```

error: required argument is not a float

Things to try:

- Twiddle the sign bit
- 1,2,4,8
- 0.5,0.25
- 2^{1023} , 2^{1024}
- 2^{-1023} , 2^{-1024}
- float("nan")

In [15]:

```
import numpy as np
```

In [16]:

```
np.binary_repr(1024)
```

Out[16]:

```
'100000000000'
```

In [18]:

```
bin(1024+1023)
```

Out[18]:

```
'0b111111111111'
```

In []: