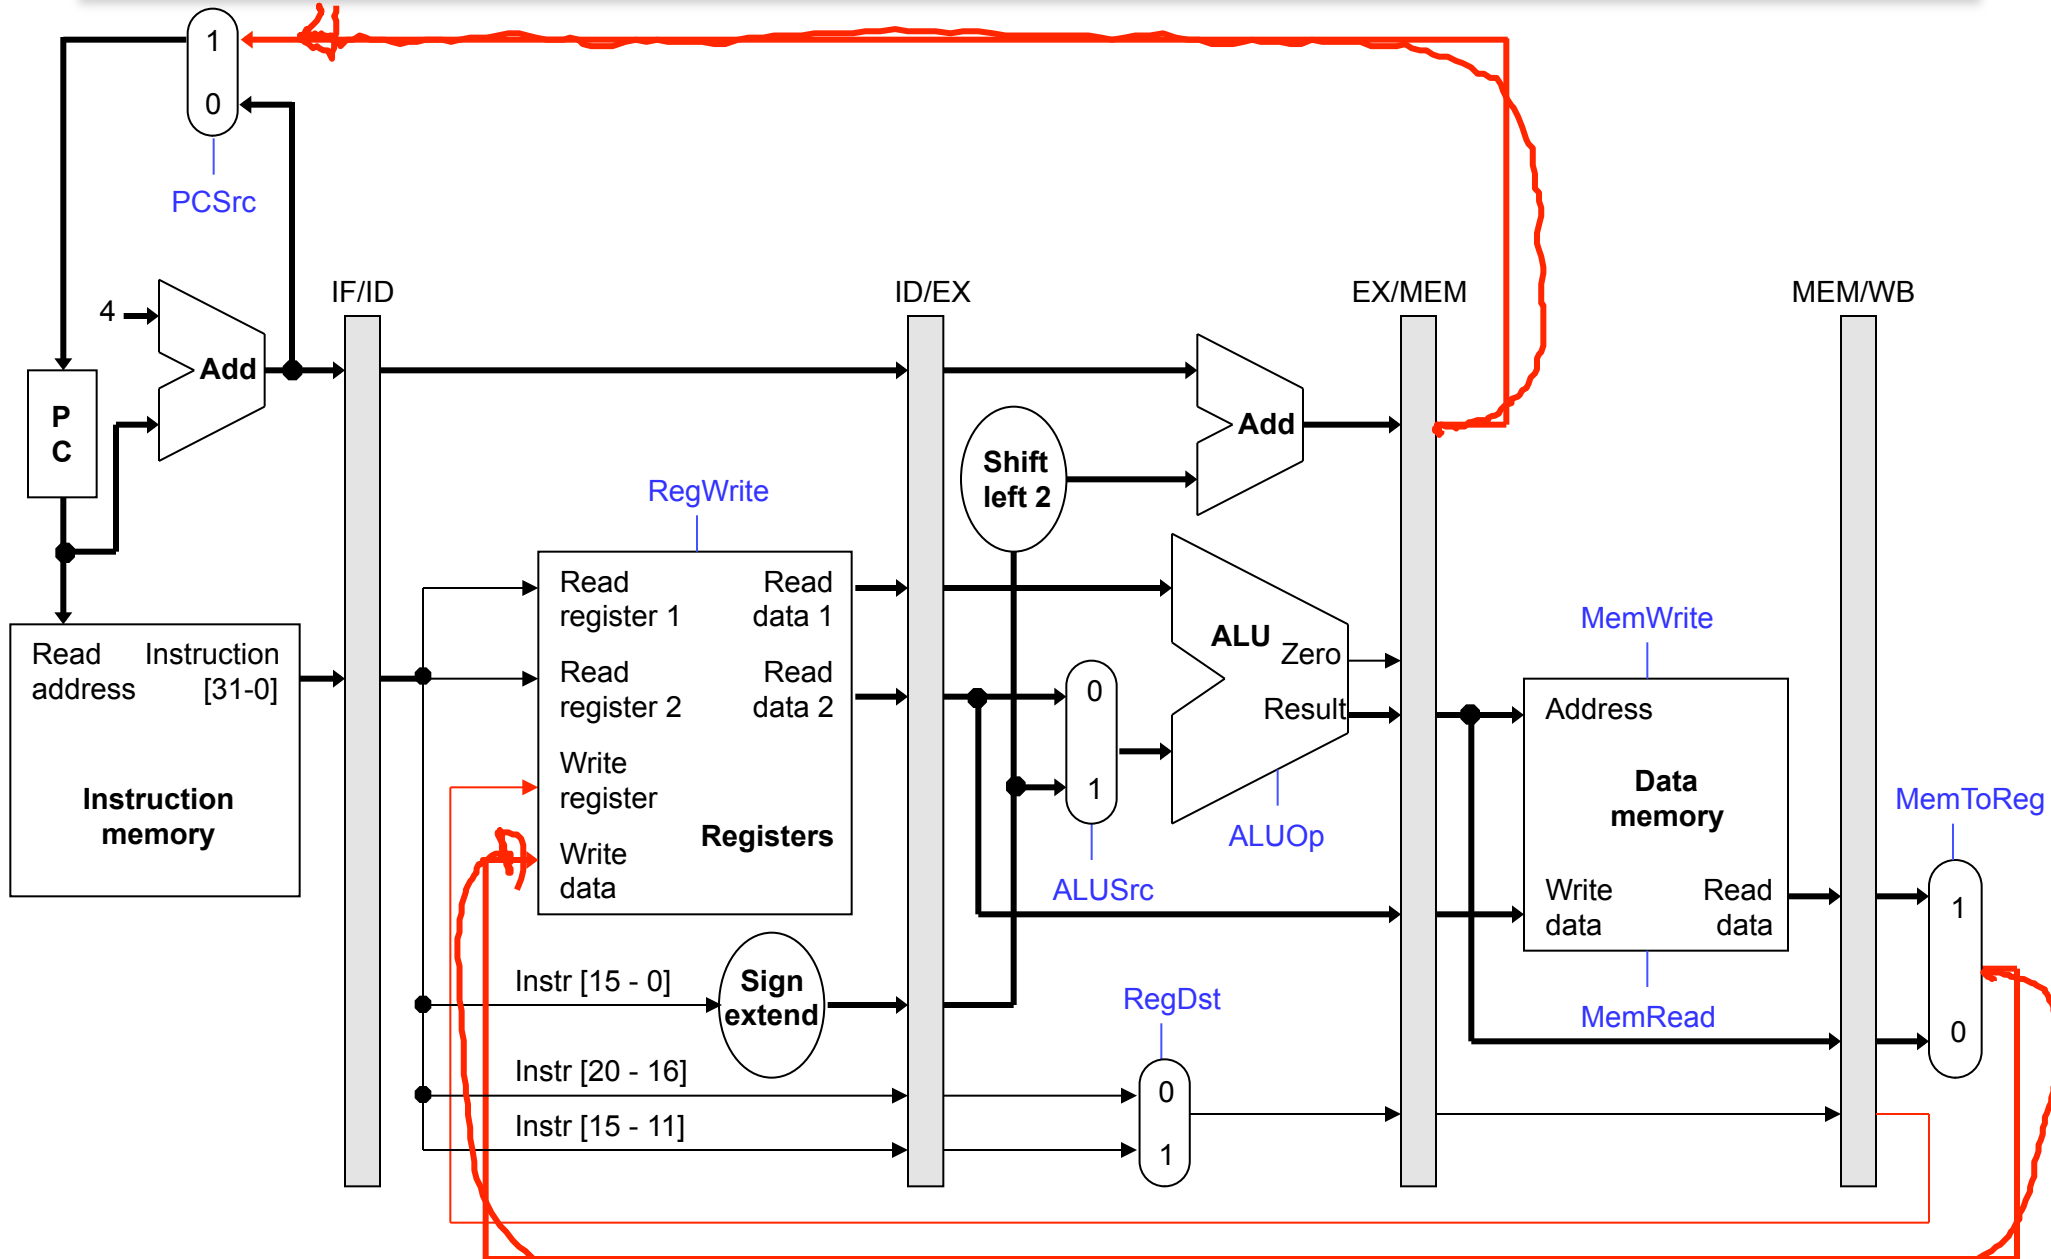

Everyone should have received an email about TB screening
If not, please contact me



Note how everything goes left to right, except ...



You should have received an email about screening for TB.

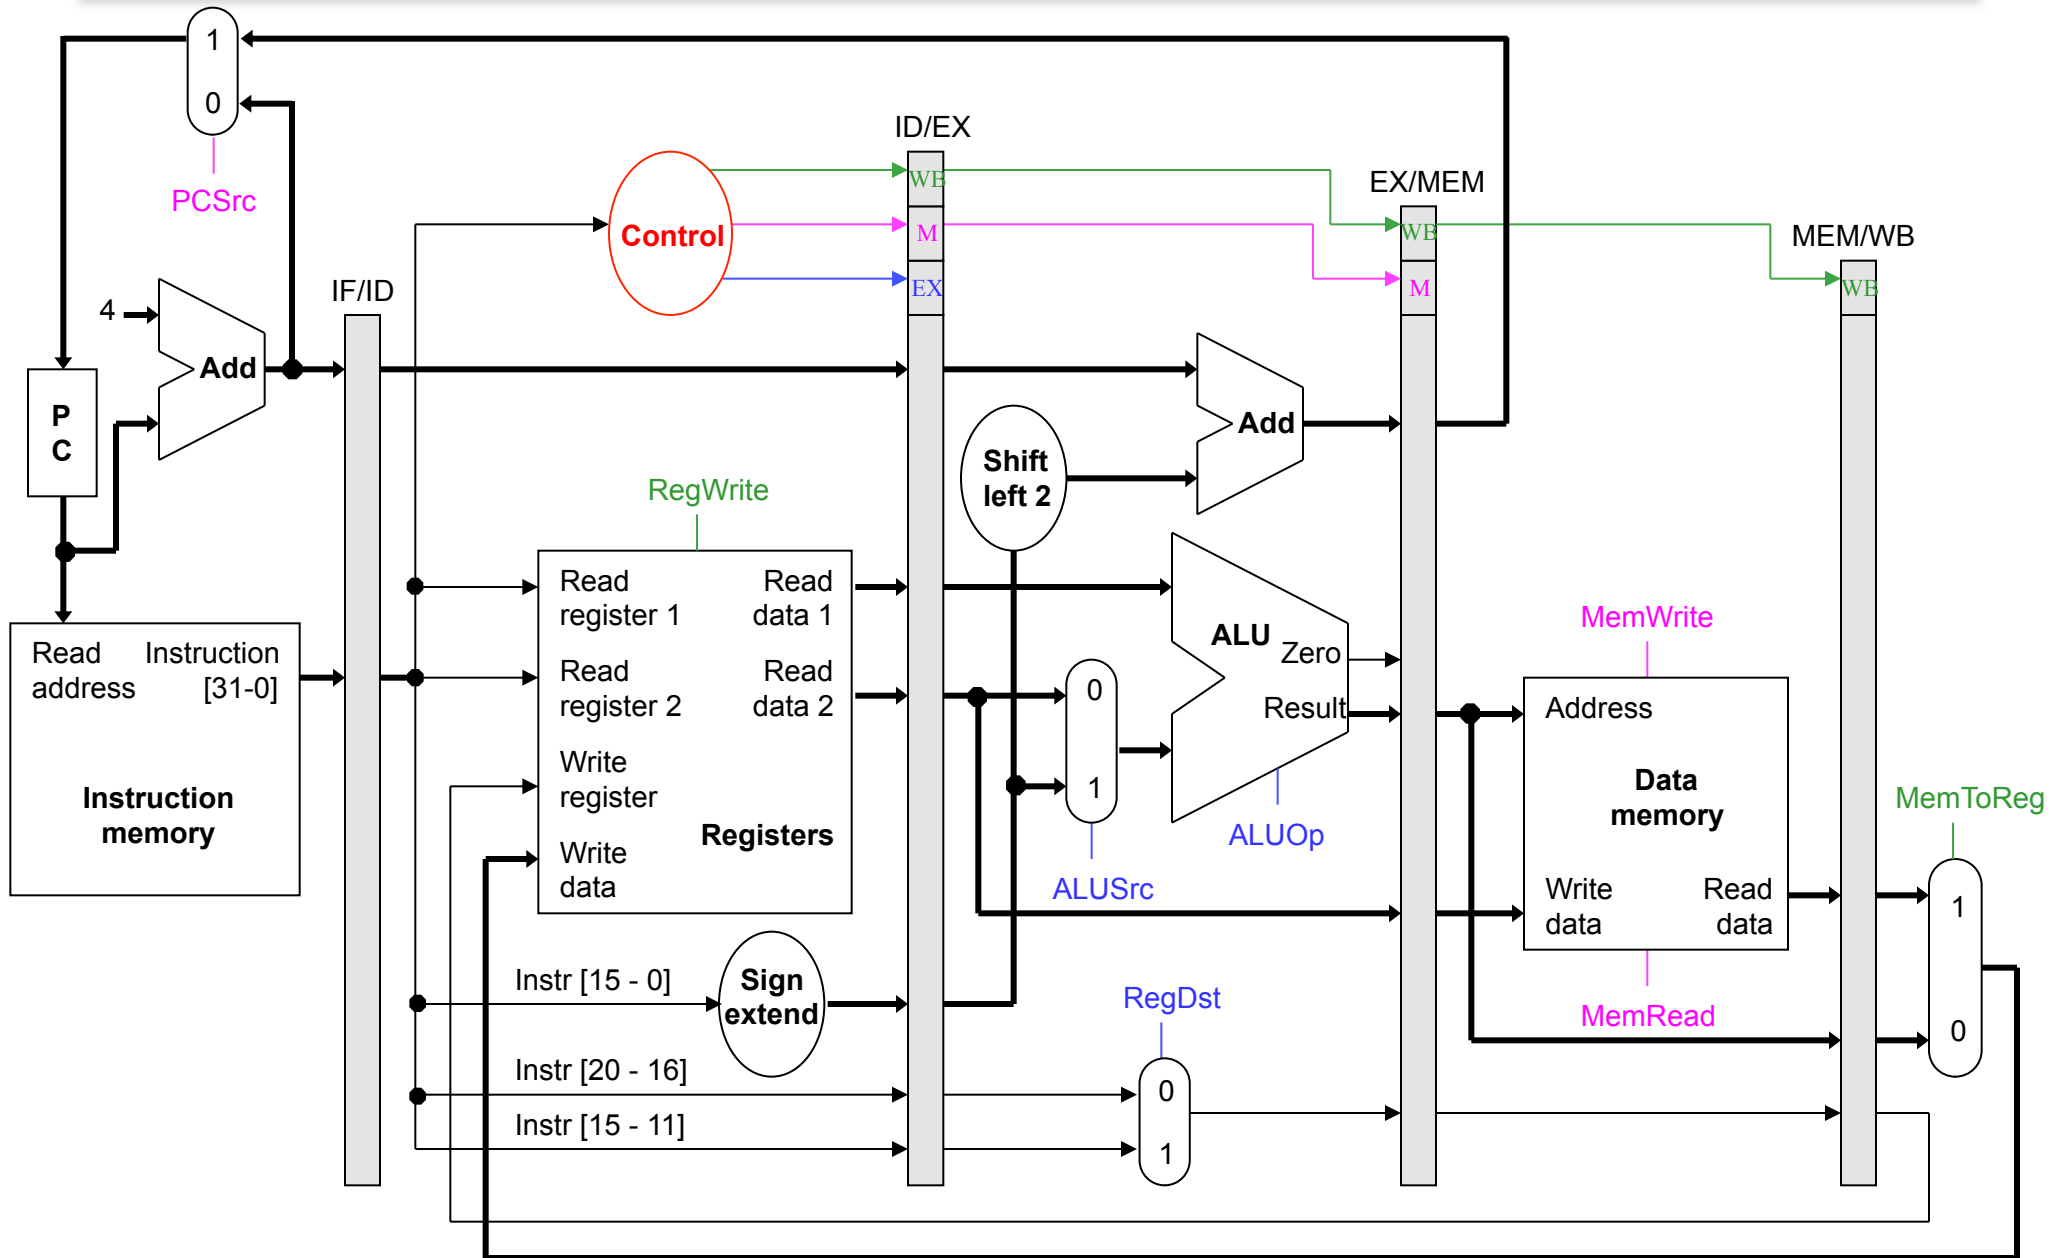
If you have not, please contact me.

Forwarding


- Previously, we introduced a **pipelined** MIPS processor which executes several instructions simultaneously.
 - Each instruction requires five stages, and five cycles, to complete.
 - Each stage uses different functional units of the datapath.
 - So we can execute up to five instructions in any clock cycle, with each instruction in a different stage and using different hardware.
- Today we'll introduce some problems that **data hazards** can cause for our pipelined processor, and show how to handle them with **forwarding**.



The pipelined datapath



Pipeline diagram review



- This diagram shows the execution of an ideal code fragment.
 - Each instruction needs a total of five cycles for execution.
 - One instruction begins on every clock cycle for the first five cycles.
 - One instruction completes on each cycle from that time on.

Our examples are too simple

- Here is the example instruction sequence used to illustrate pipelining on the previous page.

```
lw    $8, 4($29)
sub   $2, $4, $5
and   $9, $10, $11
or    $16, $17, $18
add   $13, $14, $0
```

- The instructions in this example are independent.
 - Each instruction reads and writes completely different registers.
 - Our datapath handles this sequence easily, as we saw last time.
- But most sequences of instructions are *not* independent!

An example with dependencies

Producer
Write

①	sub	\$2, \$1, \$3
②	and	\$12, \$2, \$5
③	or	\$13, \$6, \$2
④	add	\$14, \$2, \$2
⑤	sw	\$15, 100(\$2)

Consumer
Read

True dependencies


RAW
Read After Write

An example with dependencies

```
sub    $2, $1, $3
and    $12, $2, $5
or     $13, $6, $2
add    $14, $2, $2
sw     $15, 100($2)
```

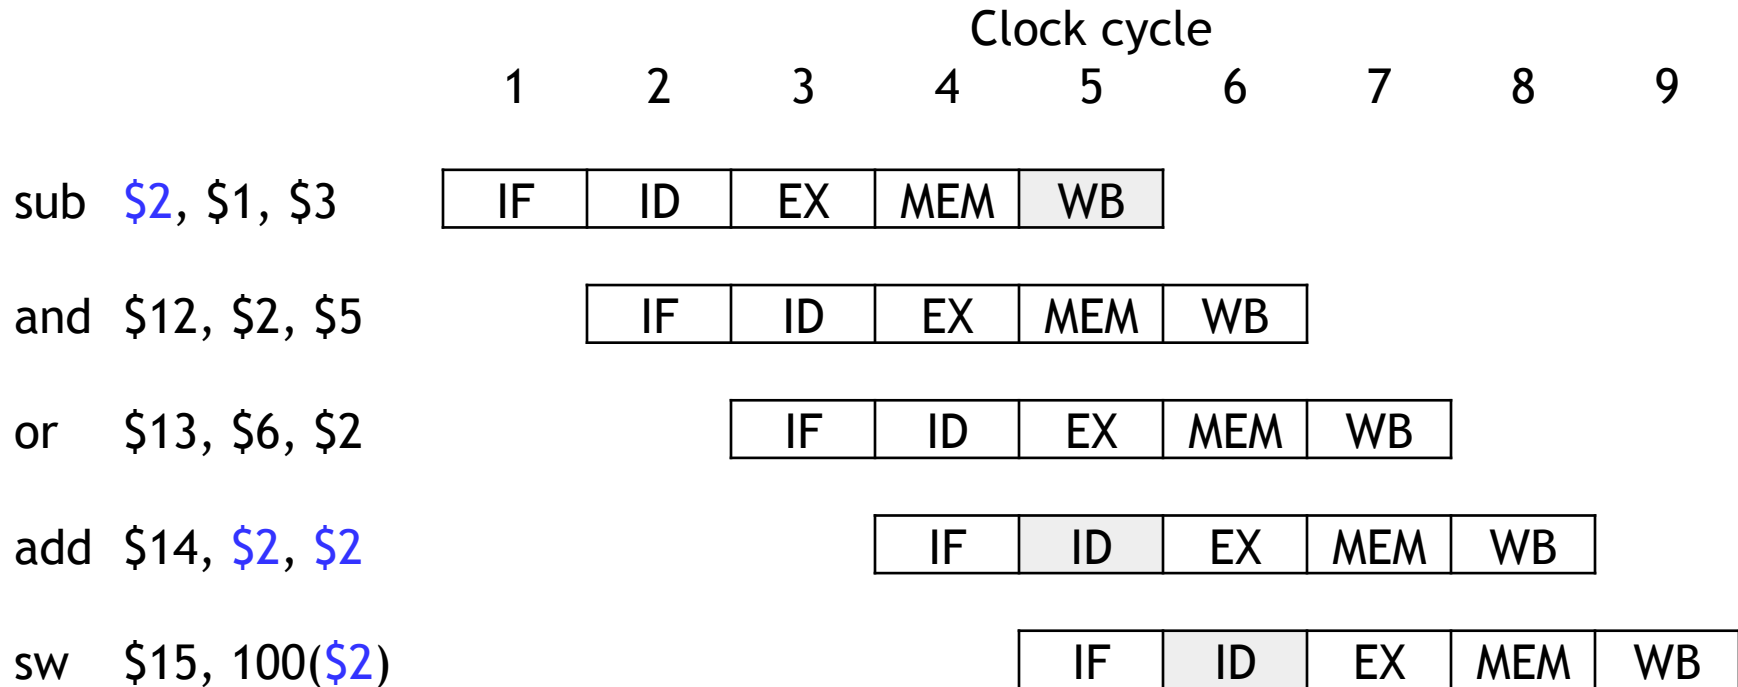
- There are several **dependencies** in this new code fragment.
 - The first instruction, SUB, stores a value into **\$2**.
 - That register is used as a source in the rest of the instructions.
- This is not a problem for the single-cycle datapath.
 - Each instruction is executed completely before the next one begins.
 - This ensures that instructions 2 through 5 above use the new value of **\$2** (the sub result), just as we expect.
- How would this code sequence fare in our pipelined datapath?

Data hazards in the pipeline diagram




- The SUB instruction does not write to register \$2 until clock cycle 5. This causes two **data hazards** in our current pipelined datapath.
 - The AND reads register \$2 in cycle 3. Since SUB hasn't modified the register yet, this will be the *old* value of \$2, not the new one.
 - Similarly, the OR instruction uses register \$2 in cycle 4, again before it's actually updated by SUB.

Things that are okay




- The ADD instruction is okay, because of the register file design.
 - Registers are written at the beginning of a clock cycle.
 - The new value will be available by the end of that cycle.
- The SW is no problem at all, since it reads \$2 after the SUB finishes.

Dependency arrows



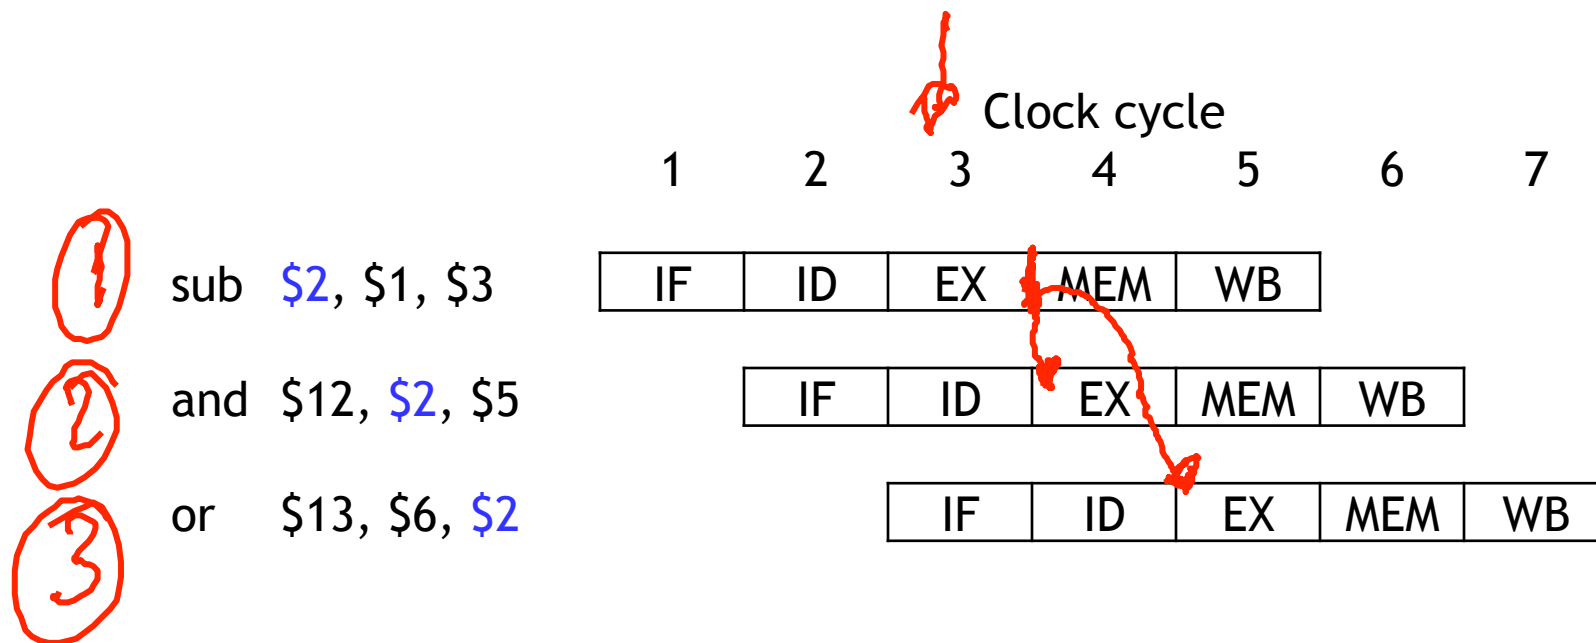
- Arrows indicate the flow of data between instructions.
 - The tails of the arrows show when register \$2 is written.
 - The heads of the arrows show when \$2 is read.
- Any arrow that points backwards in time represents a data hazard in our basic pipelined datapath. Here, hazards exist between instructions 1 & 2 and 1 & 3.

A fancier pipeline diagram




A more detailed look at the pipeline

- We have to eliminate the hazards, so the AND and OR instructions in our example will use the correct value for register \$2.
- When is the data actually produced and consumed?
- What can we do?




A more detailed look at the pipeline

- We have to eliminate the hazards, so the AND and OR instructions in our example will use the correct value for register \$2.
- Let's look at when the data is actually produced and consumed.
 - The SUB instruction produces its result in its **EX** stage, during cycle 3 in the diagram below.
 - The AND and OR need the new value of \$2 in their **EX** stages, during clock cycles 4-5 here.



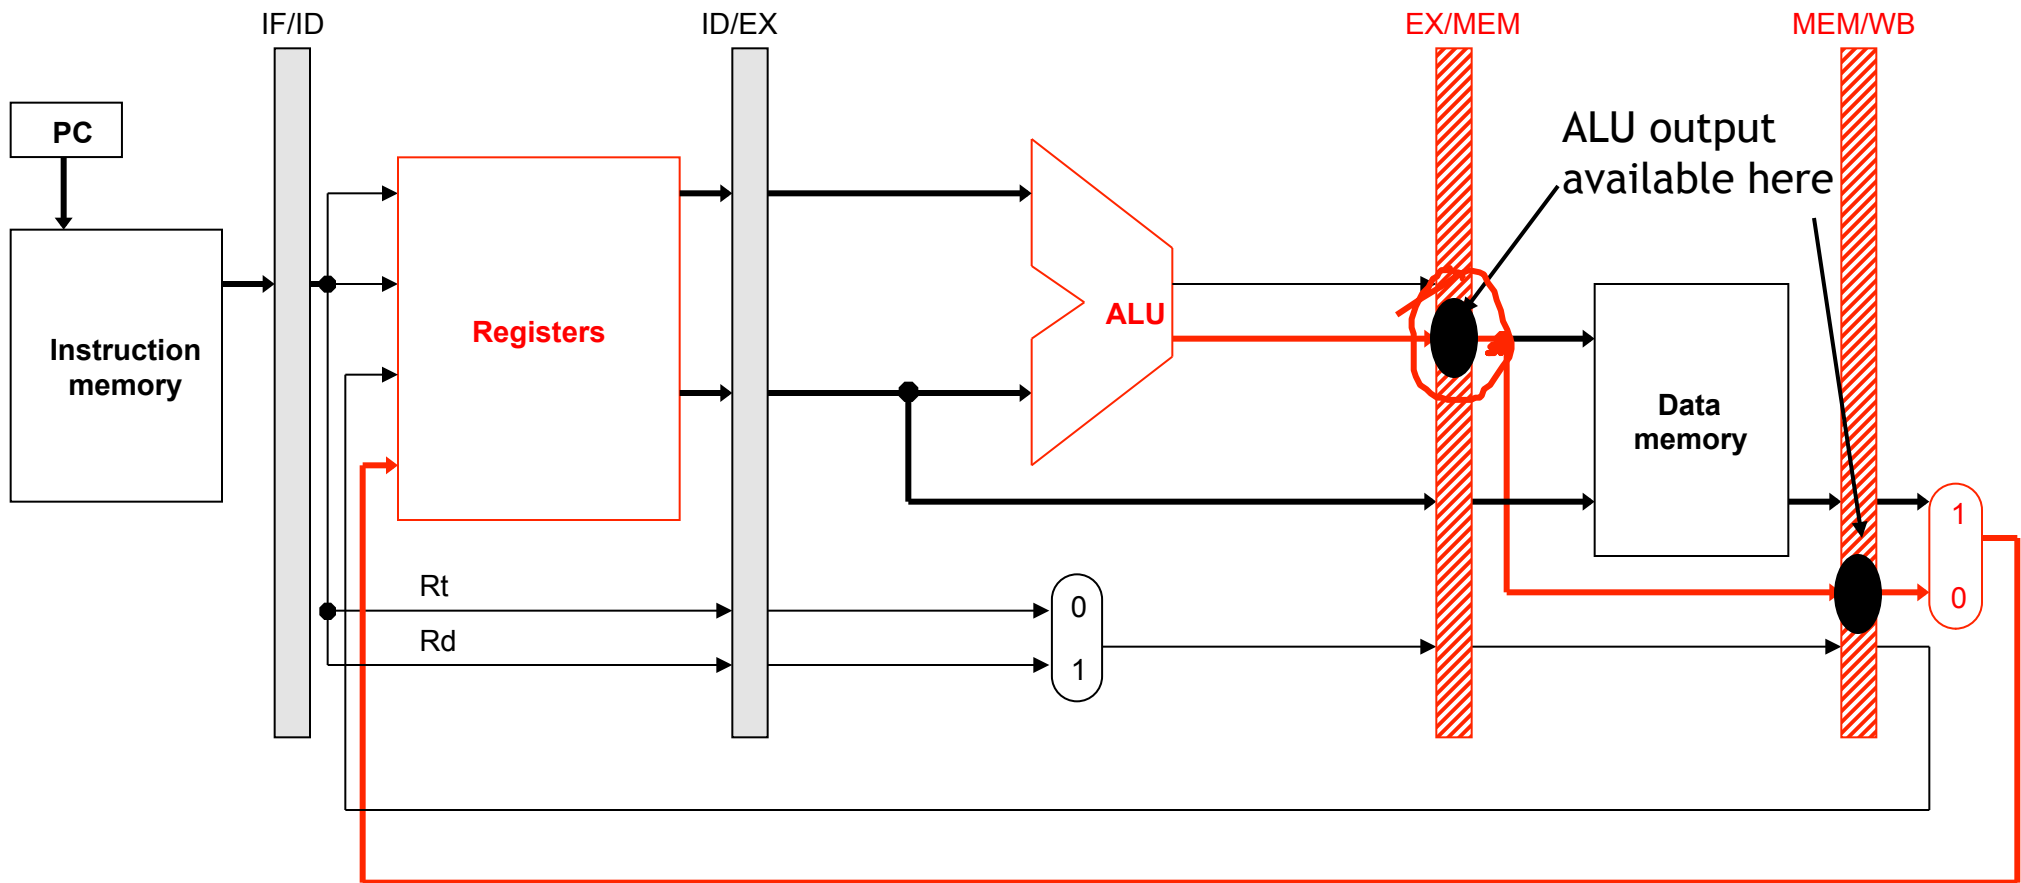
Bypassing the register file

- The actual result \$1 – \$3 is computed in clock cycle 3, *before* it's needed in cycles 4 and 5.
- If we could somehow bypass the writeback and register read stages when needed, then we can eliminate these data hazards.
 - Today we'll focus on hazards involving arithmetic instructions.
 - Next time, we'll examine the lw instruction.
- Essentially, we need to pass the ALU output from SUB directly to the AND and OR instructions, without going through the register file.



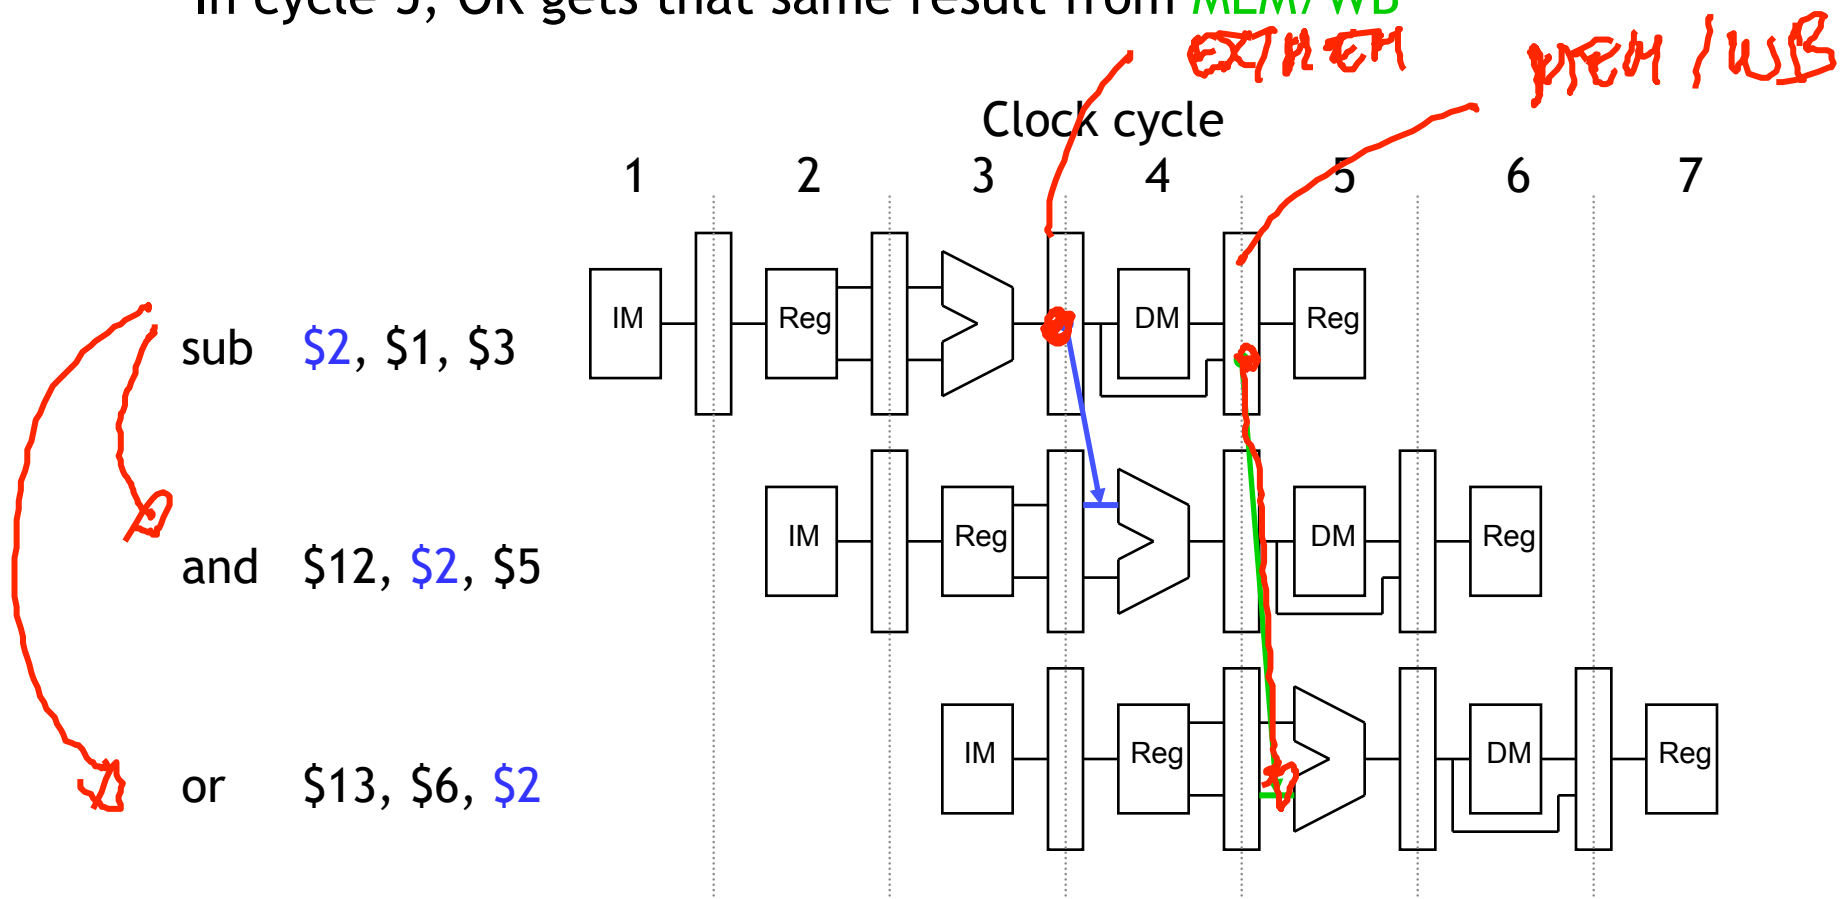
Pipeline Registers to the rescue!

- Pipeline stages communicate through pipeline registers:
IF/ID ID/EX EX/MEM MEM/WB
- We “forward” data from pipeline registers to later instructions



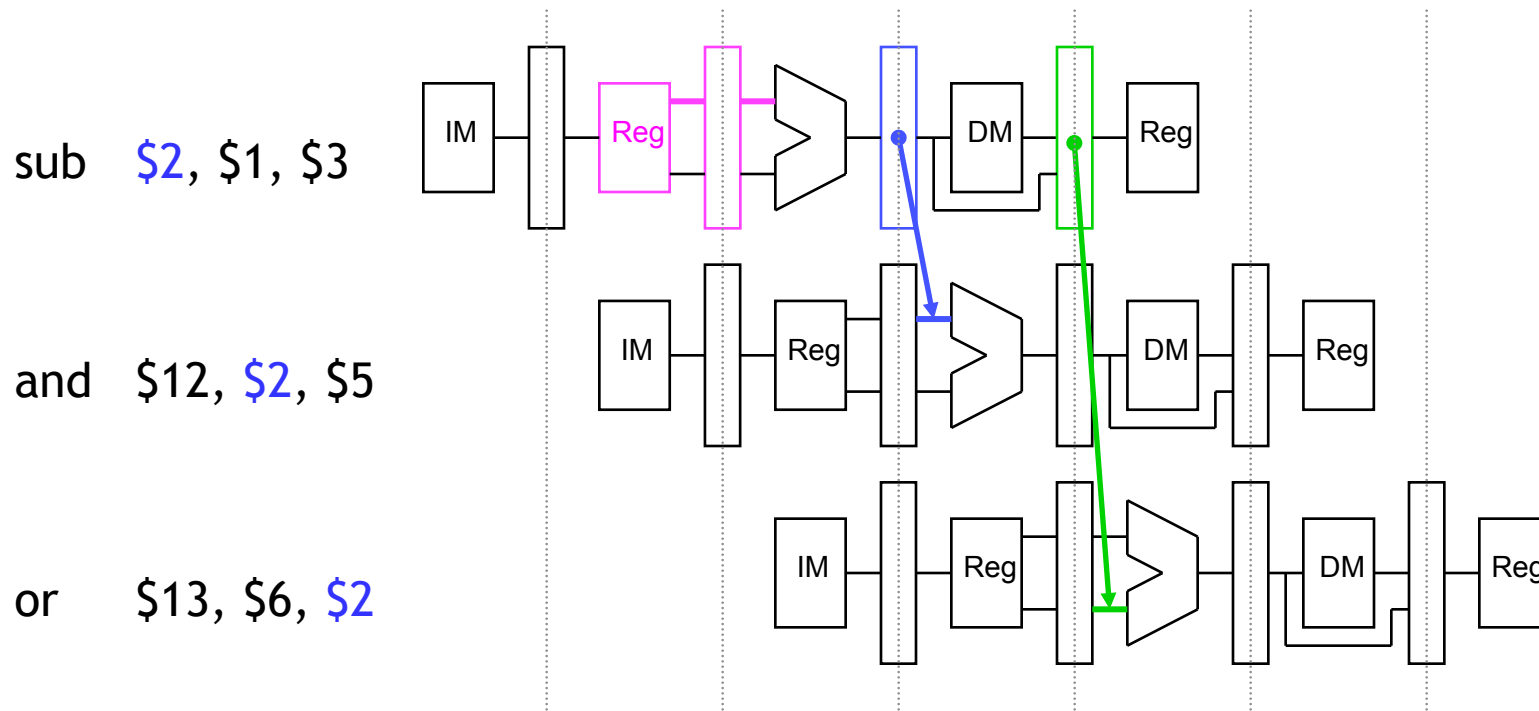
Forwarding

- The actual result $\$1 - \3 is computed in clock cycle 3, *before* it's needed in cycles 4 and 5
- We **forward** that value to later instructions, to prevent data hazards:
 - In clock cycle 4, AND gets the value $\$1 - \3 from **EX/MEM**
 - In cycle 5, OR gets that same result from **MEM/WB**

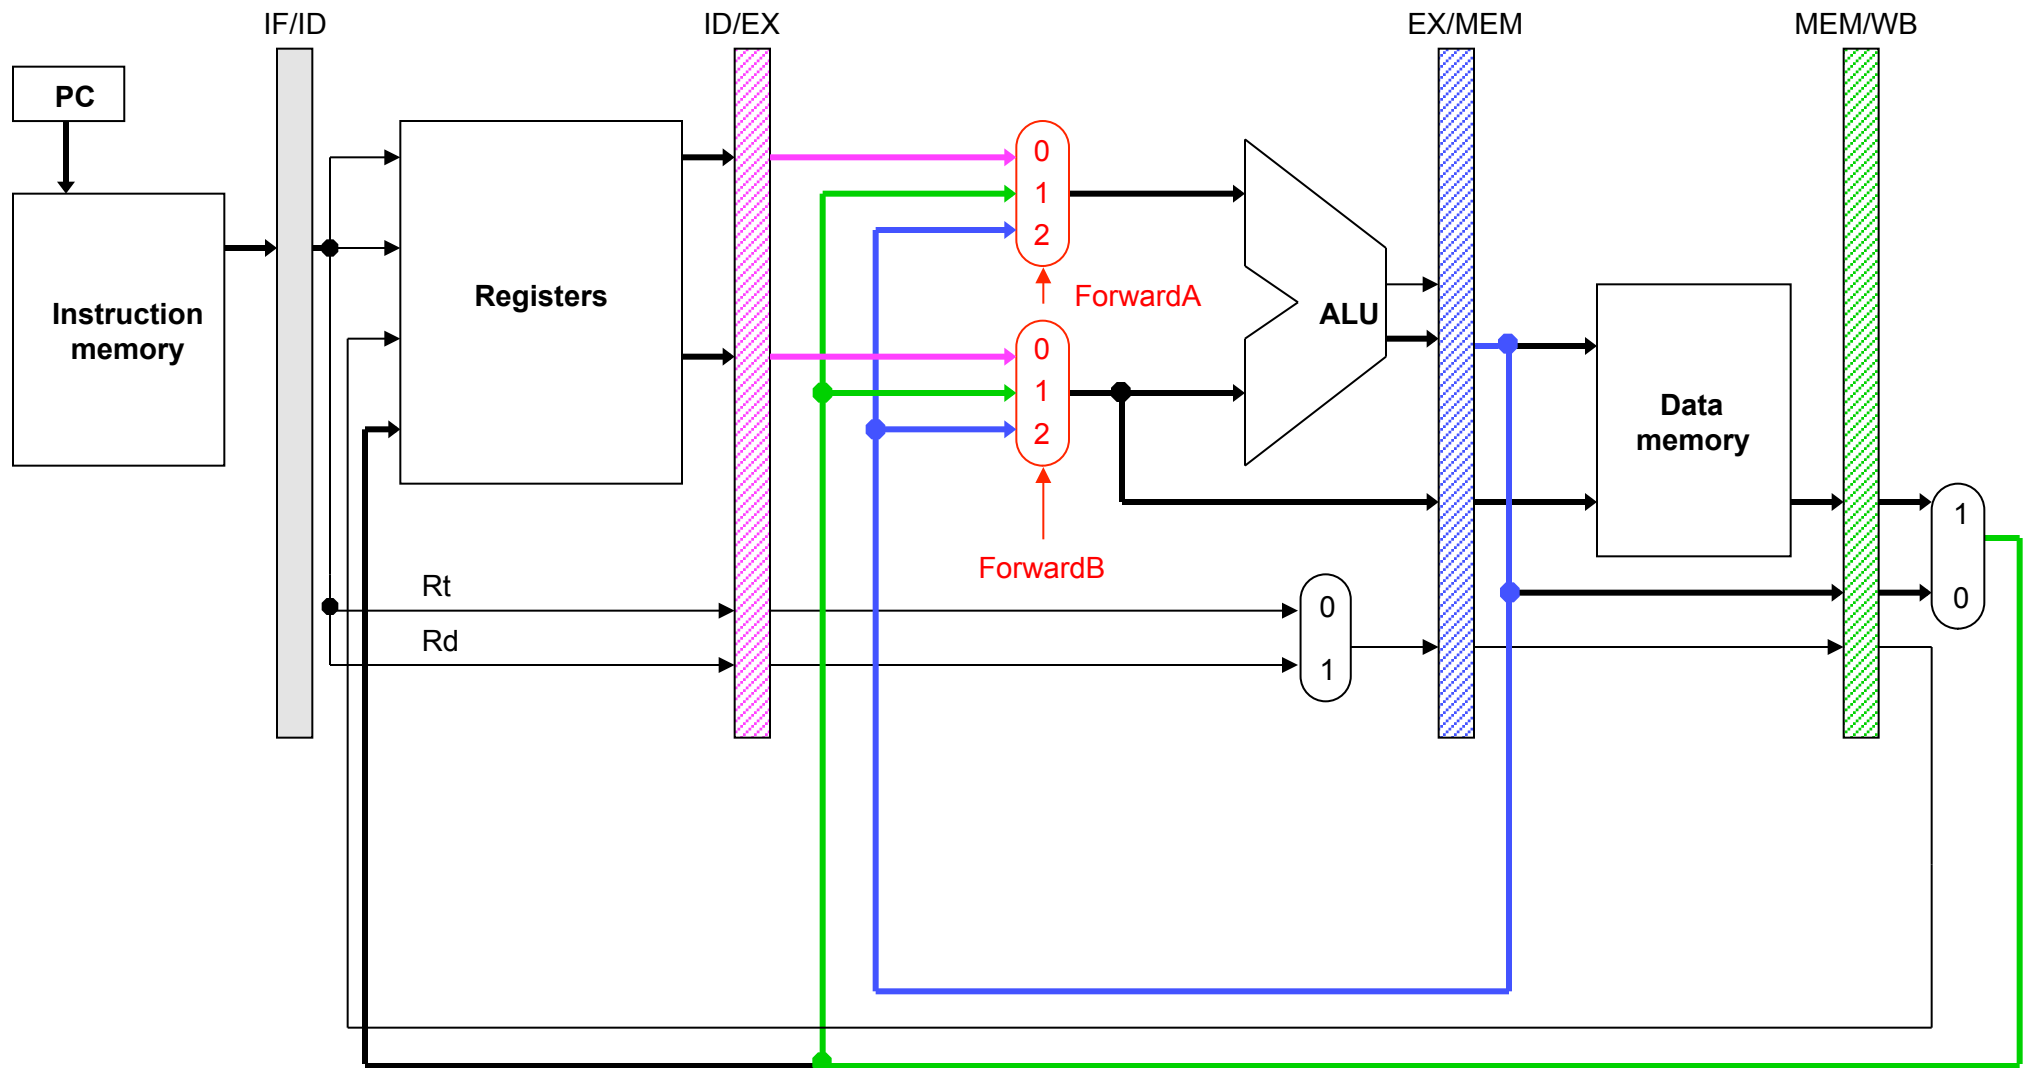


Outline of forwarding hardware

- A **forwarding unit** selects the correct ALU inputs for the EX stage:
 - No hazard: ALU's operands come from the **register file**, like normal
 - Data hazard: operands come from either the **EX/MEM** or **MEM/WB** pipeline registers instead
- The ALU sources will be selected by two new multiplexers, with control signals named **ForwardA** and **ForwardB**



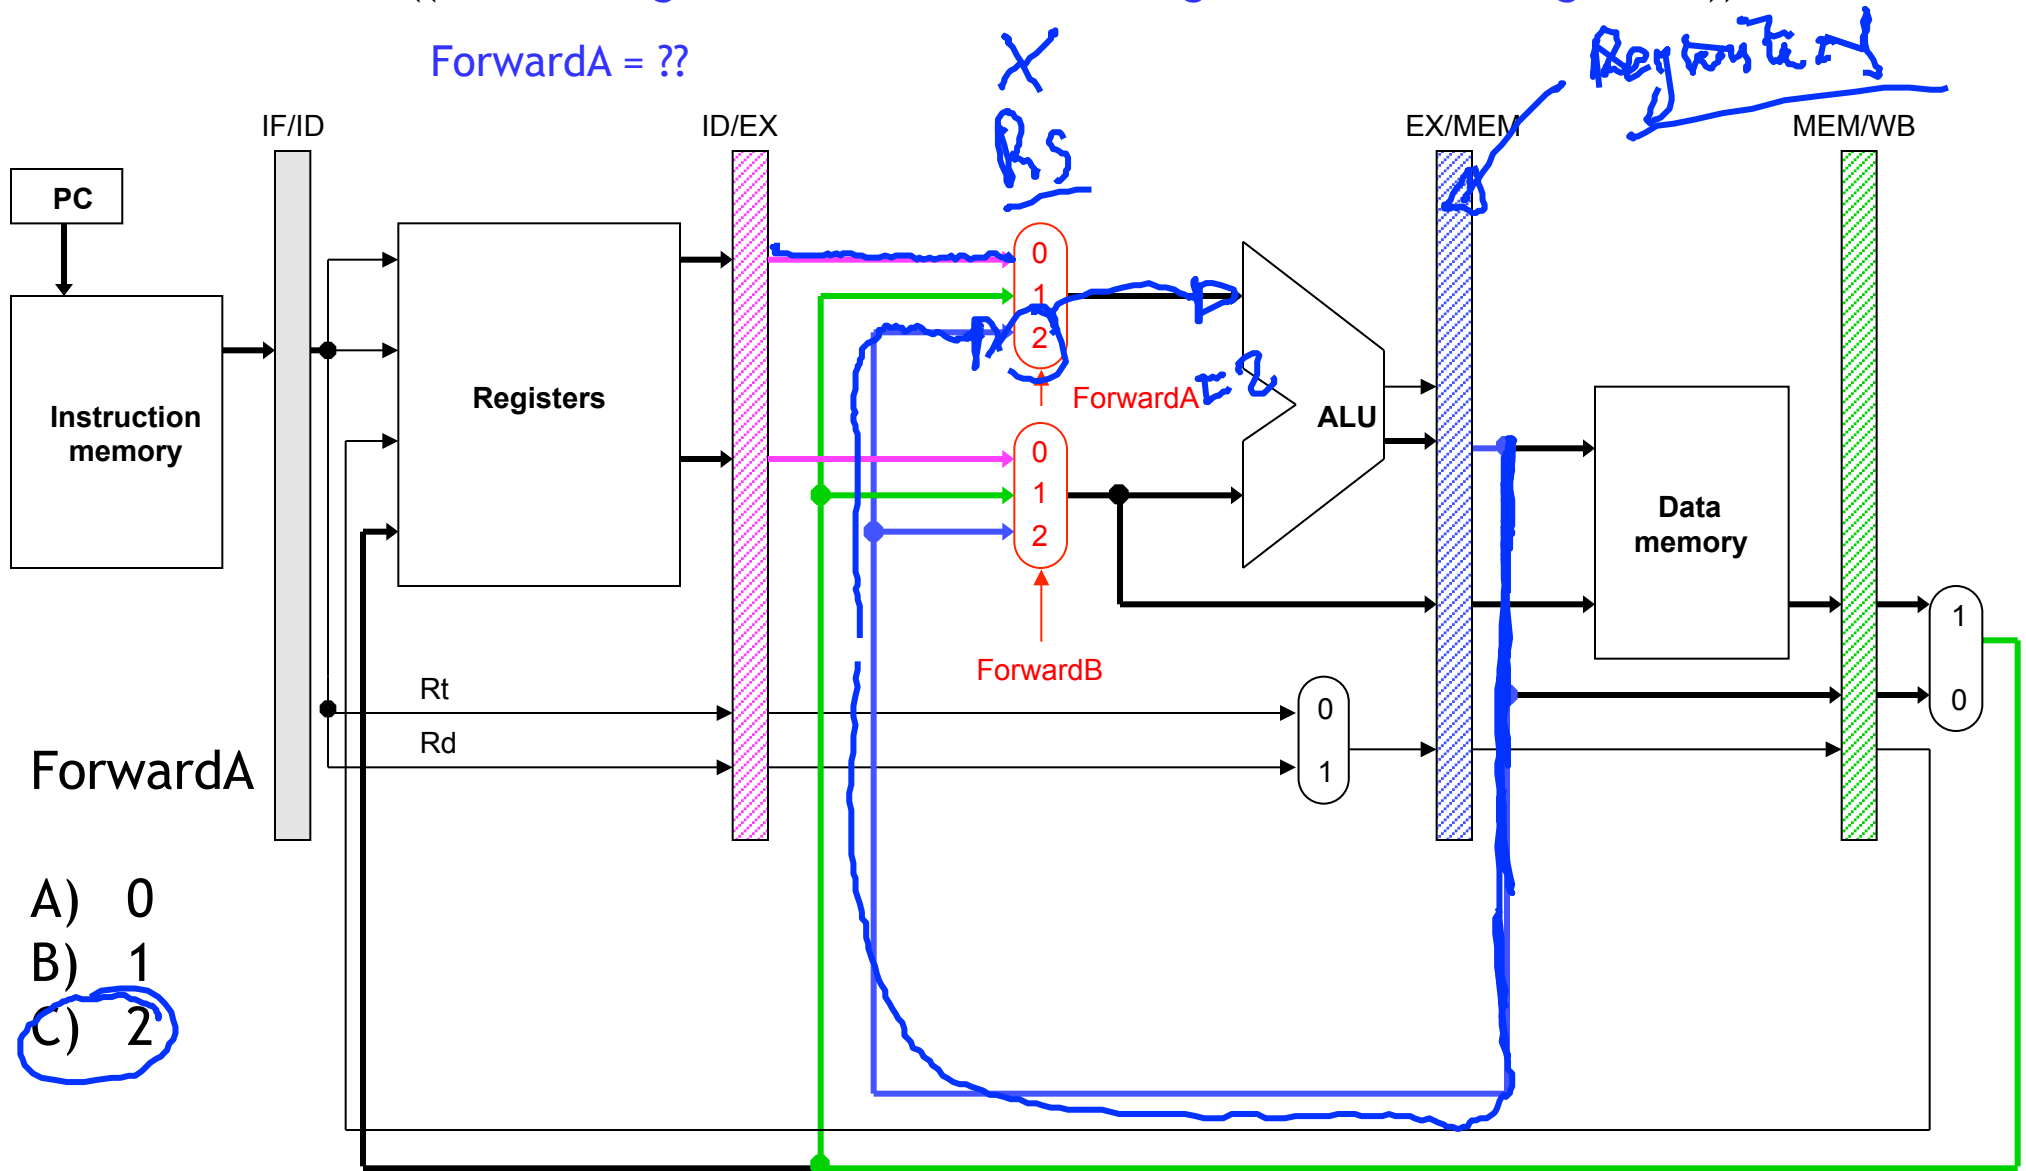
Simplified datapath with forwarding muxes



Simplified datapath with forwarding muxes

if ((EX/MEM.RegWrite = 1 and EX/MEM.RegisterRd = ID/EX.RegisterRs))

ForwardA = ??




- A) 0
- B) 1
- C) 2

Detecting EX/MEM data hazards

- When do we need to know that a hazard exists?
- So how can the hardware determine if a hazard exists?

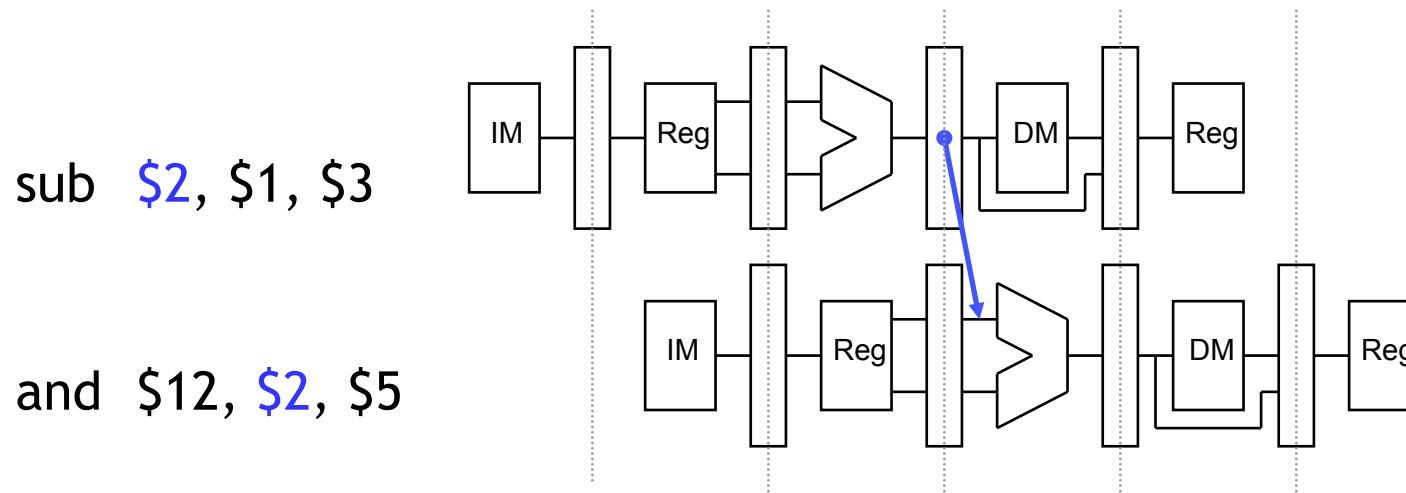
sub \$2, \$1, \$3

and \$12, \$2, \$5



Detecting EX/MEM data hazards

- So how can the hardware determine if a hazard exists?
- An **EX/MEM hazard** occurs between the instruction currently in its EX stage and the previous instruction if:
 1. The previous instruction will write to the register file, *and*
 2. The destination is one of the ALU source registers in the EX stage.
- There is an EX/MEM hazard between the two instructions below.



- Data in a pipeline register can be referenced using a class-like syntax. For example, **ID/EX.RegisterRt** refers to the rt field stored in the ID/EX pipeline.

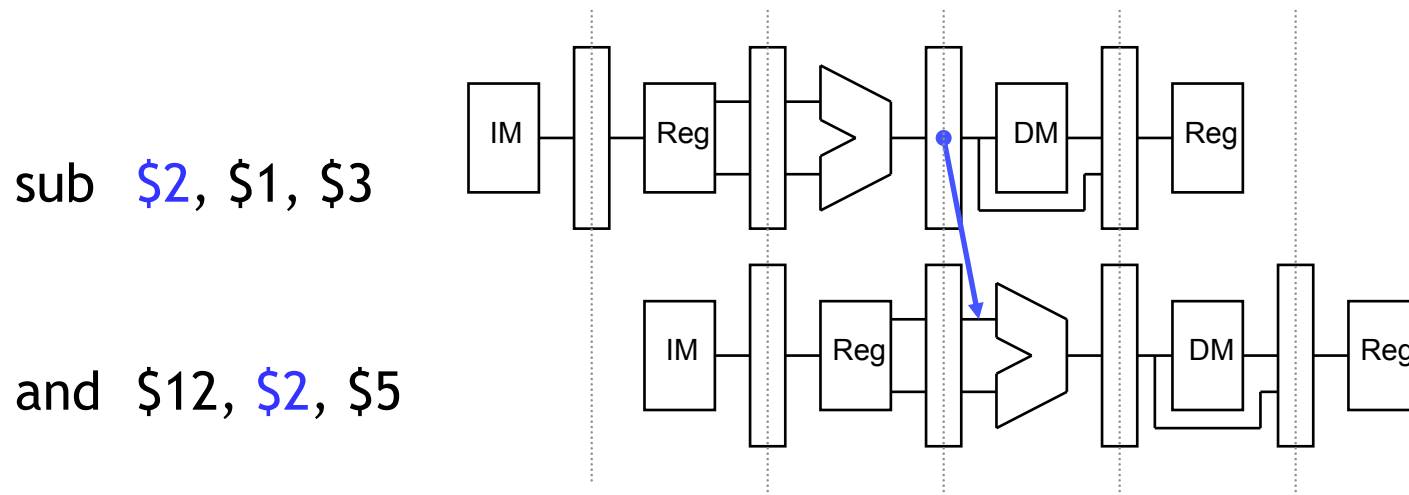
EX/MEM data hazard equations

- The first ALU source comes from the pipeline register when necessary.

if ($EX/MEM.RegWrite = 1$
and $EX/MEM.RegisterRd = ID/EX.RegisterRs$)
then $ForwardA = 2$


- The second ALU source is similar.

if ($EX/MEM.RegWrite = 1$
and $EX/MEM.RegisterRd = ID/EX.RegisterRt$)
then $ForwardB = 2$

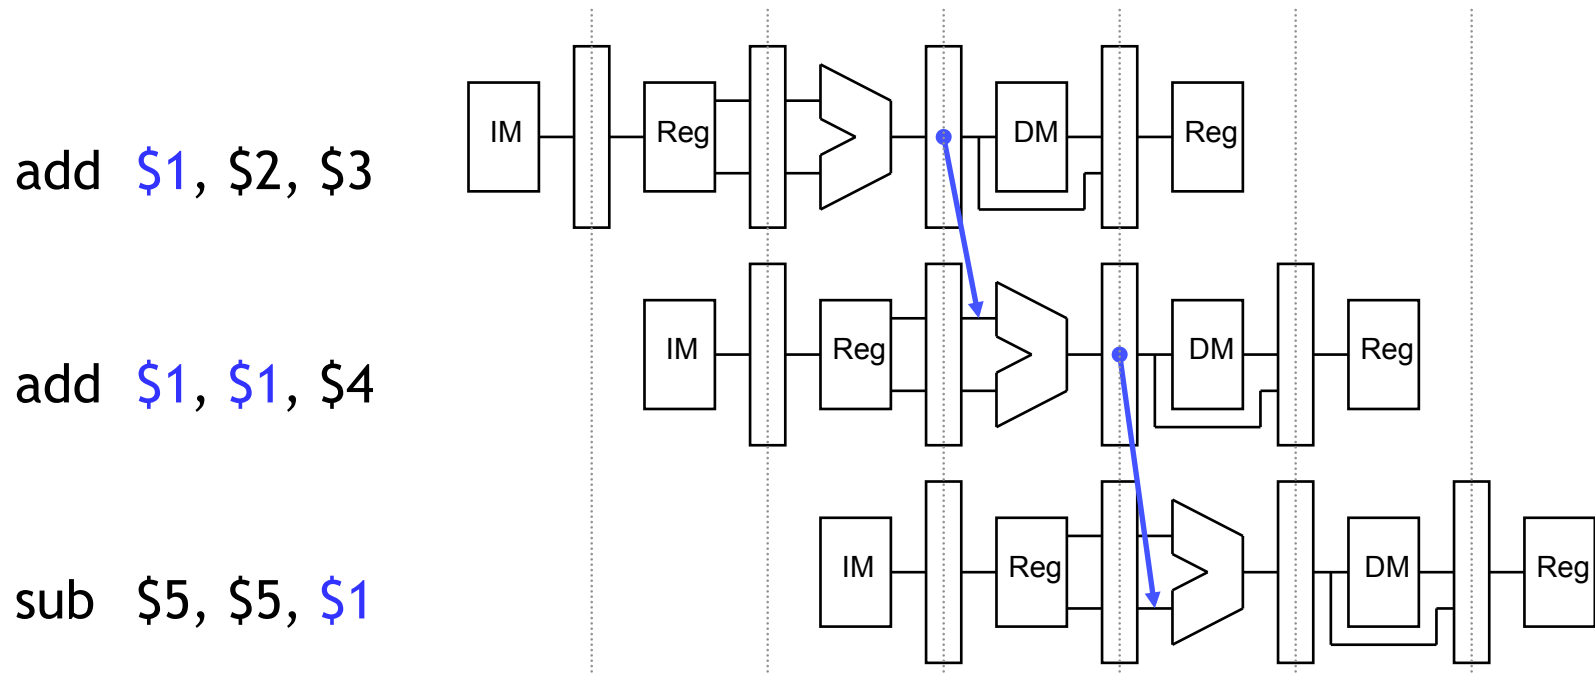


Detecting MEM/WB data hazards

- A MEM/WB hazard may occur between an instruction in the EX stage and the instruction from two cycles ago.
- One new problem is if a register is updated twice in a row.



- Register \$1 is written by *both* of the previous instructions; from which instruction should it receive its value?




Detecting MEM/WB data hazards

- A **MEM/WB hazard** may occur between an instruction in the EX stage and the instruction from *two* cycles ago.
- One new problem is if a register is updated twice in a row.

```
add  $1, $2, $3  
add  $1, $1, $4  
sub  $5, $5, $1
```

- Register \$1 is written by *both* of the previous instructions, but only the most recent result (from the second ADD) should be forwarded.



MEM/WB hazard equations

- Here is an equation for detecting and handling MEM/WB hazards for the first ALU source.

if ($\text{MEM/WB.RegWrite} = 1$
and $\text{MEM/WB.RegisterRd} = \text{ID/EX.RegisterRs}$

```
add $1, $2, $3  
add $1, $1, $4  
sub $5, $5, $1
```

then $\text{ForwardA} = 1$

- The second ALU operand is handled similarly.

if ($\text{MEM/WB.RegWrite} = 1$
and $\text{MEM/WB.RegisterRd} = \text{ID/EX.RegisterRt}$

then $\text{ForwardB} = 1$

MEM/WB hazard equations

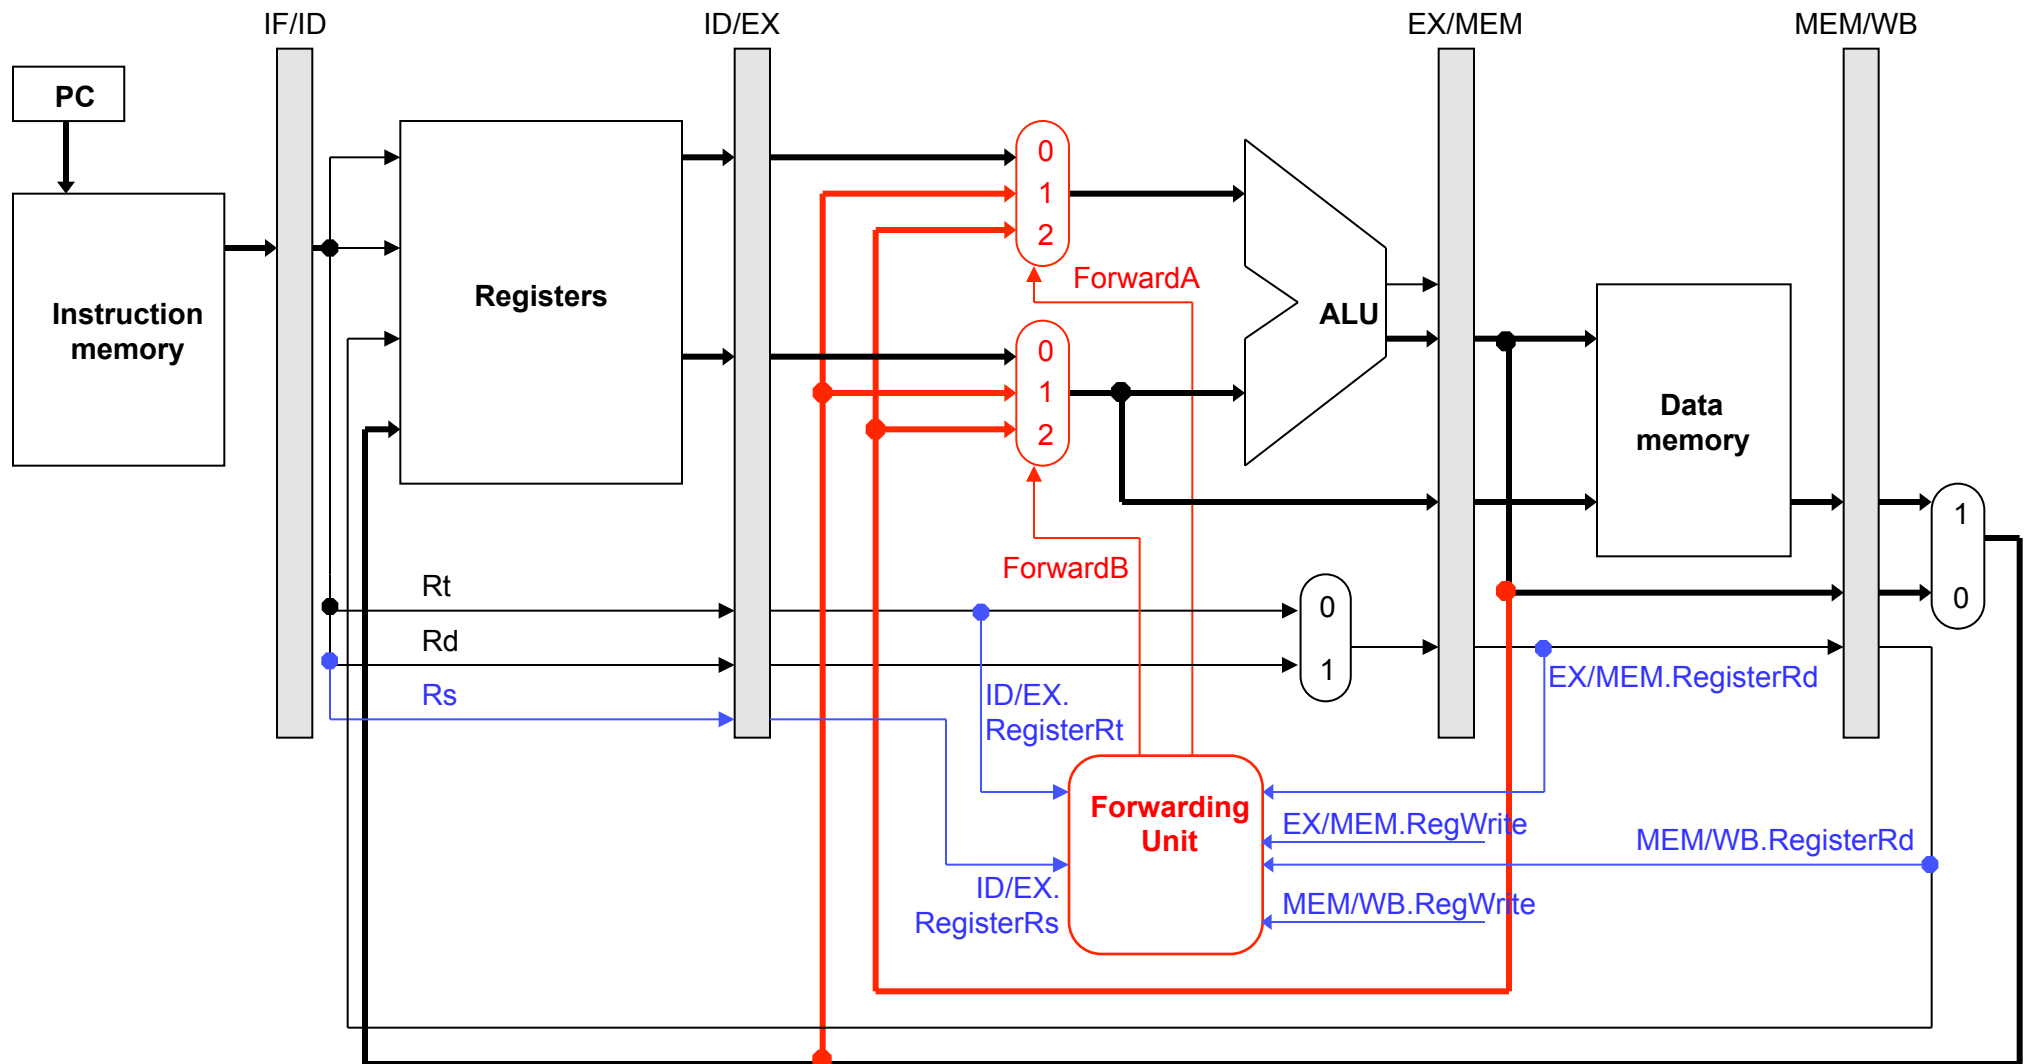
- Here is an equation for detecting and handling MEM/WB hazards for the first ALU source.

if ($\text{MEM/WB.RegWrite} = 1$
and $\text{MEM/WB.RegisterRd} = \text{ID/EX.RegisterRs}$
and ($\text{EX/MEM.RegisterRd} \neq \text{ID/EX.RegisterRs}$ or $\text{EX/MEM.RegWrite} = 0$)
then $\text{ForwardA} = 1$

- The second ALU operand is handled similarly.

if ($\text{MEM/WB.RegWrite} = 1$
and $\text{MEM/WB.RegisterRd} = \text{ID/EX.RegisterRt}$
and ($\text{EX/MEM.RegisterRd} \neq \text{ID/EX.RegisterRt}$ or $\text{EX/MEM.RegWrite} = 0$)
then $\text{ForwardB} = 1$

Simplified datapath with forwarding



The forwarding unit

- The forwarding unit has several control signals as inputs.

ID/EX.RegisterRs

EX/MEM.RegisterRd

MEM/WB.RegisterRd

ID/EX.RegisterRt

EX/MEM.RegWrite

MEM/WB.RegWrite

The forwarding unit outputs are selectors for the **ForwardA** and **ForwardB** multiplexers attached to the ALU. These outputs are generated from the inputs using the equations on the previous pages.

- Some new buses route data from pipeline registers to the new muxes.

Example

```
sub    $2, $1, $3
and    $12, $2, $5
or     $13, $6, $2
add    $14, $2, $2
sw     $15, 100($2)
```

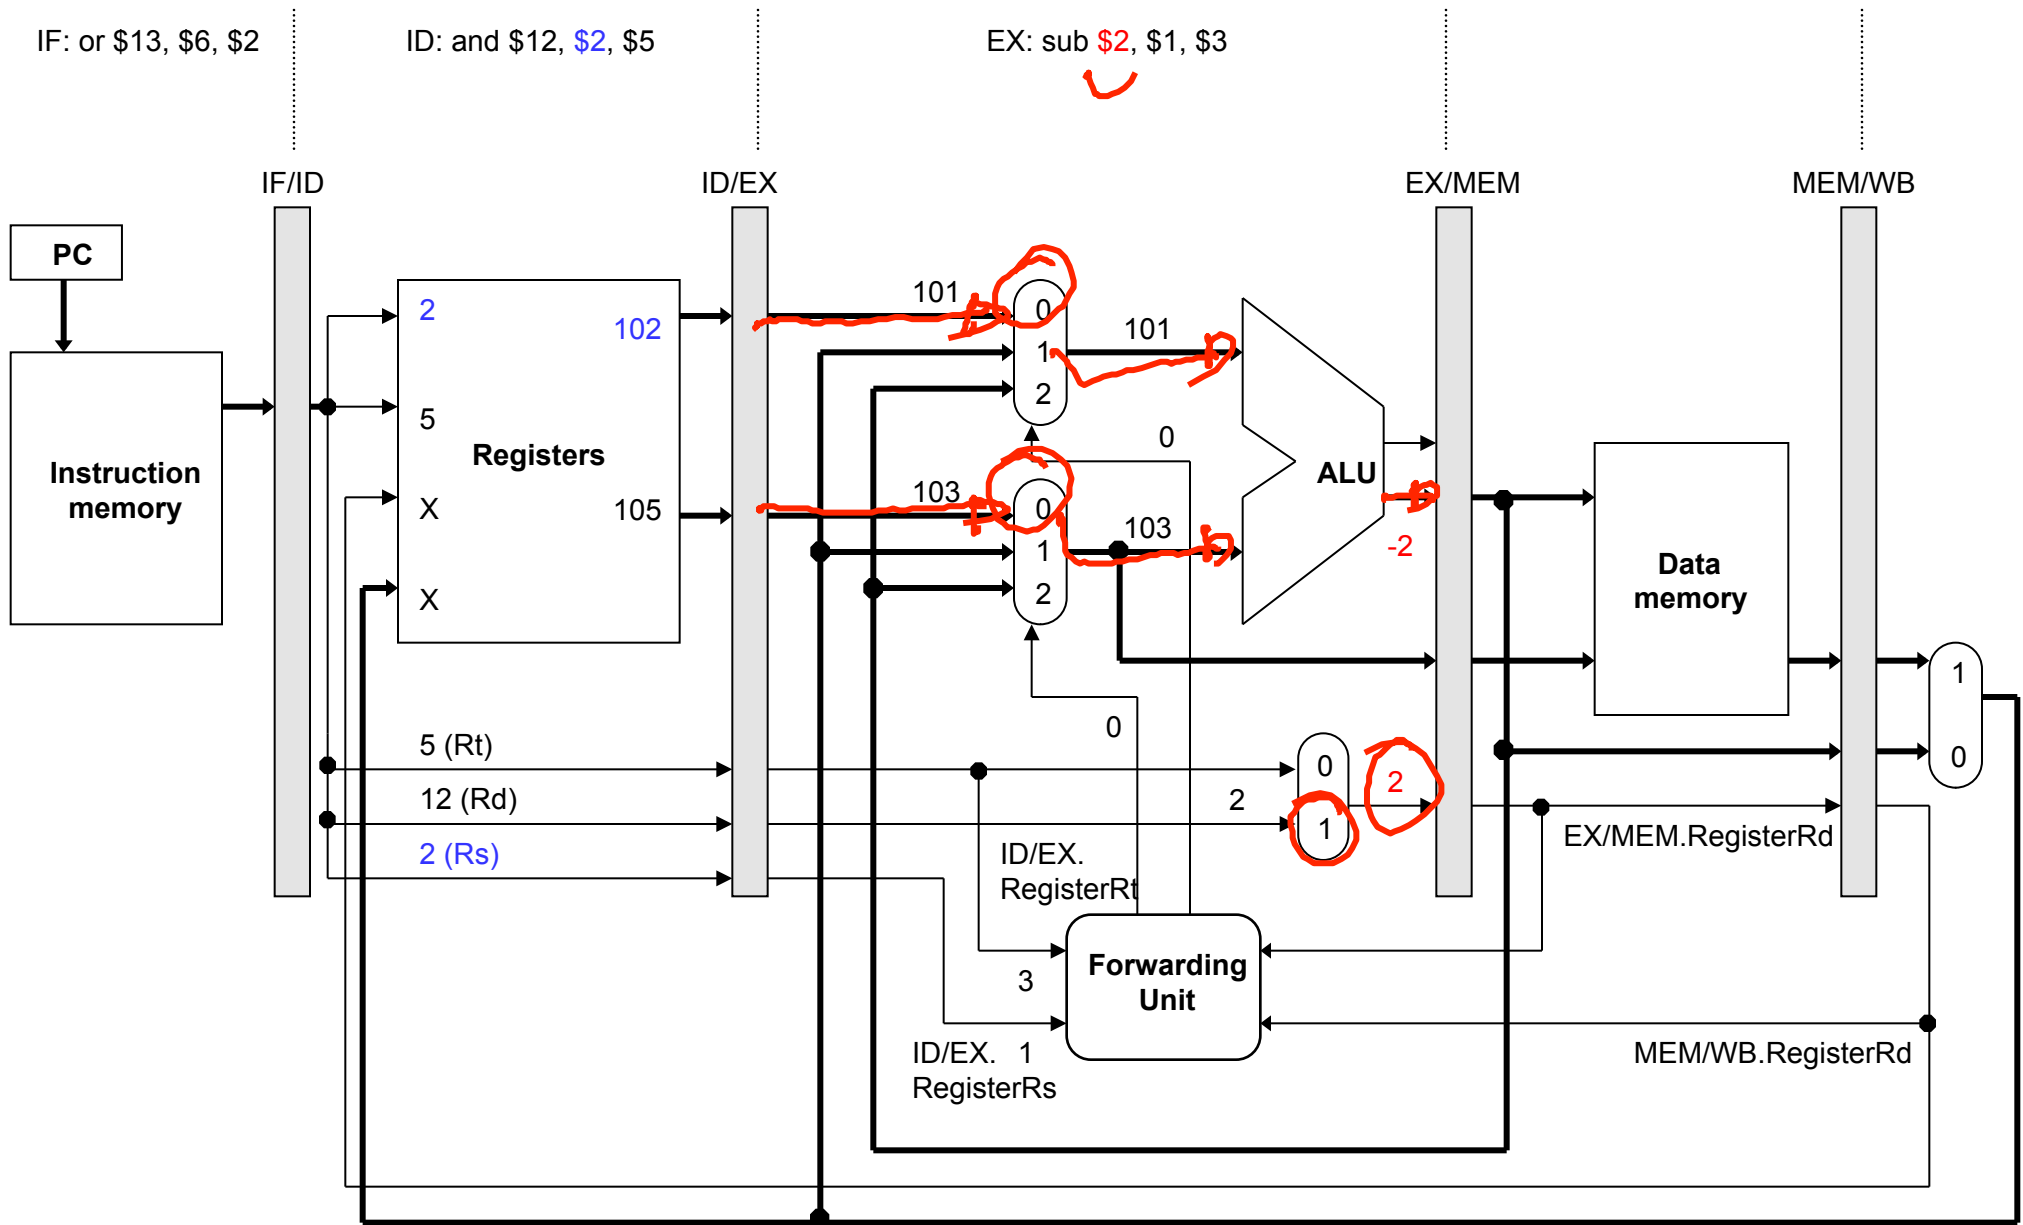
- Assume again each register initially contains its number plus 100.
 - After the first instruction, \$2 should contain -2 (101 - 103).
 - The other instructions should all use -2 as one of their operands.
- We'll try to keep the example short.
 - Assume no forwarding is needed except for register \$2.
 - We'll skip the first two cycles, since they're the same as before.

Clock cycle 3

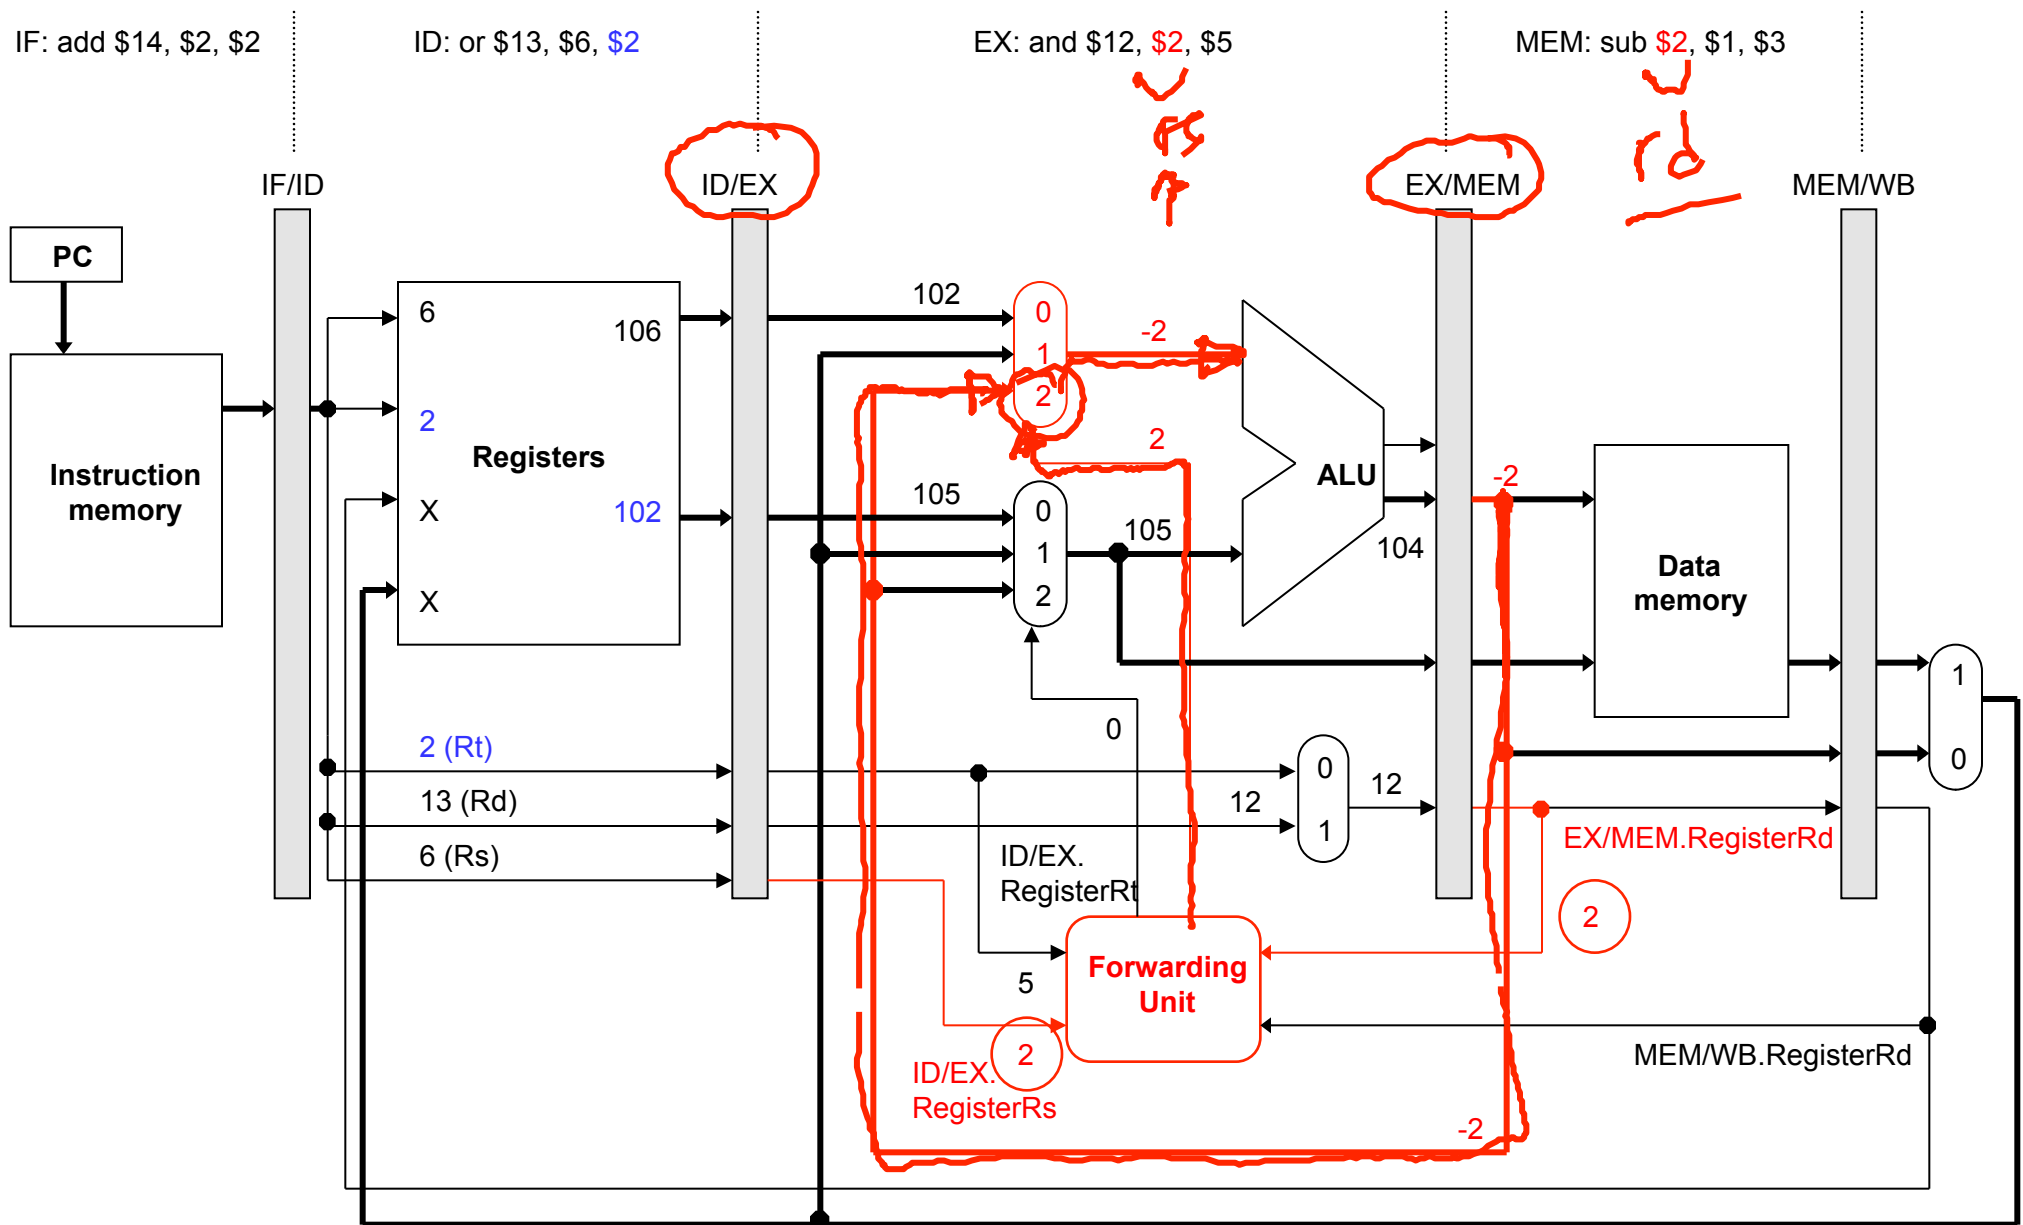
IF: or \$13, \$6, \$2

ID: and \$12, \$2, \$5

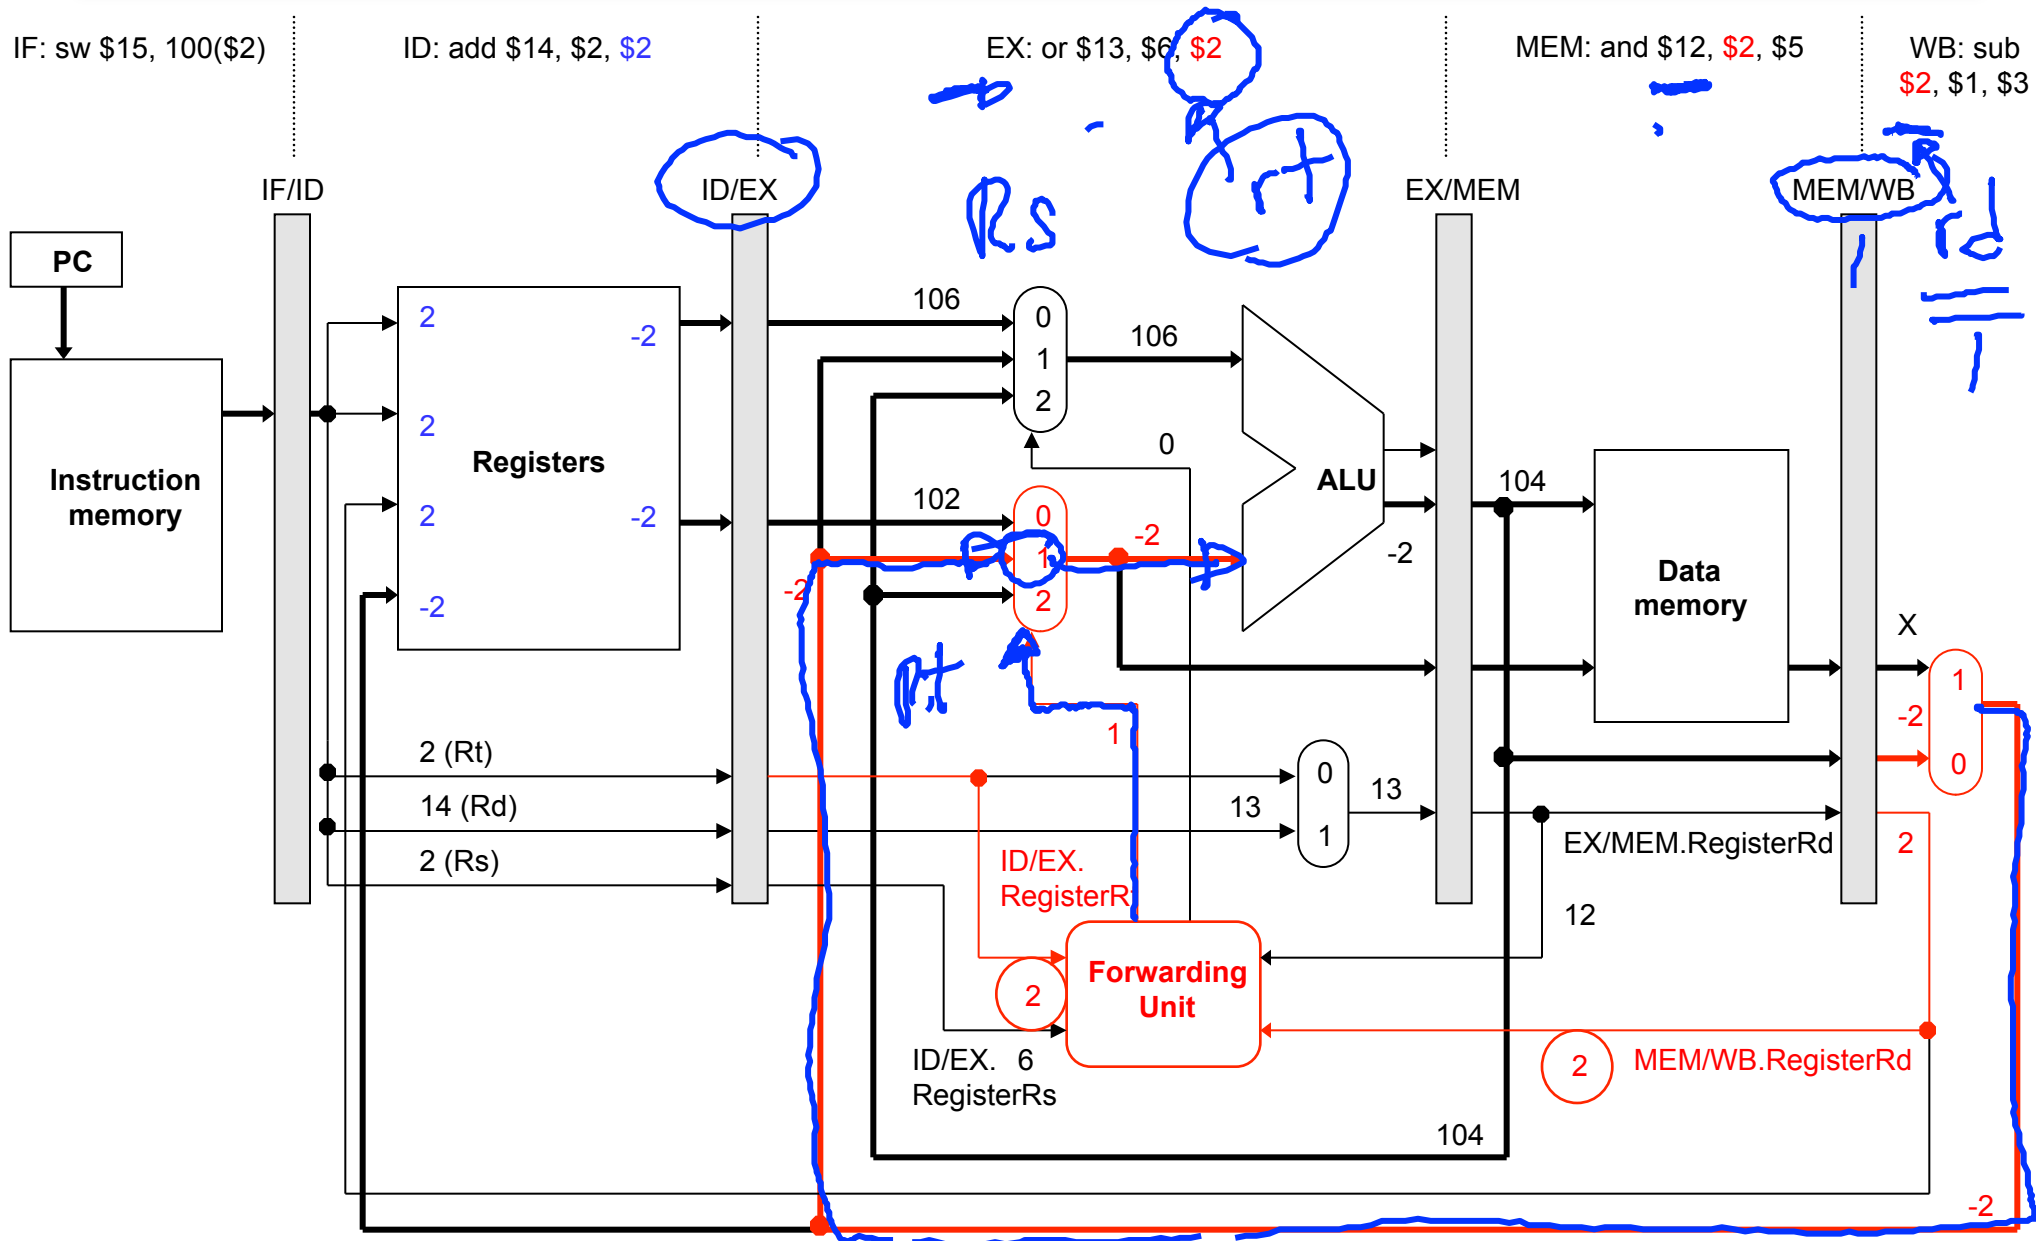
EX: sub \$2, \$1, \$3



Clock cycle 4: forwarding \$2 from EX/MEM



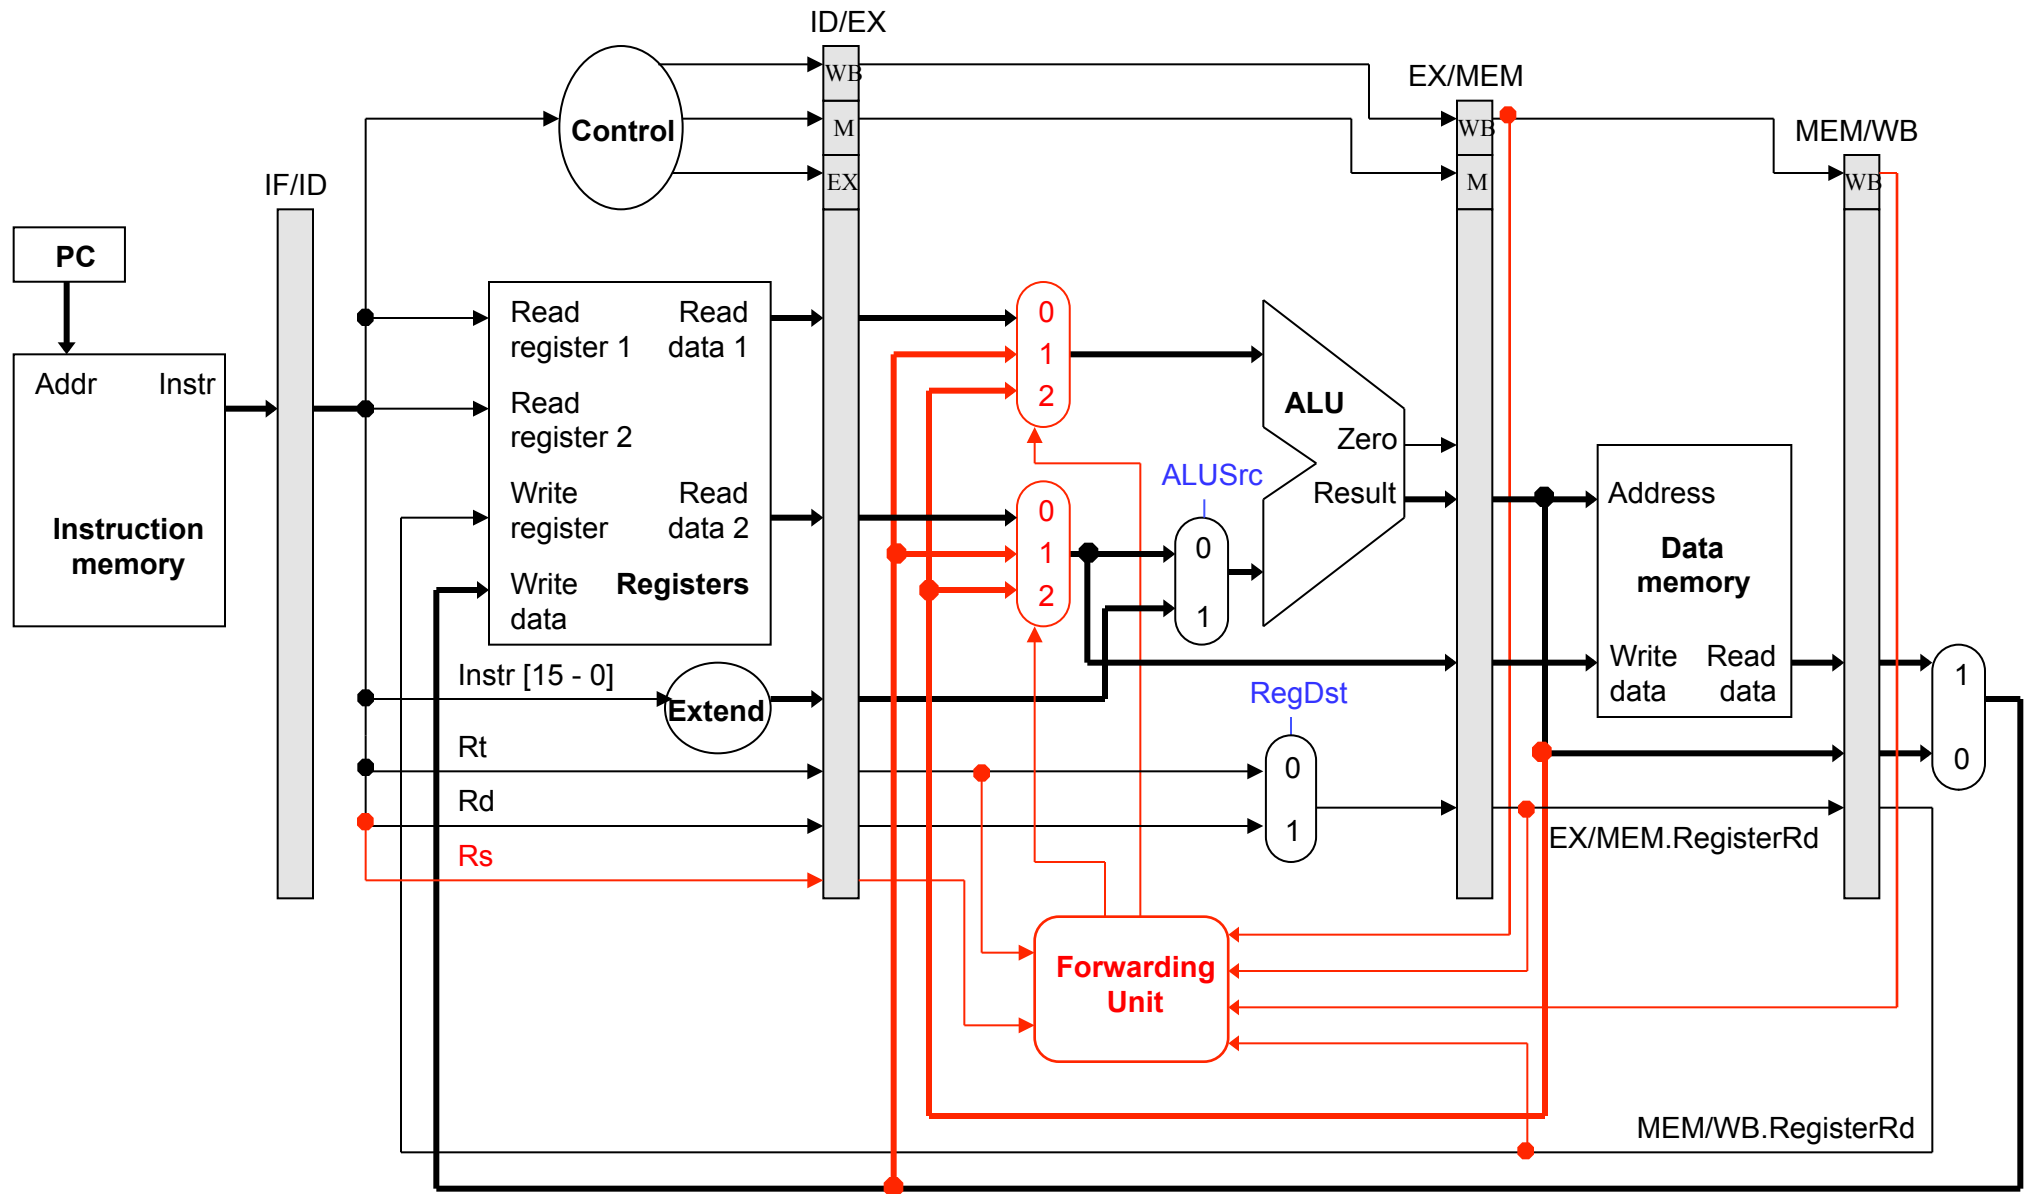
Clock cycle 5: forwarding \$2 from MEM/WB



Lots of data hazards

- The first data hazard occurs during cycle 4.
 - The forwarding unit notices that the ALU's first source register for the AND is also the destination of the SUB instruction.
 - The correct value is forwarded from the EX/MEM register, overriding the incorrect old value still in the register file.
- A second hazard occurs during clock cycle 5.
 - The ALU's second source (for OR) is the SUB destination again.
 - This time, the value has to be forwarded from the MEM/WB pipeline register instead.
- There are no other hazards involving the SUB instruction.
 - During cycle 5, SUB writes its result back into register \$2.
 - The ADD instruction can read this new value from the register file in the same cycle.

Complete pipelined datapath...so far




What about stores?

- Two “easy” cases:


add \$1, \$2, \$3

sw \$4, 0(\$1)

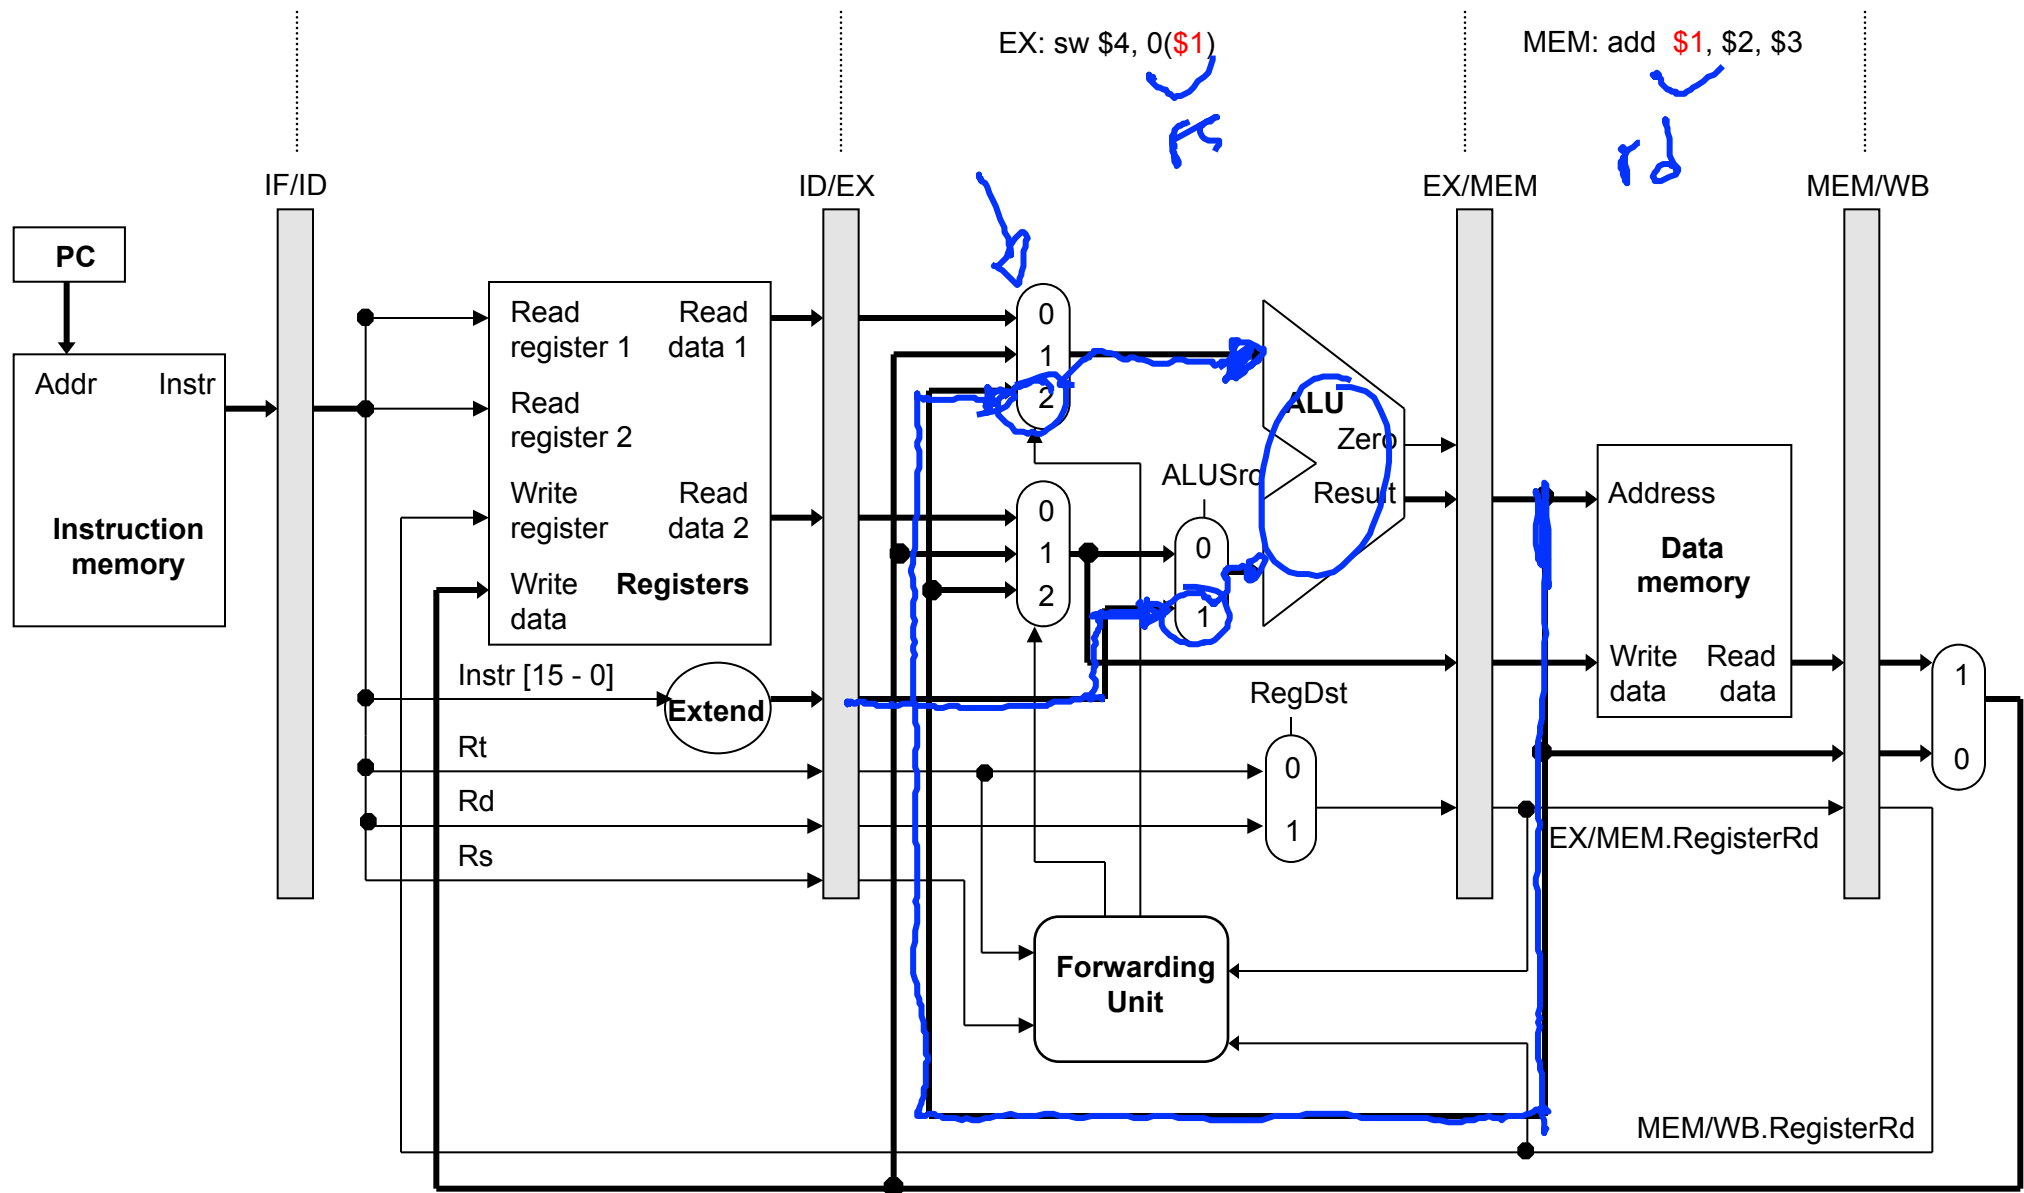


→ add \$1, \$2, \$3

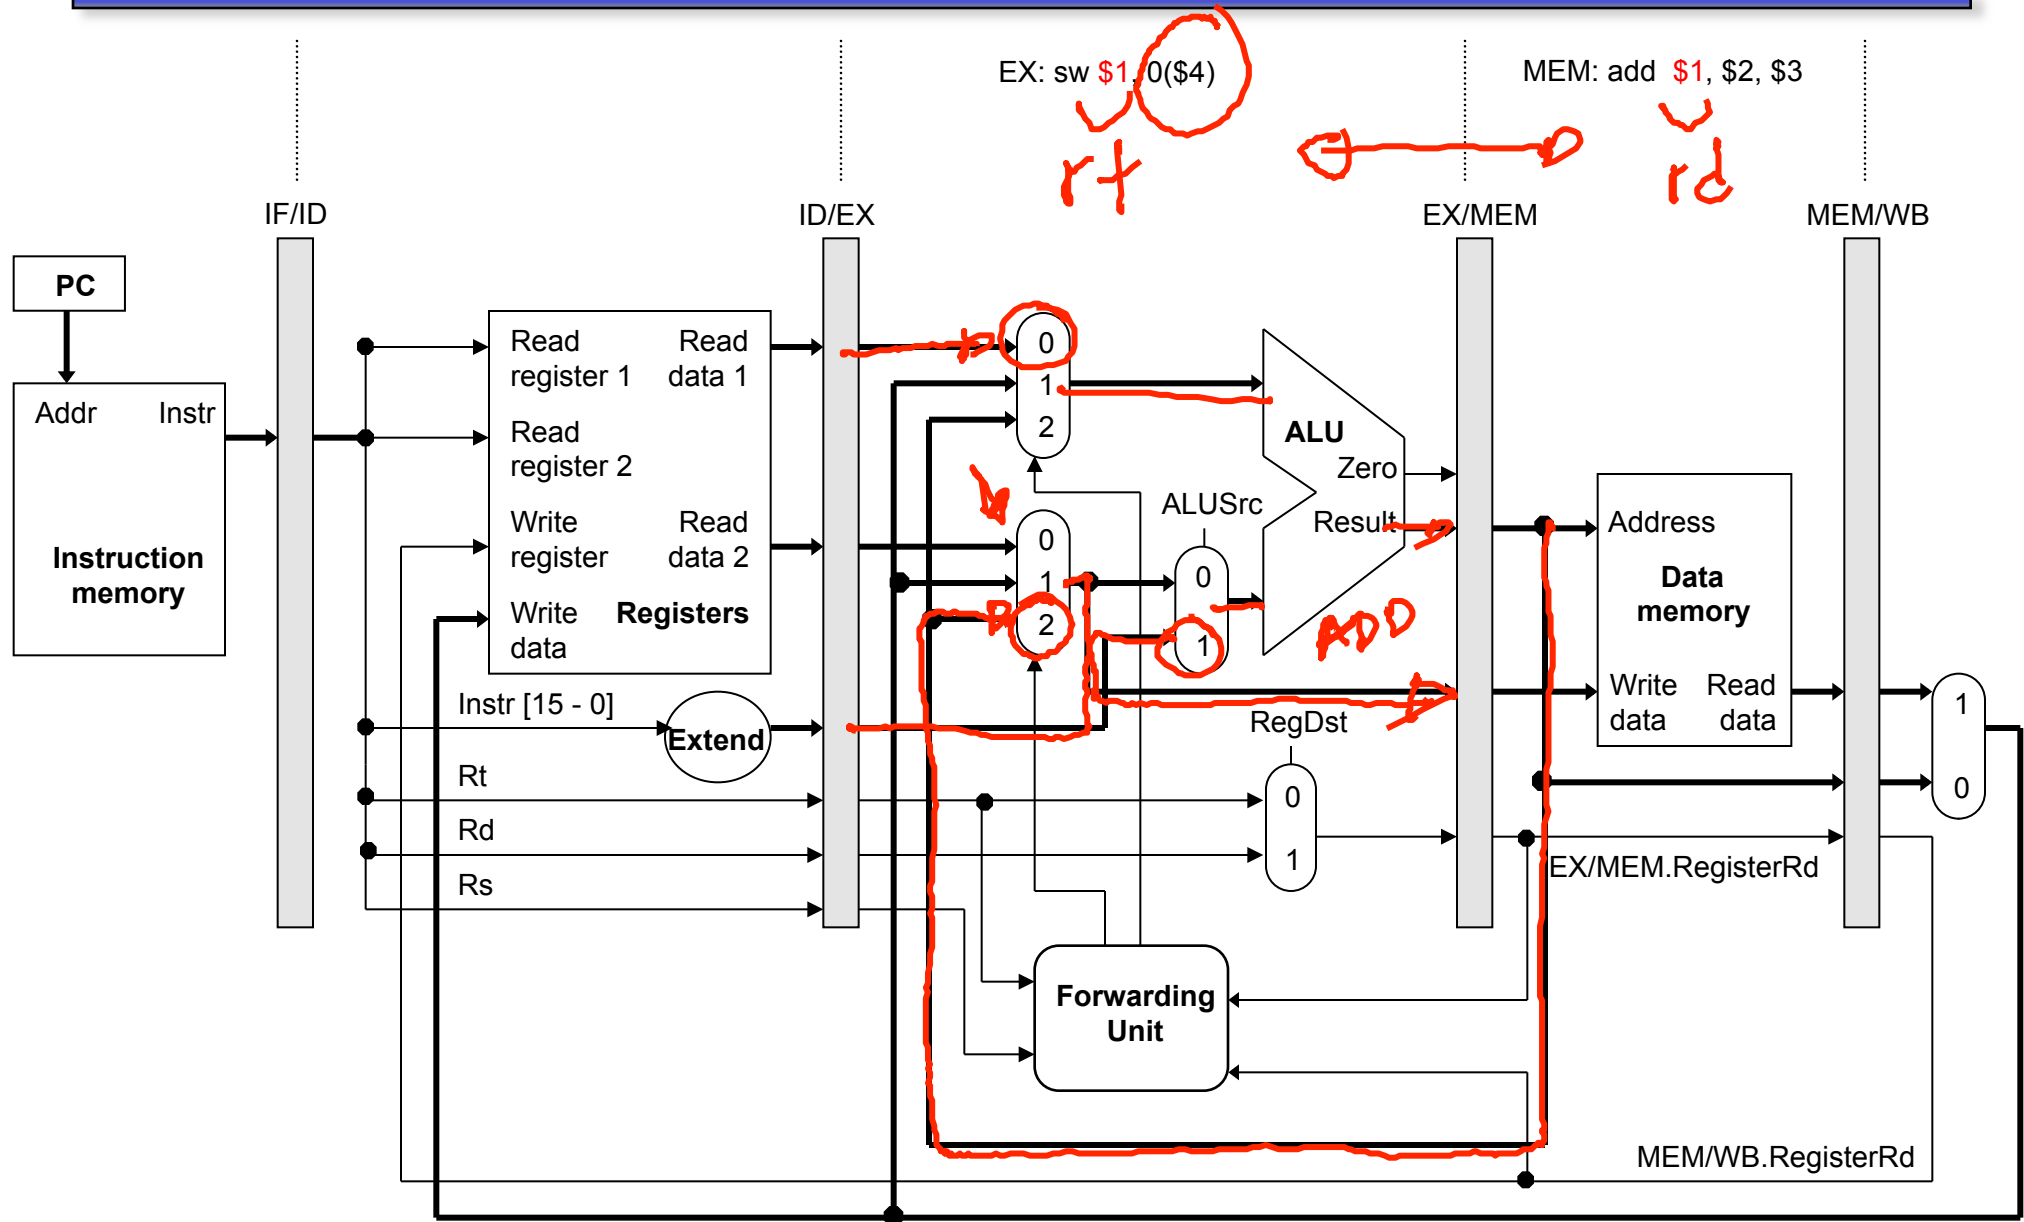
sw \$1, 0(\$4)



Store Bypassing: Version 1

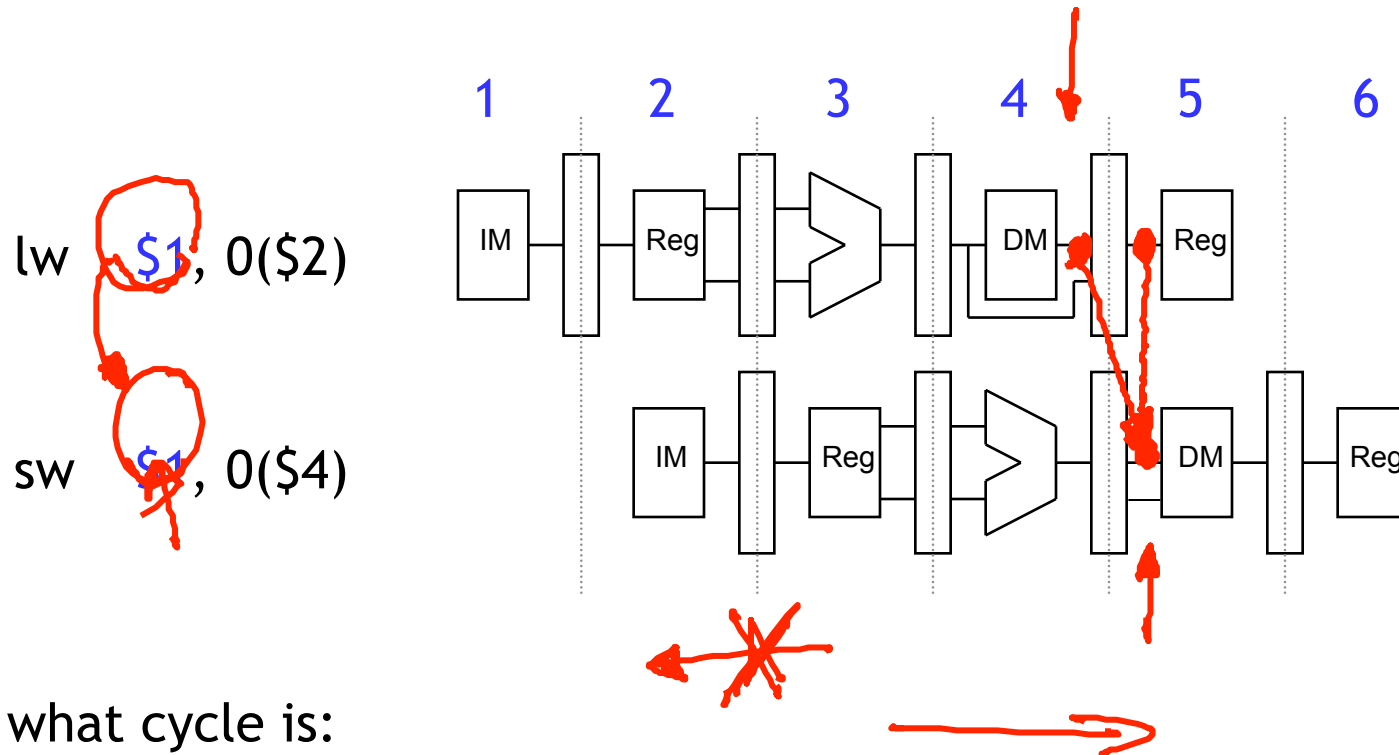


Store Bypassing: Version 2



What about stores?

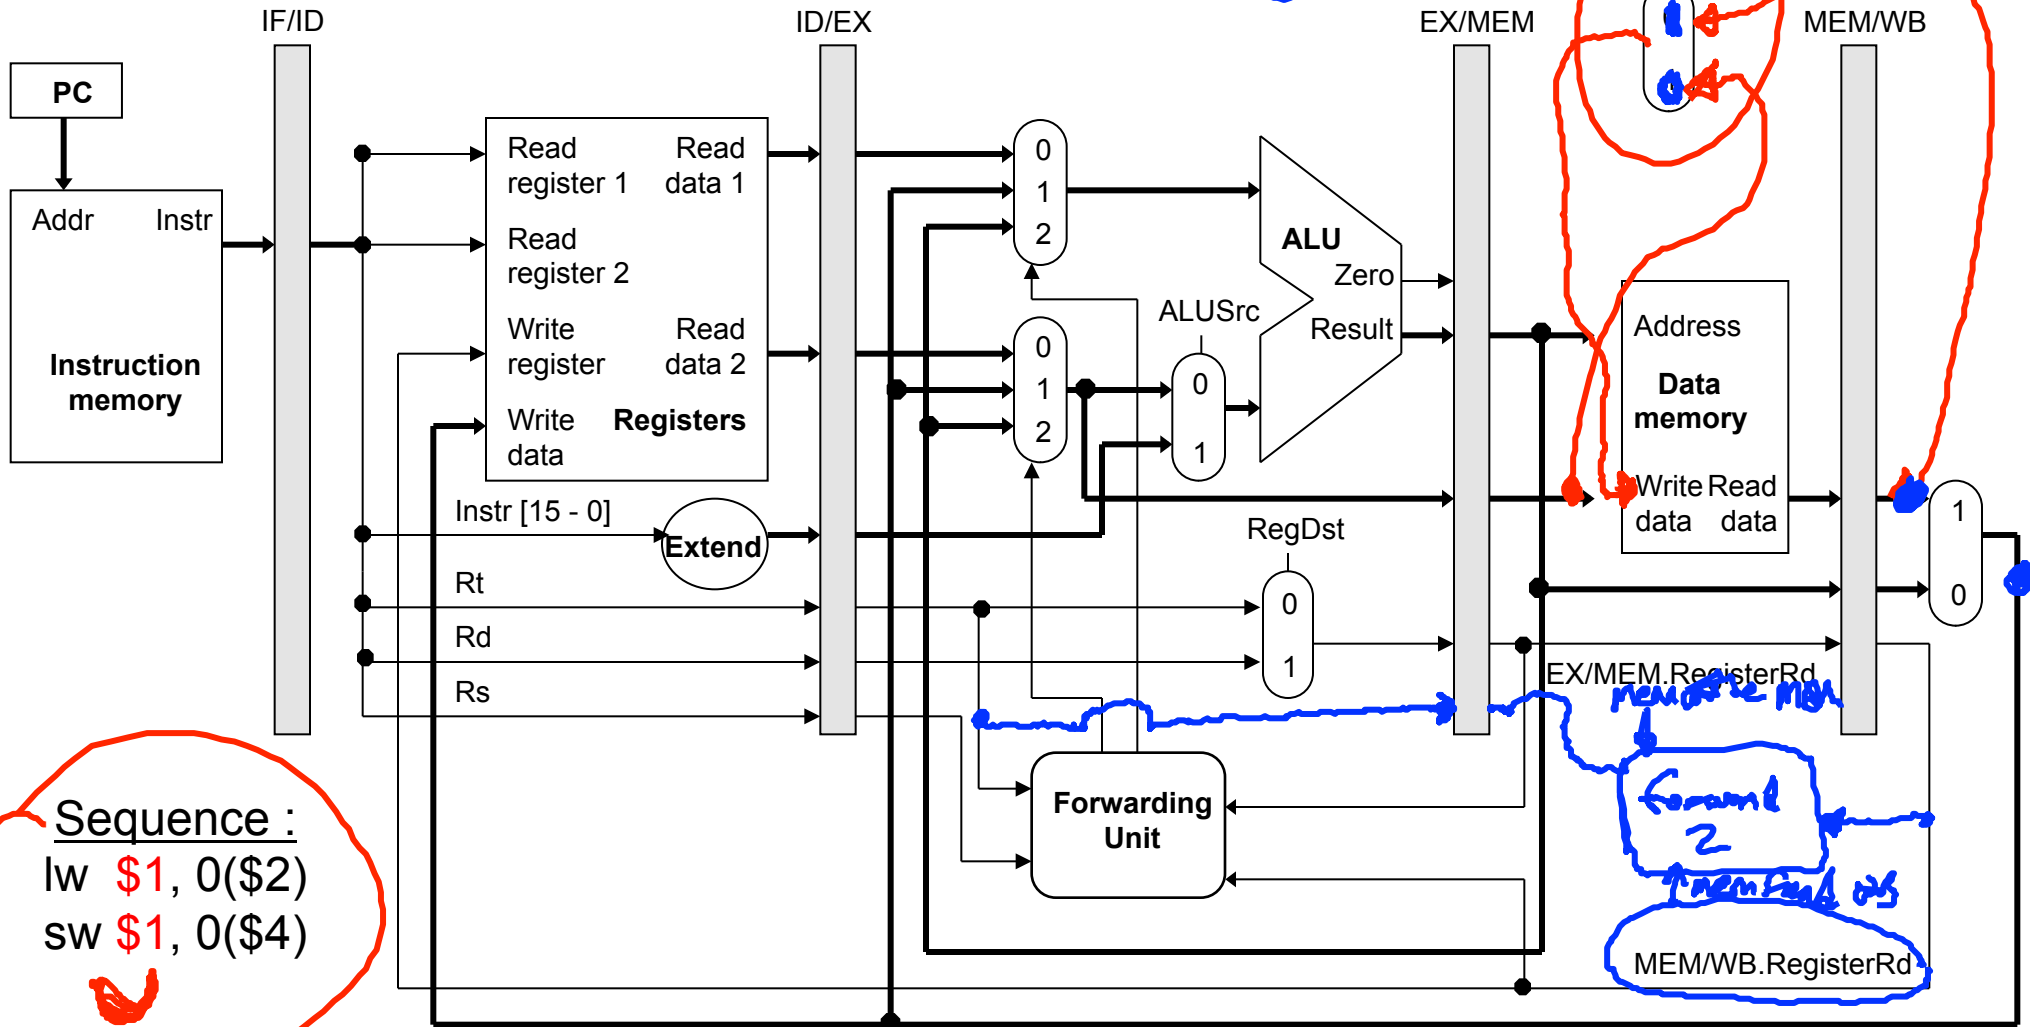
- A harder case:



- In what cycle is:
 - The load value available?
 - The store value needed?
- What do we have to add to the datapath?

Load/Store Bypassing: Extend the Datapath

FORWARD C = memWrite_MEM && MEM_READ_WB &&
 (MEM/WB.RegisterRd <= EX/MEM.RegisterRd) SW Lw



Sequence :
 lw \$1, 0(\$2)
 sw \$1, 0(\$4)

Miscellaneous comments

- Each MIPS instruction writes to at most one register.
 - This makes the forwarding hardware easier to design, since there is only one destination register that ever needs to be forwarded.
- Forwarding is especially important with deep pipelines like the ones in all current PC processors.
- Section 6.4 of the textbook has some additional material not shown here.
 - Their hazard detection equations also ensure that the source register is not \$0, which can never be modified.
 - There is a more complex example of forwarding, with several cases covered. Take a look at it!

Summary

- In real code, most instructions are dependent upon other ones.
 - This can lead to **data hazards** in our original pipelined datapath.
 - Instructions can't write back to the register file soon enough for the next two instructions to read.
- **Forwarding** eliminates data hazards involving arithmetic instructions.
 - The forwarding unit detects hazards by comparing the destination registers of previous instructions to the source registers of the current instruction.
 - Hazards are avoided by grabbing results from the pipeline registers *before* they are written back to the register file.
- Next time we'll finish up pipelining.
 - Forwarding can't save us in some cases involving lw.
 - We still haven't talked about branches for the pipelined datapath.

