


# Threads and Cache Coherence in Hardware

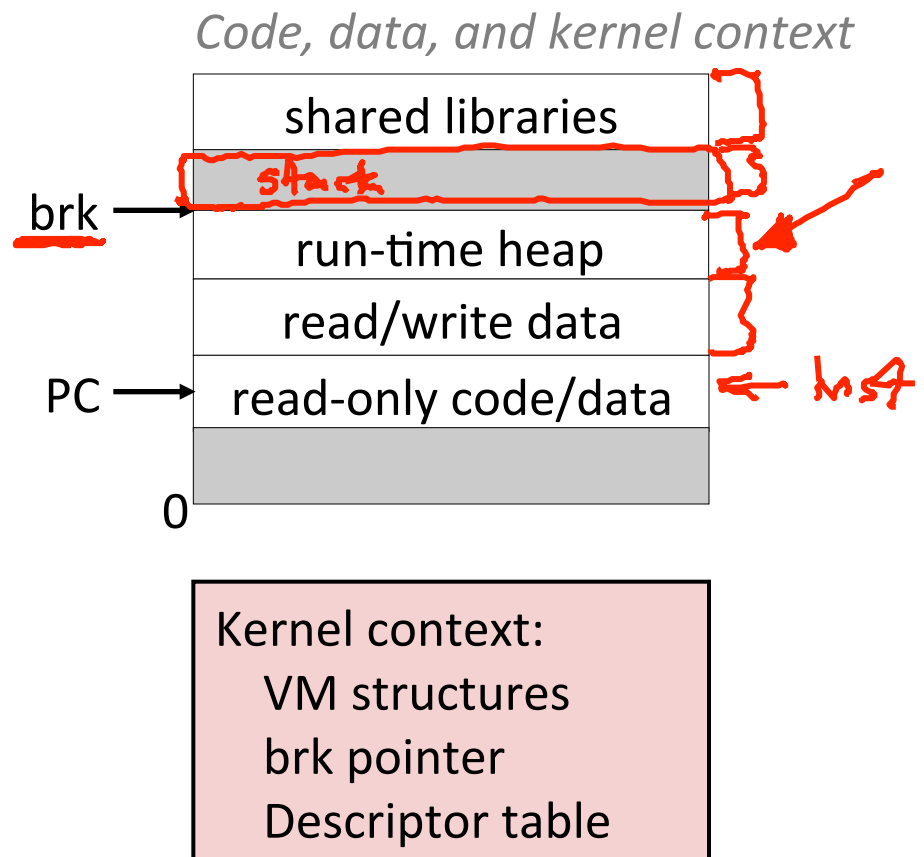
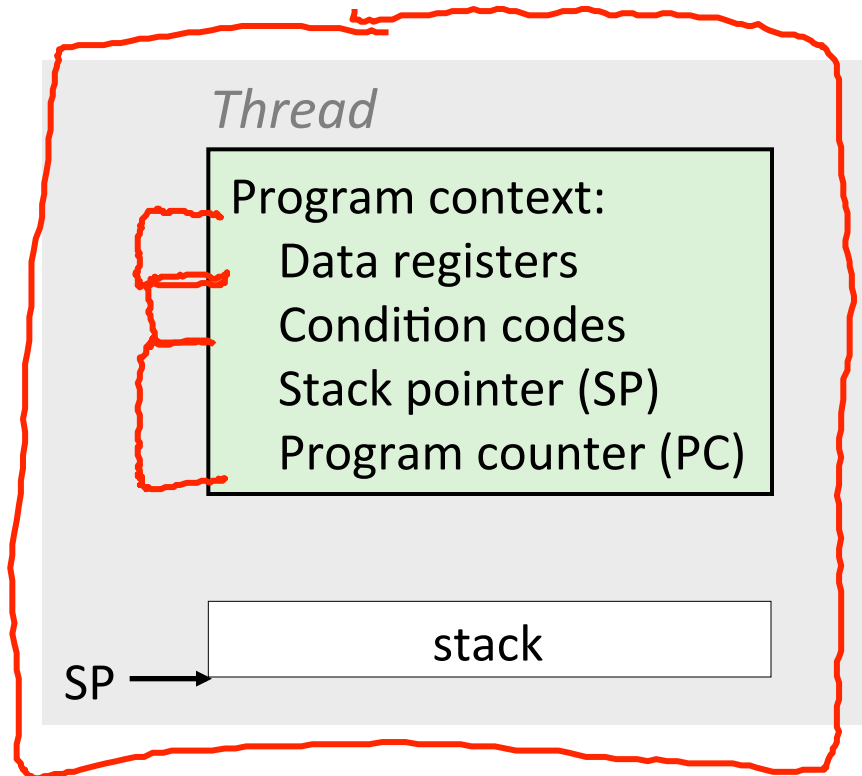
- Previously, we introduced multi-cores.
  - Today we'll look at issues related to multi-core memory systems:
    1. Threads
    2. Cache coherence



Intel Core i7

# Process View

- Process = thread + code, data, and kernel context



# Process with Two Threads

## Thread 1

Program context:

Data registers  
Condition codes  
Stack pointer (SP)  
Program counter (PC)




## Thread 2

Program context:

Data registers  
Condition codes  
Stack pointer (SP)  
Program counter (PC)




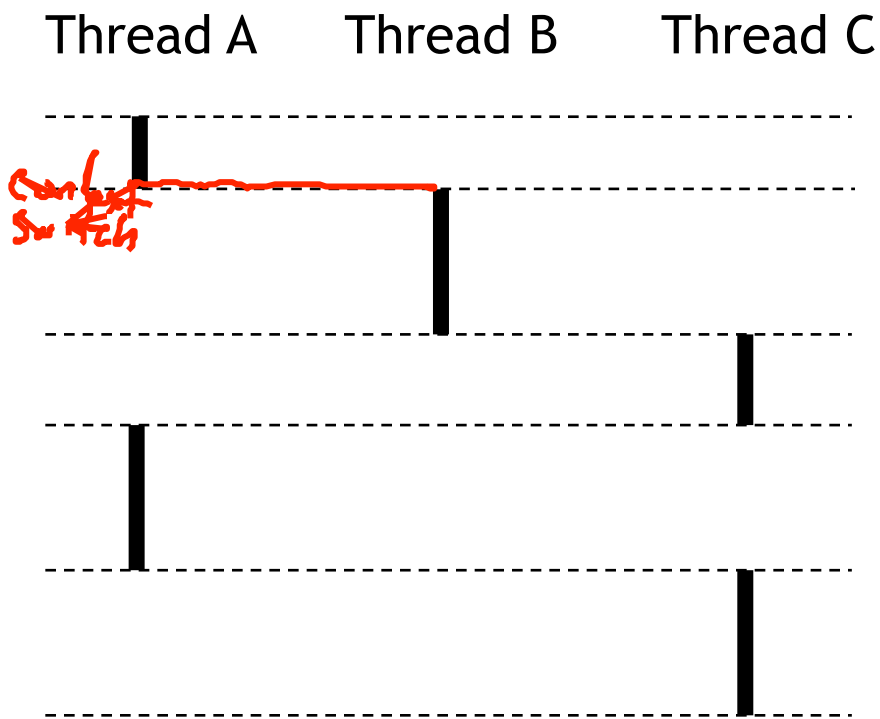
## Code, data, and kernel context



Kernel context:  
VM structures  
brk pointer  
Descriptor table

# Thread Execution

- Single Core Processor
  - Simulate concurrency by time slicing
- Multi-Core Processor
  - Can have true concurrency



# How do threads from the same process communicate?

---

- A) Their registers are kept in sync; when one thread writes a register, the written value is copied to the corresponding register in all cores that are running threads from the same process.
- B) Threads have distinct register files, but can read and write each other register files directly.
- C) Threads have distinct register files and must communicate through memory by one thread writing a memory location and later another thread reading that same memory location.
- D) Threads have distinct registers and memory and must communicate through network protocols like TCP/IP.

E) telepathy

# (Hardware) Shared Memory

---


- The programming model for multi-core CPUs

VA  $\rightarrow$  PA


What is it?

- Memory system with a single global physical address space
  - So processors can communicate by reading and writing same memory
- Design Goal 1: Minimize memory latency
  - Use co-location & caches
- Design Goal 2: Maximize memory bandwidth
  - Use parallelism & caches


# Some example memory-system designs




(a) Shared cache



(b) Bus-based shared memory

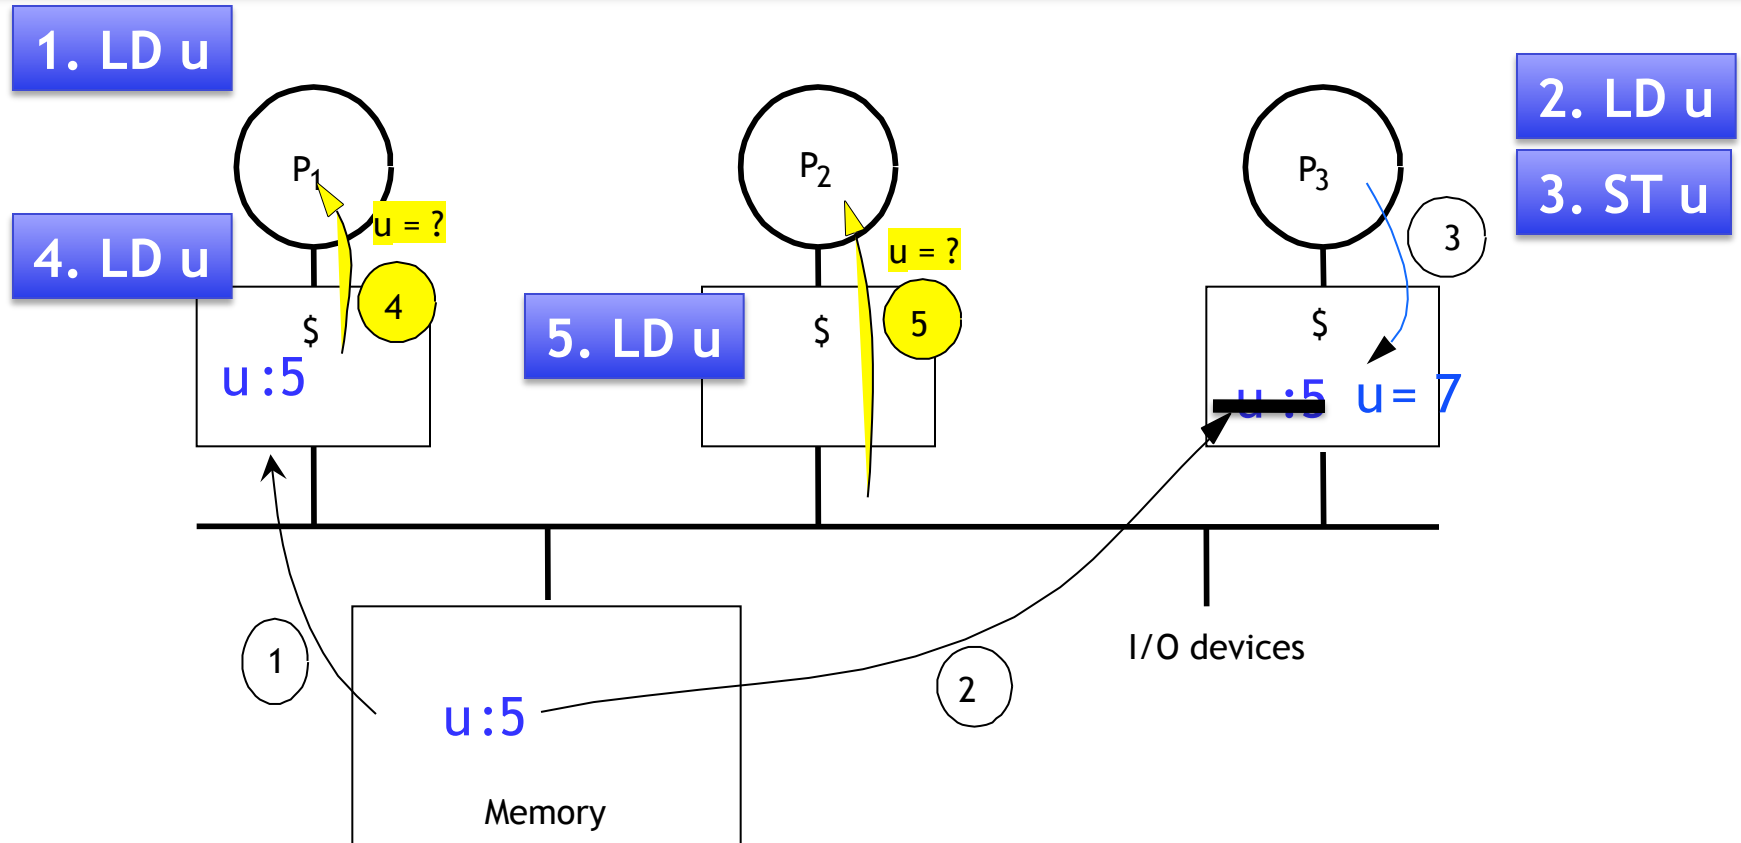


(c) Dancehall



(d) Distributed-memory

# The Cache Coherence Problem



- Multiple copies of a block can easily get inconsistent
  - Processor writes; I/O writes
- Processors could see different values for  $u$  after event 3



# HW Cache Coherence

- According to Webster's dictionary ...
    - **Cache**: a secure place of storage
    - **Coherent**: logically consistent
  - **Cache Coherence**: **keep storage logically consistent**
    - Coherence requires enforcement of 2 properties
1. Write propagation
    - All writes eventually become visible to other processors
  2. Write serialization
    - All processors see writes to same block in same order

P1  
x = 2

P2  
x = ?

P3  
x = 9

P4  
x = ?

# Cache Coherence Invariant

---

- Each block of memory is in exactly one of these 3 states:
  1. **Uncached:** Memory has the only copy
  2. **Writable:** Exactly 1 cache has the block and only that processor can write to it.
  3. **Read-only:** Any number of caches can hold the block, and their processors can read it.


only

**invariant** | in<sup>l</sup>ve(ə)rēənt |

noun Mathematics

a function, quantity, or property that remains unchanged when a specified transformation is applied.

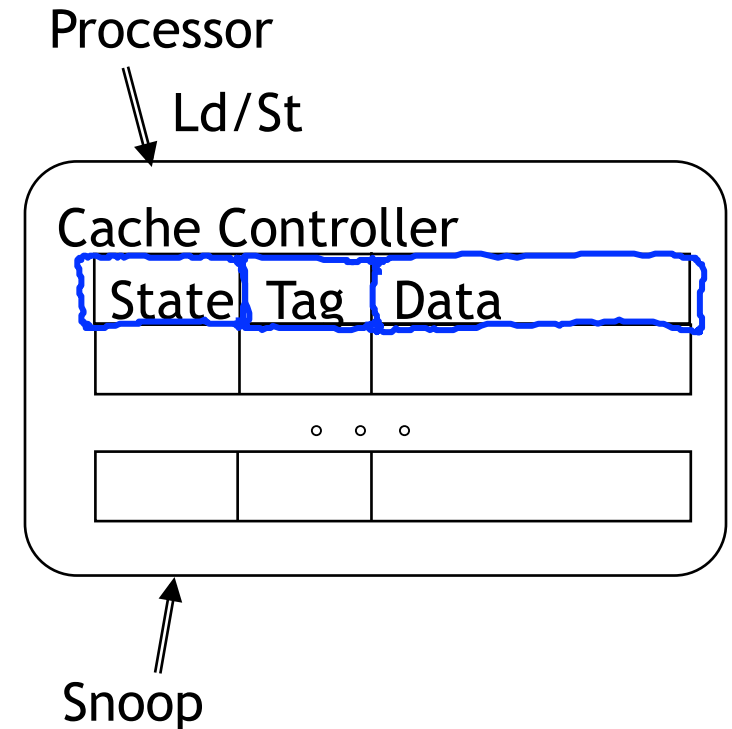
# Snoopy Cache Coherence Schemes



- Bus is a broadcast medium & caches know what they have
  - Cache controller “snoops” all transactions on the shared bus
    - Relevant transaction if for a block it contains
    - Take action to ensure coherence
      - Invalidate or supply value
- Depends on state of the block and the protocol

# Maintain the invariant by tracking “state”

- Every cache block has an associated state
  - This will supplant the valid and dirty bits
- A cache controller updates the state of blocks in response to processor and snoop events and generates bus transactions
- Snoopy protocol
  - set of states
  - state-transition diagram
  - actions



# MSI protocol

---

This is the simplest possible protocol, corresponding directly to the 3 options in our invariant

- **Invalid State**: the data in the cache is not valid.
- **Shared State**: multiple caches potentially have copies of this data; they will all have it in **shared** state. Memory has a copy that is consistent with the cached copy.
- **Dirty or Modified**: only 1 cache has a copy. Memory has a copy that is inconsistent with the cached copy. Memory needs to be updated when the data is displaced from the cache or another processor wants to read the same data.

# Actions

---

## Processor Actions:

- Load **A**
- Store **A**
- Eviction: processor wants to replace cache block

## Bus Actions:

- **GETS**: request to get data in shared state
- **GETX**: request for data in modified state (*i.e.*, eXclusive access)
- **UPGRADE**: request for exclusive access to data owned in shared state

## Cache Controller Actions:

- **Source Data**: this cache provides the data to the requesting cache
- **Writeback**: this cache updates the block in memory

# MSI Protocol

