

schema element \Rightarrow create table R (...)
Example: Trigger Definition

E CREATE TRIGGER BeerTrig
AFTER INSERT ON Sells
REFERENCING NEW ROW AS NewTuple
FOR EACH ROW
C WHEN (NewTuple.beer NOT IN
(SELECT name FROM Beers))
A INSERT INTO Beers(name)
VALUES(NewTuple.beer);

The event

The condition

The action

Sells

Beers

Bud	NULL

Tuple-level What happen?
if no, \Rightarrow reject.

Options: The Event

- AFTER can be BEFORE.
 - Also, INSTEAD OF if the relation is a view.
 - A great way to execute view modifications: have triggers translate them to appropriate modifications on the base tables.
- INSERT can be DELETE or UPDATE.
 - And UPDATE can be UPDATE . . . ON a particular attribute.

insert into sells values (..., ..., ...)

Options: **FOR EACH ROW**

Tuple-level

- Triggers are either *row-level* or *statement-level*.
- FOR EACH ROW indicates row-level; its absence indicates statement-level.
- Row level triggers are executed once for each modified tuple.

if no

- **Statement-level** triggers execute once for an SQL statement, regardless of how many tuples are modified.

the SQL stmt.

Stmt-level:

INSERT/DELETE/UPDATE = OLD/NEW table.

Options: REFERENCING
Tuple-Level (for each row)

- INSERT statements imply a new tuple ~~(for row-level) or new set of tuples (for statement level)~~.
- DELETE implies an old tuple ~~or table~~.
- UPDATE implies both. old/new
- Refer to these by

[NEW OLD][~~TUPLE~~ TABLE] AS <name>

referencing ^{ROW} OLD table as OLDT.

Options: The Condition

- Any boolean-valued condition is appropriate.
- It is evaluated before or after the triggering event, depending on whether BEFORE or AFTER is used in the event.
- Access the new/old tuple or set of tuples through the names declared in the REFERENCING clause. (or fixed by “OLD”, “NEW” in MySQL.)

+ tables

Options: The Action

- There can be more than one SQL statement in the action.
 - Surround by BEGIN . . . END if there is more than one.
- But queries make no sense in an action, so we are really limited to modifications.

Another Example

- Using Sells(bar, beer, price) and a unary relation RipoffBars(bar) created for the purpose, maintain a list of bars that raise the price of any beer by more than \$1.

The Trigger

```
CREATE TRIGGER PriceTrig
```

```
AFTER UPDATE OF price ON Sells
```

The event –
only changes
to prices

```
REFERENCING
```

```
OLD ROW as old
```

```
NEW ROW as new
```

Updates let us
talk about old
and new tuples

Condition:
a raise in
price > \$1

```
FOR EACH ROW
```

We need to consider
each price change

```
WHEN (new.price > old.price + 1.00)
```

```
INSERT INTO RipoffBars
```

```
VALUES(new.bar);
```

When the price change
is great enough, add
the bar to RipoffBars

Behind the Scene: Why Trigger was invented?

Aspects of a trigger subsystem in an integrated database system. Proceedings of the 2nd international conference on Software engineering.
1976.

1. Extended assertions. (why?)
2. ??

Behind the Scene: This is why...

Aspects of a Trigger Subsystem in an Integrated Database System

by

Kapali P. Eswaran
IBM Research Laboratory
San Jose



①

ABSTRACT. This paper considers the specifications and design of a trigger subsystem in a database management system. The use of triggers as extended assertions and as a means to materialize virtual data objects are discussed. The functional requirements of a trigger subsystem and different implementation issues are studied. We also examine the relationships between a trigger subsystem and the rest of the database system, in particular the authorization and locking subsystems.

②

Triggers on Views

- Generally, it is impossible to modify a view, because it doesn't exist.
- But an **INSTEAD OF** trigger lets us interpret view modifications in a way that makes sense.
- Example: We'll design a view **Synergy** that has (drinker, beer, bar) triples such that the bar serves the beer, the drinker frequents the bar and likes the beer.

('Barb', 'hello', 'joe bar')

Example: The View

(drinker, beer, bar)

CREATE VIEW Synergy AS

SELECT Likes.drinker, Likes.beer, Sells.bar

FROM Likes, Sells, Frequents
WHERE Likes.drinker = Frequents.drinker
AND Likes.beer = Sells.beer
AND Sells.bar = Frequents.bar;

Natural join of Likes,
Sells, and Frequents

Likes
(Dr, Bar)

Freq
(Dr, Bar)

Sells
(Bar, Beer, price)

Pick one copy of each attribute

//
?

Interpreting a View Insertion

- We cannot insert into Synergy --- it is a view.
- But we can use an INSTEAD OF trigger to turn a (drinker, beer, bar) triple into three insertions of projected pairs, one for each of Likes, Sells, and Frequents.
 - The Sells.price will have to be NULL.

The Trigger

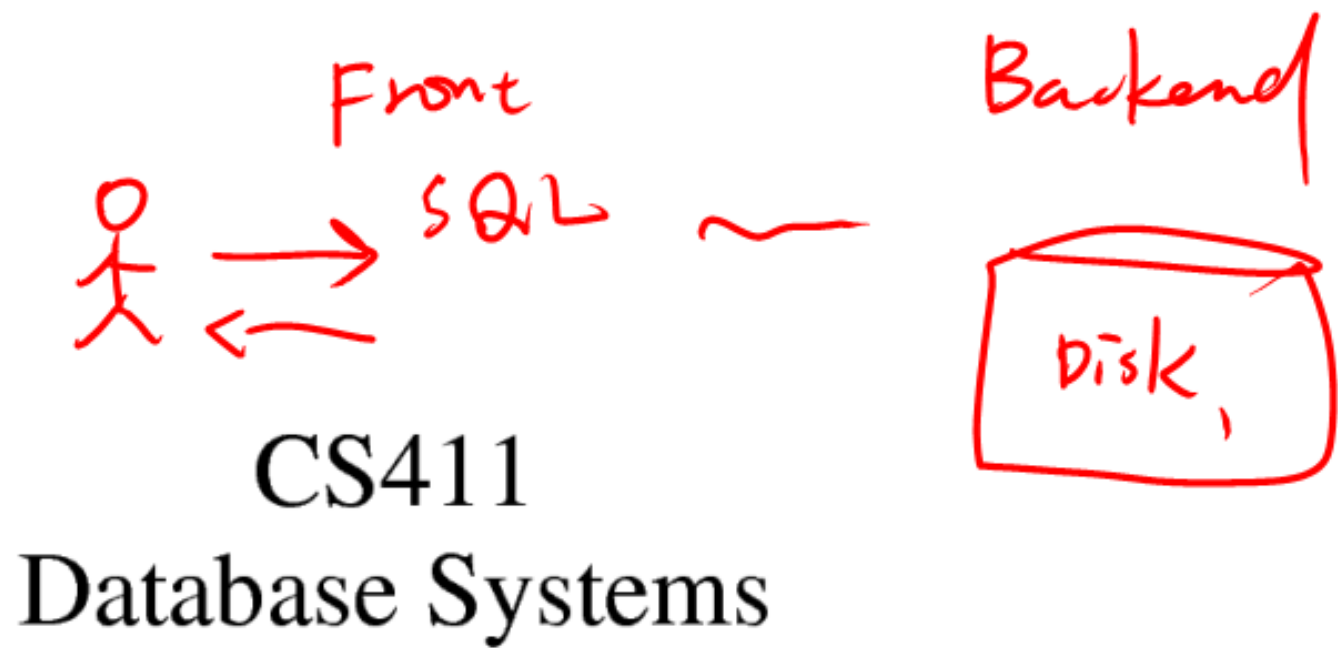
insert into synergy values ('Barb', 'Hello', joe
bar)

CREATE TRIGGER ViewTrig

E ✓ INSTEAD OF INSERT ON Synergy
REFERENCING NEW ROW AS n
FOR EACH ROW

BEGIN

{ INSERT INTO LIKES VALUES(n.drinker, n.beer);
INSERT INTO SELLS(bar, beer) VALUES(n.bar, n.beer); price = Null)
INSERT INTO FREQUENTS VALUES(n.drinker, n.bar);
END;



07: Indexing

Announce

~ One week off.

After that :

~ Special Topics:

Final ^{on} { Wed : Anatomy of SQLite.
Fri : SQL Tuning.

query does
sel/fro/where

Why Do We Learn This?

tuples,

of course,

Find out thing quickly

Index

✓ by values ↔ pth

GPA > 3.0
dept = "CS"
logical
declarative

phy
procedural

Indexing


- Indexing
 - types of indexes
 - B+ trees
 - hash tables

Q: What is “indexing”?

- To build an index.
- But what is an index?

*an organization of labeled ptrs to
collection of items.*

- Examples in the real world?



label
value

What is "indexing"?

("food", Exit 151)



Indexes

GPA 73.0
Dept = 'CS'

- An index on a file speeds up selections on the search key field(s)
- Search key = any subset of the fields of a relation *used in query cond.*
 - Search key is **not** the same as key (minimal set of fields that uniquely identify a record in a relation).
- Entries in an index: (k, r), where:
 - k = the key *3.2*
 - r = the record OR record id OR record ids *, r*


Types of Indexes

- Clustered/unclustered
 - Clustered = records sorted in the key order
 - Unclustered = no
- Dense/sparse
 - Dense = each record has an entry in the index
 - Sparse = only some records have
- Primary/secondary
 - Primary = on the primary key
 - Secondary = on any key
 - Some textbooks interpret these differently
- B+ tree / Hash table / ...

Sequential File before 1972 (B-tree invented)

Ex: Clustered, Dense Index

- Clustered: File is sorted on the index attribute
- Dense: every age value appear in Index: sequence of (key, pointer) pairs




sorted by Age

Clustered Sparse Index

- Sparse index: one key per data block

select
from
where age = 25

Dense




need more time to check.

Alex 20
Bob 20
:
20 yrs.

How if duplicate keys?

Clustered Index with Duplicate Keys

- **Dense** index: point to the first record with that key



Clustered Index with Duplicate Keys

- **Sparse** index: pointer to lowest search key in each block:

age = 10?


? ←

age = 20?

two way!


- OK?

Try search for 20



Clustered Index with Duplicate Keys

- Better: pointer to lowest new search key in each block (overflow)
- Search for 20



Unclustered Indexes

- Often for indexing other attributes than primary key
- Always dense (why ?) stud

age=20

✓
✓
x
x

10	
10	
20	
20	

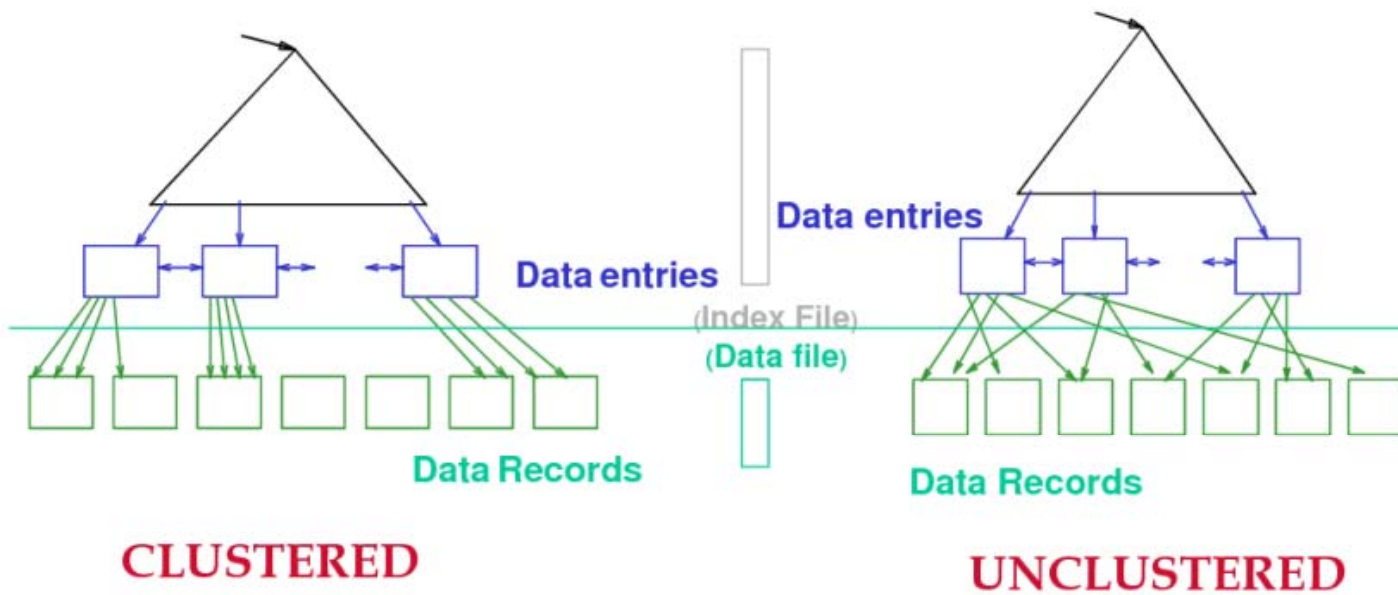
20	
30	
30	
30	

age Name

20	Alex
30	Bob
30	Carl
20	
10	
20	
10	
30	

locality
is
lost,

Summary Clustered vs. Unclustered Index




Composite Search Keys

- *Composite Search Keys*: Search on a combination of fields.

- Equality query: Every field value is equal to a constant value. E.g. wrt $\langle \text{sal}, \text{age} \rangle$ index:
 - age=20 and sal =75
- Range query: Some field value is not a constant. E.g.:
 - age =20; or age=20 and sal > 10


Examples of composite key indexes using lexicographic order.



Q: Our textbook as example: Indexes?



front Back
2 TOC Topic.



- How many indexes? Where?
- What are keys? What are records?
- Clustered?
- Dense?

X • Primary?

	TOC	Topic
key	chap/sect#	topic
records	page	page
clustered	✓	✗
Dense	✓	✓ ← must be dense