

CS411

*Introduction to Database
Programming*

Mark Overholt

Announcements

- Homework #1 Due Tonight! 11:59pm on Compass!
- **Project Track 1:**if you still need a project group or that your group still needs more people, please let Hanna know **ASAP**. We'll send out your group assignment by Friday. Group assignment is first come first serve.
So, email Hanna the following information:
 - If you still need a project group; or
 - If your group needs more people, indicate the number of people needed.
- There are 6 Feedback questions regarding this status of this class. PLEASE respond to these survey questions, so that we may improve the class.
 - <https://agora.cs.illinois.edu/display/cs411sp11/Feedback>

Today's Objectives

- Quick overview of the Projects server and cPanel
 - Logging in
 - Transferring Files to your Server
 - Creating Databases and DB Users
 - PHPMyAdmin
- DB Programming
 - Using SQL in PHP
 - Using SQL in Java
 - Using Stored Procedures

cPanel

- Everyone should have access to the projects.cs.illinois.edu server for their project.
 - If you don't...let us know!!
- Demo of cPanel
 - Logging on - <http://accounts.cs.illinois.edu/projects/>
 - Adding/Modifying Files
 - Adding DB
 - Adding User
 - PHPMyAdmin

DB Programming

- SQL is a very high-level language.
- **Not** intended for general-purpose computations.
- Solutions:
 - Outside DBMS: use SQL together with general-purpose programming languages.
 - Inside DBMS: augment SQL with constructs from general-purpose programming languages.

All these methods follow the same basic paradigm

1. Connect to a DB server.
2. Say what database you want to use.
3. Assemble a string containing an SQL statement.
4. Get the DBMS to prepare a plan for executing the statement.
5. Execute the statement.
6. Extract the results into variables in the local programming language.

Database specific API

- Designed and implemented by the DBMS vendor for a specific programming language.
- We will go over MySQL API for PHP.

Connection

- `mysql_connect` opens a connection to the DBMS.
- It gets the DBMS and login information and returns a *connection* resource. The *connection* resource is used in future calls.
- `mysql_select_db` selects the desired database
- `mysql_close` closes the connection at the end. It is called automatically at the end of the script.

Connection Example in PHP

```
$username = "USERNAME";  
$password = "PASSWORD";  
$host = "localhost";  
$dbname = "DatabaseName";  
$connect = mysql_connect($host, $username,  
    $password);  
mysql_select_db($dbname, $connect);  
    //Do Stuff with the DB Connection  
mysql_close($connect);
```

Executing Statements

- `mysql_query(S, C)` causes the SQL statement represented by `S` to be executed on connection `C`.
- We can detect DBMS initiated errors using `mysql_error()`.
- `mysql_query` returns a handle to the query result set.
 - It returns `TRUE/FALSE` for `DELETE`, `UPDATE`, and `INSERT` statements.

Query Example

```
$query = "SELECT * FROM Beer";
```

```
$result = mysql_query($query, $connect);
```

```
$query = "INSERT INTO Beer (manf, beer)  
VALUES ("Miller", "High Life");
```

```
if (mysql_query($query, $connect)) {
```

```
    echo "Insert Successful";
```

```
} else {
```

```
    echo "Insert Failed: " . mysql_error();
```

```
}
```

Fetching Tuples

- When the SQL statement executed is a query, we need to fetch the tuples of the result.
- `mysql_fetch_array(H)` gets the tuples from the result set *H* and stores them in an associative array.
- `mysql_free_result(H)` frees the result set. It is called automatically at the end of the script

Fetching Tuples Example

```
$query = "SELECT * FROM Beer";
```

```
$result = mysql_query($query, $connect);
```

```
While ($row = mysql_fetch_array($result)) {
```

```
    $manf = $row['manf'];
```

```
    $beer = $row['beer'];
```

```
    echo "The beer " . $beer . " is made by " . $manf;
```

```
}
```

```
mysql_free_result($result);
```

Helpful PHP Mysql Commands

- `mysql_num_rows(r);`
 - Checks the *Result* handler to see how many rows were returned
- `mysql_affected_rows();`
 - Useful on DELETE or UPDATE commands to see how many rows were changed
- `mysql_insert_id();`
 - Useful for AUTO_INCREMENT fields. Get the last ID (PRIMARY KEY) you just inserted.
- More functions at:
 - <http://us.php.net/mysql>

Open interface: JDBC

- Database specific API makes the program dependent to one DBMS.
- Open interfaces solve this problem:
 - Designed by a third party like the creator of the language.
 - Implemented by DBMS vendors.
 - Used by DB programmers.
- Programmer do not have to change their code to work with another DBMS.

Connections and Statements

- The same progression from connections to statements that we saw in CLI appears in JDBC.
- A *connection object* is obtained from the environment in a somewhat implementation-dependent way.
- We'll start by assuming we have myCon, a connection object.

Statements

- JDBC provides two classes:
 1. Statement = an object that can accept a string that is an SQL statement and can execute such a string.
 2. PreparedStatement = an object that has an associated SQL statement ready to execute.
- Why? Performance
 - PreparedStatement SQL command is stored in DBMS after the first call.

Creating Statements

- The Connection class has methods to create Statements and PreparedStatement.

```
Statement stat1 = myCon.createStatement();
```

```
PreparedStatement stat2 =
```

```
myCon.prepareStatement(
```

```
    "SELECT Address FROM Supplier" +
```

```
    "WHERE Supplier_Name = 'John Smith'"
```

```
);
```

Java trick: +
concatenates
strings.



Executing SQL Statements

- Programmer handles errors using SQLException class.
- JDBC distinguishes queries from modifications, which it calls “updates.”
- Statement and PreparedStatement each have methods executeQuery and executeUpdate.
 - For Statements, these methods have one argument: the query or modification to be executed.
 - For PreparedStatements: no argument.

Example: Update

- stat1 is a Statement.
- We can use it to insert a tuple as:

```
stat1.executeUpdate(  
    “INSERT INTO Supplier” +  
    “VALUES(‘S4’, ‘Mary’, ‘12 Goodwin St.’)”  
);
```

Example: Query

- stat2 is a PreparedStatement holding the query “SELECT Address FROM Supplier WHERE Supplier_Name = ‘John Smith’”.
- executeQuery returns an object of class ResultSet --- we’ll examine it later.
- The query:

```
ResultSet Menu = stat2.executeQuery();
```

Accessing the ResultSet

- An object of type ResultSet is something like a cursor.
- Method Next() advances the “cursor” to the next tuple.
 - The first time Next() is applied, it gets the first tuple.
 - If there are no more tuples, Next() returns the value FALSE.

Accessing Components of Tuples

- When a ResultSet is referring to a tuple, we can get the components of that tuple by applying certain methods to the ResultSet.
- Method `getX (i)`, where X is some type, and i is the component number, returns the value of that component.
 - The value must have type X .

Example: Accessing Components

- List is the ResultSet for the query “SELECT Address FROM Supplier WHERE Supplier_Name = ‘John Smith’”.

- Access the address from each tuple by:

```
while ( List.Next() ) {  
    theAddress = List.getString(1);  
    /* do something */  
}
```

- More: <http://java.sun.com/products/jdbc/overview.html>

Programming Inside the DBMS

Stored Procedures

- An extension to SQL, called SQL/PSM, or “persistent, stored modules,” allows us to store procedures as database schema elements.
- The programming style is a mixture of conventional statements (if, while, etc.) and SQL.
- Lets us do things we cannot do in SQL alone.
- Why? Performance.
- They are harder to develop and maintain.

Basic PSM Format

```
CREATE PROCEDURE <name> (  
    <parameter list> )
```

```
    <optional local declarations>
```

```
    <body>;
```

- Function alternative:

```
CREATE FUNCTION <name> (
```

```
    <parameter list> ) RETURNS <type>
```

Example: compute square footage of house lot; then you can query on it

Parameters in PSM

- Unlike the usual name-type pairs in languages like C, PSM uses mode-name-type triples, where the *mode* can be:
 - IN = procedure uses value, does not change value.
 - OUT = procedure changes, does not use.
 - INOUT = both.

Example: Stored Procedure

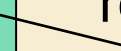
- Let's write a procedure that takes three arguments i , n , and a , and adds a tuple to Supplier that has Supplier_ID = i , Supplier_Name = n , and Address = a .
 - Used to add new Supplier more easily.

The Procedure

```
CREATE PROCEDURE newSup (
```

```
IN i    VARCHAR(20),  
IN n    VARCHAR(50),  
IN a    VARCHAR(100)
```


Parameters are all
read-only, not changed



```
)
```

```
INSERT INTO Supplier  
VALUES(i, n, a);
```

The body ---
a single insertion



Invoking Procedures

- Use SQL/PSM statement CALL, with the name of the desired procedure and arguments.

- Example:

CALL newSup('s5', 'Bob', '1 Main St.);

- Functions used in SQL expressions where a value of their return type is appropriate.

Types of PSM statements -- 1

- RETURN <expression> sets the return value of a function.
 - Unlike C, etc., RETURN *does not* terminate function execution.
- DECLARE <name> <type> used to declare local variables.
- BEGIN . . . END for groups of statements.
 - Separate by semicolons.

Types of PSM Statements -- 2

- Assignment statements:

SET <variable> = <expression>;

– Example: SET b = 'Bud';

- Statement labels: give a statement a label by prefixing a name and a colon.

Example: IF

- Let's rate suppliers by how many parts they have
 - ≤ 0 parts: 'bad'.
 - < 2 parts: 'average'.
 - ≥ 2 parts: 'good'.
- Function Rate(s) rates supplier s.

Example: IF (continued)

```
CREATE FUNCTION Rate (IN b VARCHAR(50) )
```

```
    RETURNS CHAR(10)
```

```
    DECLARE pCnt INTEGER;
```

```
    BEGIN
```

```
        SET pCnt = (SELECT COUNT(*) FROM Supplier,  
                    Catalog WHERE Supplier.Supplier_ID =  
                    Catalog.Supplier_ID and  
                    Supplier_Name= b);
```

Number of
the parts of the
supplier b

```
        IF pCnt < 0 THEN RETURN 'bad'  
        ELSEIF pCnt < 2 THEN RETURN 'average'  
        ELSE RETURN 'good'  
        END IF;
```

```
    END;
```

Return occurs here, not at
one of the RETURN statements

Nested
IF statement

Example: Exiting a Loop

loop1: LOOP

...

LEAVE loop1; ← If this statement is executed . . .

...

END LOOP; ← Control winds up here

Other Loop Forms

- WHILE <condition>
 DO <statements>
END WHILE;
- REPEAT <statements>
UNTIL <condition>
END REPEAT;

Queries

- General SELECT-FROM-WHERE queries are *not* permitted in PSM.
- There are three ways to get the effect of a query:
 1. Queries producing one value can be the expression in an assignment.
 2. Single-row SELECT . . . INTO.
 3. Cursors.

Example: Assignment/Query

- If a is a local variable and Supplier(Supplier_ID, Supplier_Name, Address) the usual relation, we can get the address of 'John Smith' by:

```
SET a = (SELECT Address FROM Supplier  
        WHERE Supplier_Name = 'John  
        Smith' );
```

SELECT . . . INTO

- An equivalent way to get the value of a query that is guaranteed to return a single tuple is by placing INTO <variable> after the SELECT clause.
- Example:

```
SELECT Address INTO a FROM  
Supplier WHERE Supplier_Name =  
'John Smith';
```


Cursors

- A *cursor* is essentially a tuple-variable that ranges over all tuples in the result of some query.
- Declare a cursor *c* by:

DECLARE *c* CURSOR FOR <query>;

Opening and Closing Cursors

- To use cursor c , we must issue the command:
 OPEN c ;
 - The query of c is evaluated, and c is set to point to the first tuple of the result.
- When finished with c , issue command:
 CLOSE c ;

Fetching Tuples From a Cursor

- To get the next tuple from cursor c , issue command:

FETCH FROM c INTO x_1, x_2, \dots, x_n ;

- The x 's are a list of variables, one for each component of the tuples referred to by c .
- c is moved automatically to the next tuple.

Breaking Cursor Loops -- 1

- The usual way to use a cursor is to create a loop with a FETCH statement, and do something with each tuple fetched.
- A tricky point is how we get out of the loop when the cursor has no more tuples to deliver.

Breaking Cursor Loops -- 2

- Each SQL operation returns a *status*, which is a 5-digit number.
 - For example, 00000 = “Everything OK,” and 02000 = “Failed to find a tuple.”
- In PSM, we can get the value of the status in a variable called SQLSTATE.

Breaking Cursor Loops -- 3

- We may declare a condition, which is a boolean variable that is true if and only if SQLSTATE has a particular value.
- Example: We can declare condition NotFound to represent 02000 by:

```
DECLARE NotFound CONDITION FOR  
SQLSTATE '02000';
```

Breaking Cursor Loops -- 4

- The structure of a cursor loop is thus:

```
cursorLoop: LOOP
```

```
...
```

```
FETCH c INTO ... ;
```

```
IF NotFound THEN LEAVE cursorLoop;
```

```
END IF;
```

```
...
```

```
END LOOP;
```

Example: Cursor

- Let's write a procedure that examines Supplier and Catalog and raises by \$1 the cost of all parts in the 'John Smith' inventory that are under \$10.
 - Yes, we could write this as an UPDATE, but the details are instructive anyway.

The Needed Declarations

```
CREATE PROCEDURE raiseCost( )
```

```
DECLARE theID VARCHAR(50);  
DECLARE theCost REAL;
```

Used to hold
ID-cost pairs
when fetching
through cursor c

```
DECLARE NotFound CONDITION FOR  
SQLSTATE '02000';
```

```
DECLARE c CURSOR FOR
```

Returns John's parts costs

```
(SELECT Supplier.Supplier_ID, Cost FROM  
Supplier, Catalog WHERE Supplier.Supplier_ID =  
Catalog.Supplier_ID and Supplier_Name = 'John  
Smith');
```

The Procedure Body


```
BEGIN
```

```
  OPEN c;
```

```
  menuLoop: LOOP
```

```
    FETCH c INTO theID, theCost;
```

Check if the recent
FETCH failed to
get a tuple



```
    IF NotFound THEN LEAVE menuLoop END IF;
```

```
    IF theCost < 10.00 THEN
```

```
      UPDATE Catalog SET Cost = theCost+1.00
```

```
      WHERE Supplier_ID = theID;
```


```
    END IF;
```

```
  END LOOP;
```

```
  CLOSE c;
```

```
END;
```

If John charges less than \$10 for
the part raise it's cost by \$1.



Further Readings

- <http://dev.mysql.com/doc/refman/5.0/en/stored-routines.html>
- http://www.oracle.com/technology/tech/pl_sql/index.html