UNIVERSITY OF ILLINOIS
AT URBANA-CHAMPAIGN

CS411 - Welcome!

illinois.edu

About me

- Ryan “Not Dr.” Cunningham
  - A PhD student in bioinformatics
    - Dabbled in NLP, information retrieval, security, machine learning
  - Worked for DoD, DoE, Agribusiness, and Telecoms
  - MS in Computer Science from Central Florida
  - BS in Computer Engineering from Cincinnati
  - Taught in 125, 225, 232, 233, and 373

About me

- Ryan “Not Dr.” Cunningham
  - A PhD student in bioinformatics
    - Dabbled in NLP, information retrieval, security, machine learning
  - Worked for DoD, DoE, Agribusiness, and Telecoms
  - MS in Computer Science from Central Florida
  - BS in Computer Engineering from Cincinnati
  - Taught in 125, 225, 232, 233, and 373

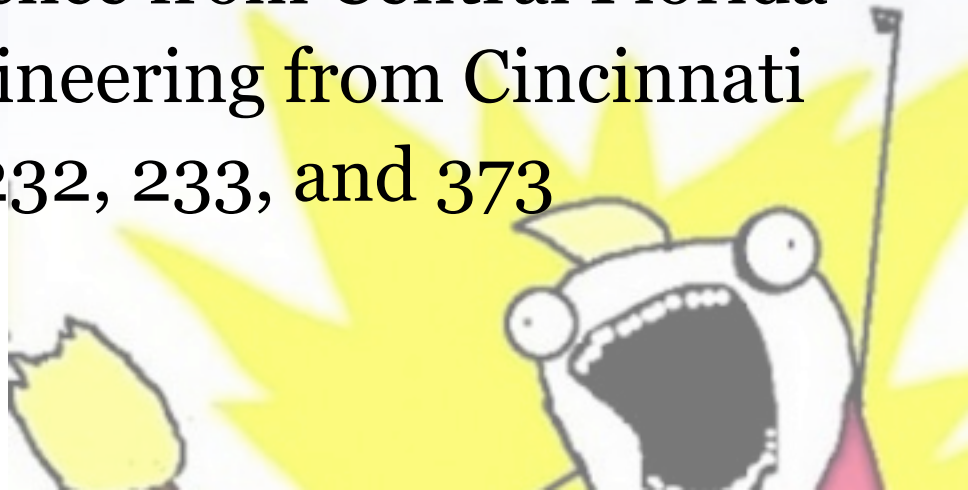

About staff

- Nikita Spirin
  - Course projects and online students
- Khuram Shahzad
  - Course project track 1 and Piazza
- Magnesh Bendre
  - MP design
- Rui Wang
  - Written assignments

Course Website

- https://wiki.engr.illinois.edu/display/cs411sp13/Overview
- Syllabus, assignments, etc.
- All official course policies posted here

Grading Policy

- Homework 25%
- Projects 30%
- Midterm 20%
- Final Exam 25%

Course Projects

- Track 1: Database Web Application
  - Teams of 3-4 (form by Feb 6th)
  - Semester long project with several stages
  - Opportunity to be creative and ambitious
  - Start brainstorming!

Course Projects

- Track 2: Literature Survey or Research Extension
  - Required for those registered for 4 credits
  - Optional extra credit for others
  - Groups of 1-2
  - Either do a high quality literature review or expand your semester project into a serious research project

Homework

- 4-5 written assignments
  - Meant to reinforce concepts and prepare you for the midterm and final
- 3-4 programming assignments
  - Meant to help you understand the complexities of implementing a DBMS

Piazza

- https://piazza.com/class#spring2013/cs411
- A web forum where you can post questions
- Sign up ASAP so you don't miss out!
  - If you prefer your UIUC email, just sign up
  - If you prefer to use another account, please send me an email

Texbook

- *Database Systems: The Complete Book, Second Edition* by Hector Garcia-Molina, Jeffrey D. Ullman, and Jennifer Widom

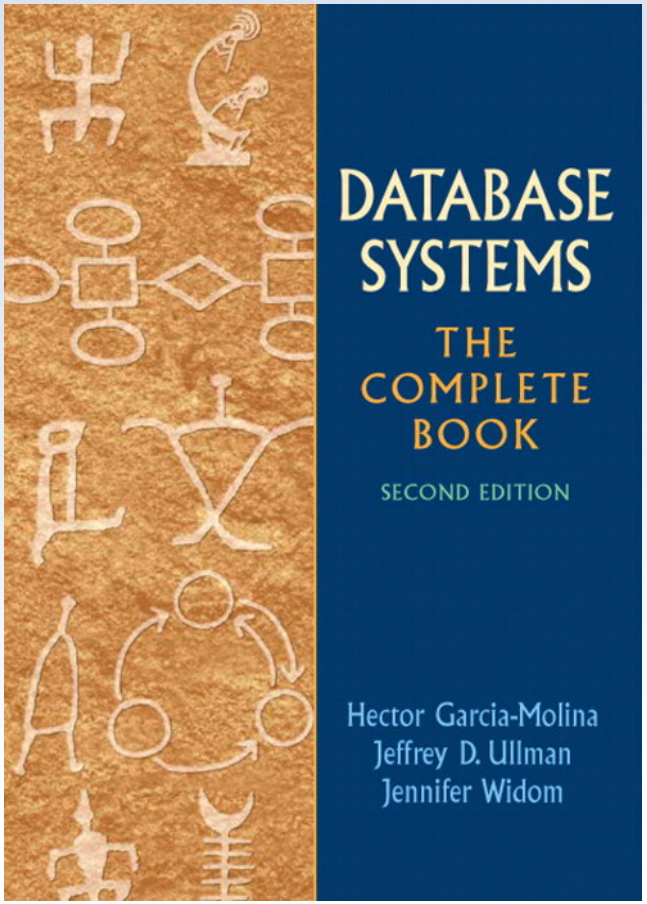

Why study databases?

- Most computer science assumes we can manipulate data in RAM
- What to do data is much larger than RAM?
- This is very common:
  - credit card transactions, mobile phones, search engines
  - Google operates on *petabytes* of data

Why study databases?

- Without them, our current way of life would be impossible.
  - No Google, iPhone, Facebook, or Amazon!
- Database systems are crucial for our infrastructure and economy

Why study databases?

- Concepts are extremely useful in other domains

What makes databases different?

- Can't restructure the data for each computation (only one *schema*)
- Efficiently use the entire system (CPU, RAM, Disk, *and* Network)
- Data should be *persistent* and continuously updated
- Multiple concurrent users

What is a database?

- A *database management system* (DBMS)
  1. Allows users to specify *schema (*logical structure) of their data with a *data-definition language* (DDL)
  2. Allows users to *query* the data (perform computation on the data) with a *data-manipulation language* (DML)

What is a database?

- A *database management system* (DBMS)
  3. Supports *persistent storage* of large amounts of data in a way that supports 1 and 2 above
  4. Enables *durability* in the face of failures
  5. Controls accesses by multiple users, ensuring
     - *isolation* (user's access is independent of others)
     - *atomicity* (an action is never performed partially)

History

- Problems first encountered in the 1960's
  - Banking systems
  - Airline reservations (surprisingly important)
  - Corporate records
- Essentially, people were building *ad hoc* systems on top of file systems
- Each query required a *custom program!*

History

- In 1970's Ted Codd wrote "A relational model for large shared data banks"
  - Proposed a *relational model* of data
  - Data storage abstracted from user
  - Supported high level query language
- Through the 1980's and 1990's, this model became standard and widely adopted

History

- From 2000's to present:
  - Codd's model is still the core paradigm of the DBMS infrastructure
  - But much more data that is less organized
    - Images, video, social networking
  - Peer-to-peer and parallel systems developed
  - Extended and supplemented relational model in light of these developments

An overview

- How does a DBMS work?
  - Here's an overview

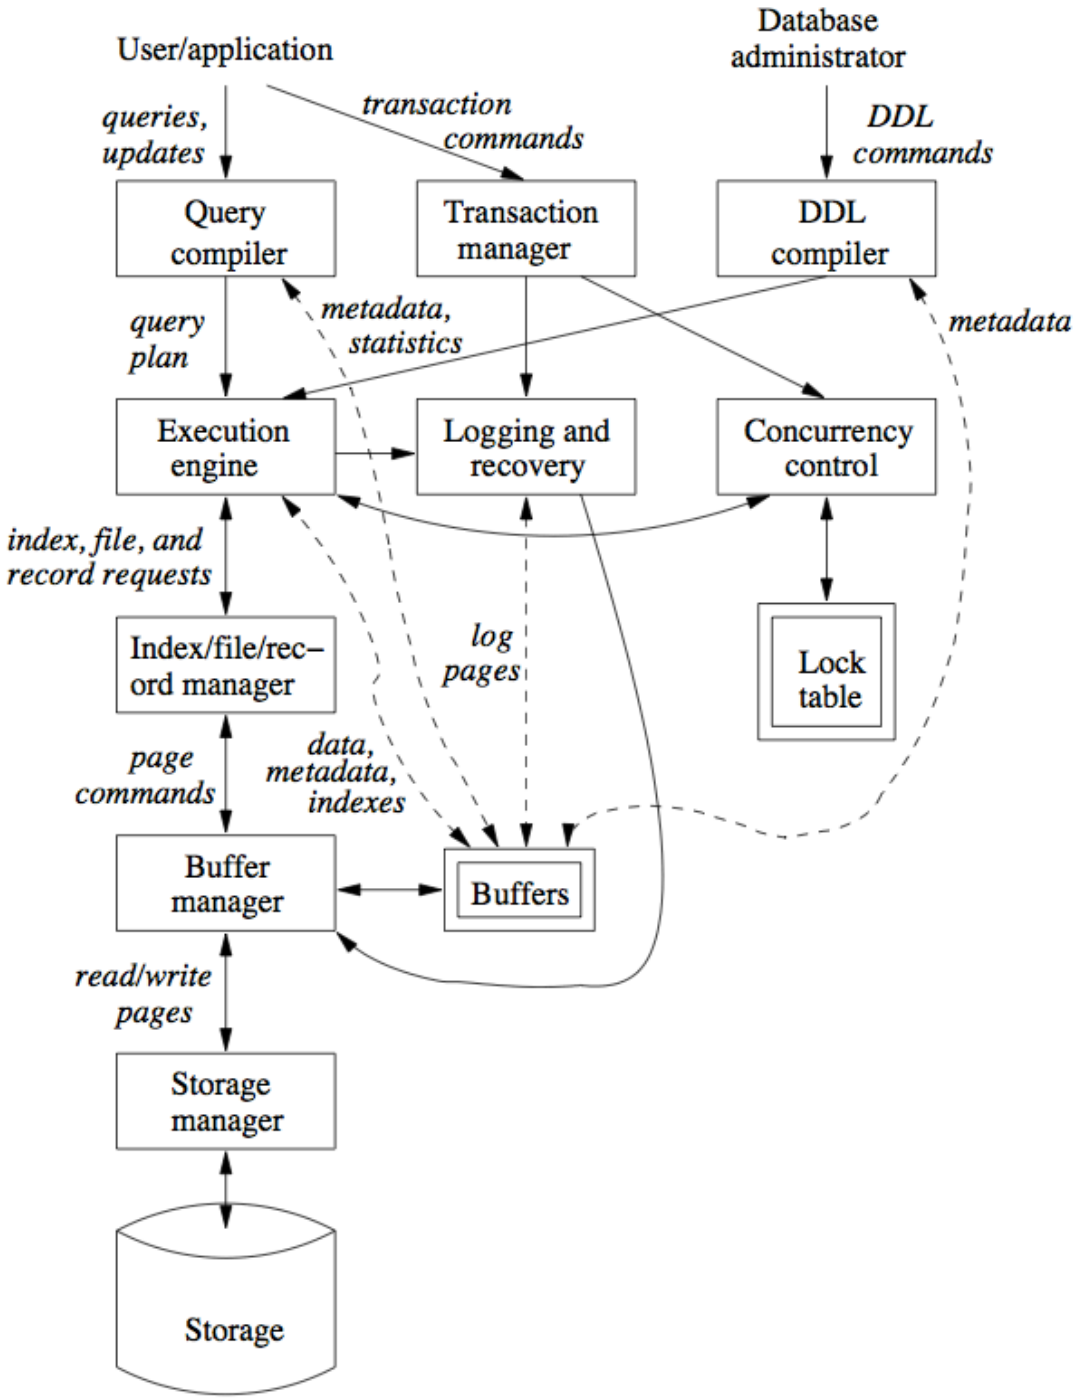

illinois.edu

Did you get all that?

- We'll spend all semester learning about these systems
- But let's break this down to get a little preview

Interacting with the DBMS

- Two ways to interact with the DBMS
  1. As a "user" interacting with the data
  2. As an "administrator" modifying the structure of the data

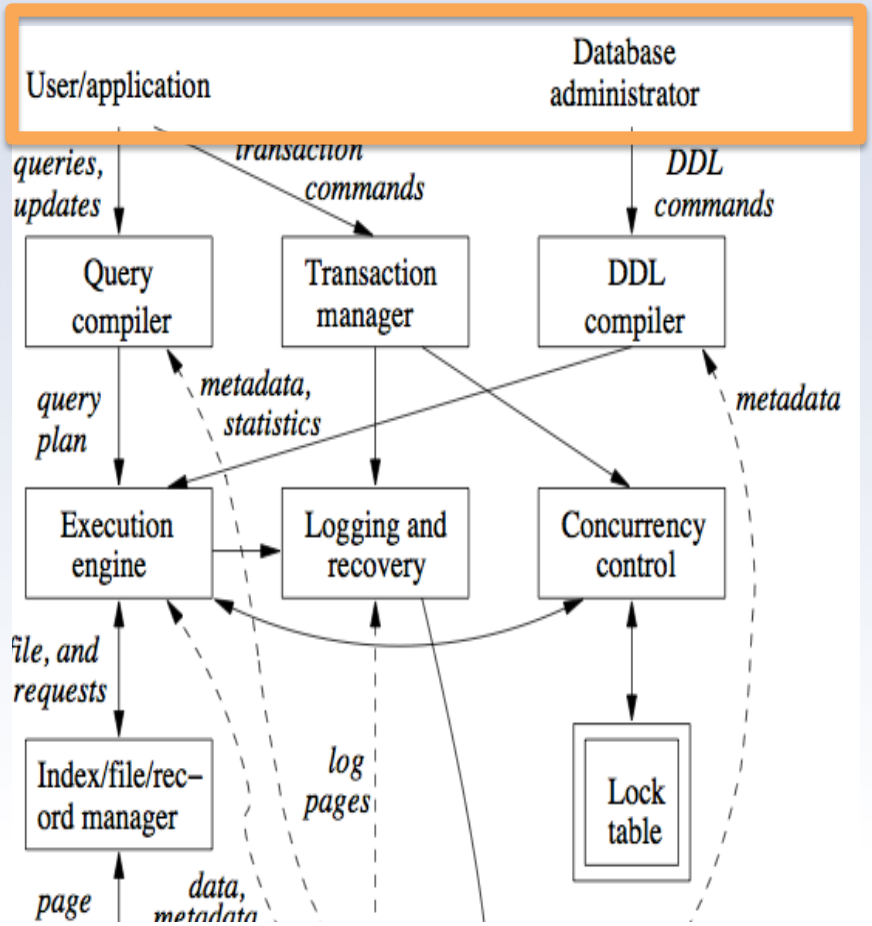

Focusing on the user

- Users submit queries to the query compiler in a data manipulation language (DML)
- Parsed by the query compiler into a query plan

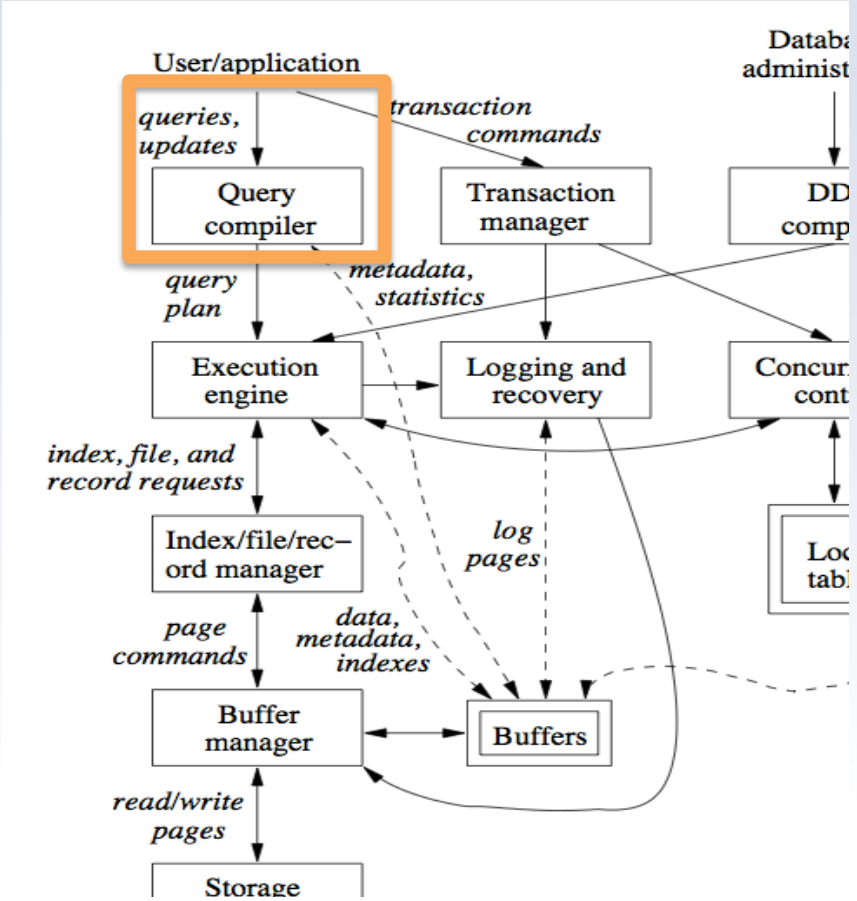

Focusing on the user

- Query plan is executed by the execution engine
- Sends specific low level requests to the index/record manager to get the data

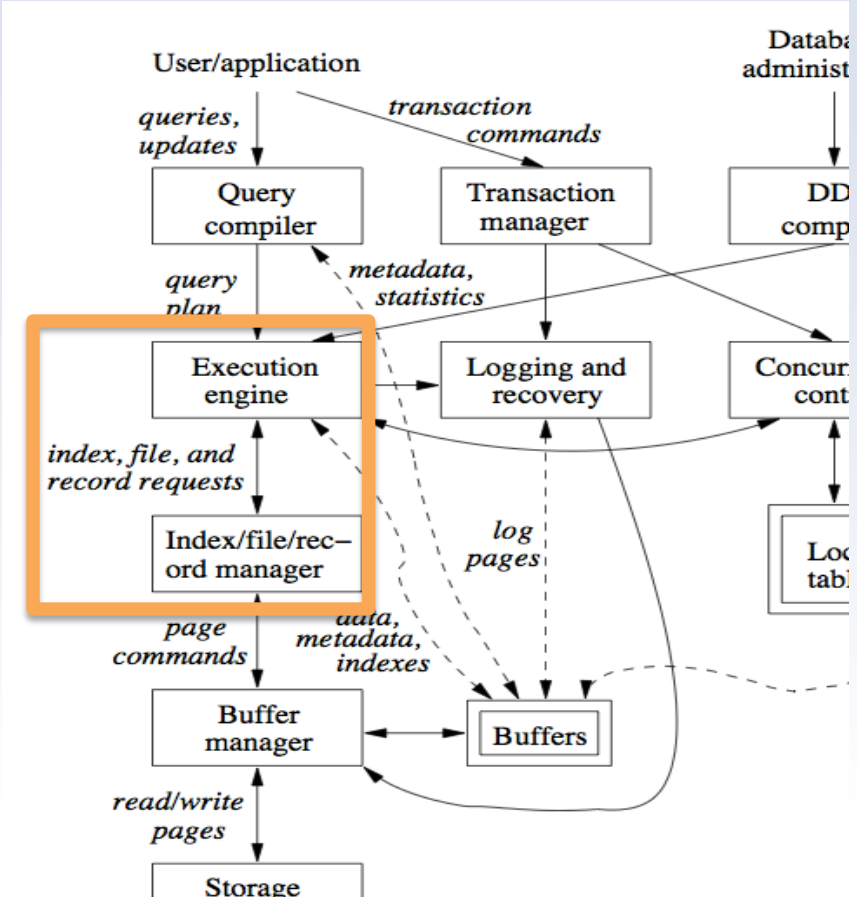

Focusing on the admin

- The database administrator (DBA) sends data definition language DDL commands to the DDL compiler
- Also sent to the execution engine

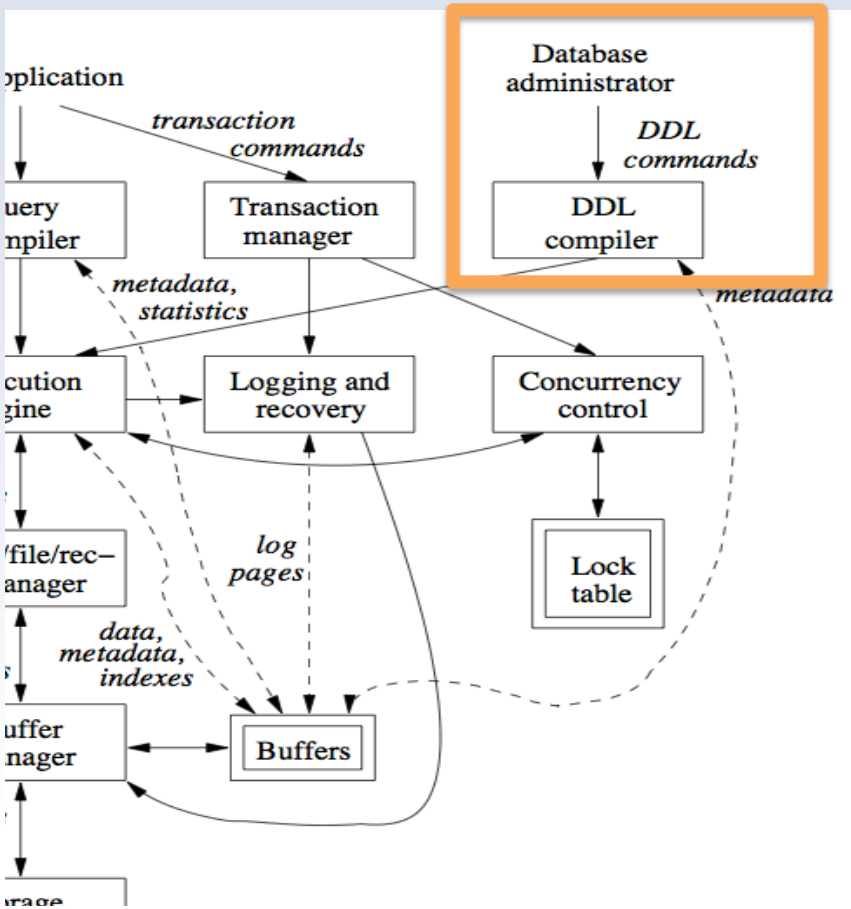

Storage and Buffer Management

- Storage manager keeps track of where the data is
  - Stored in large chunks so we can access it in bulk
  - Transferred in and out of RAM in *pages*

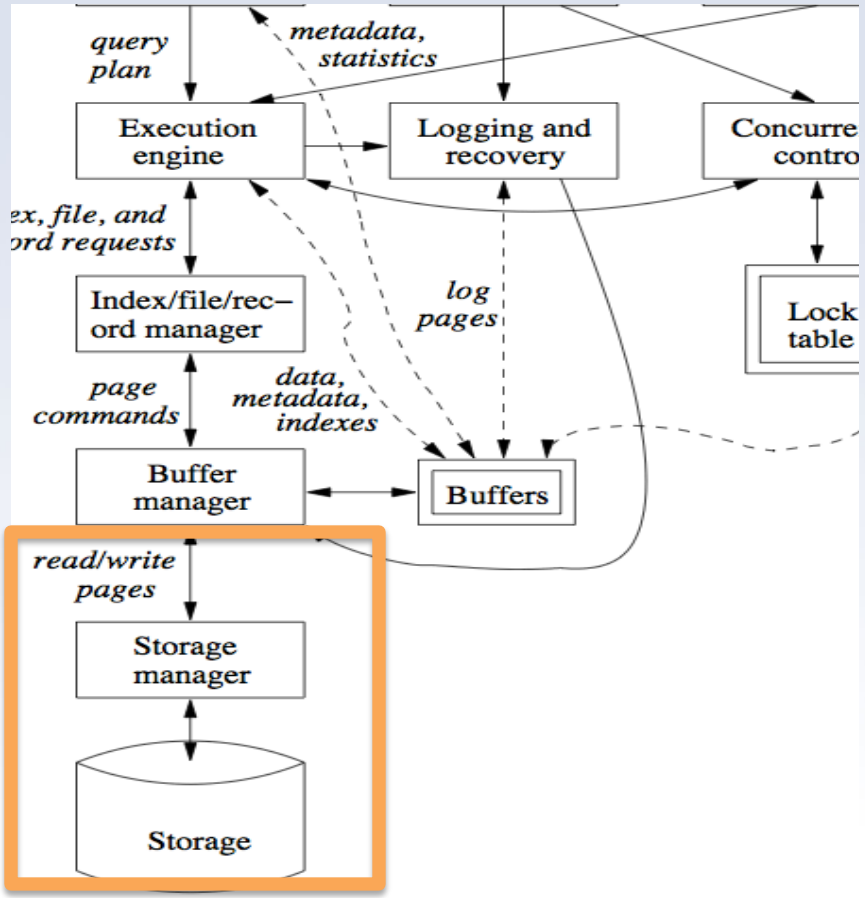

Storage and Buffer Management

- Buffer manager partitions RAM into *buffers*
  - essentially keeps data in page sized chunks that we can perform computation on

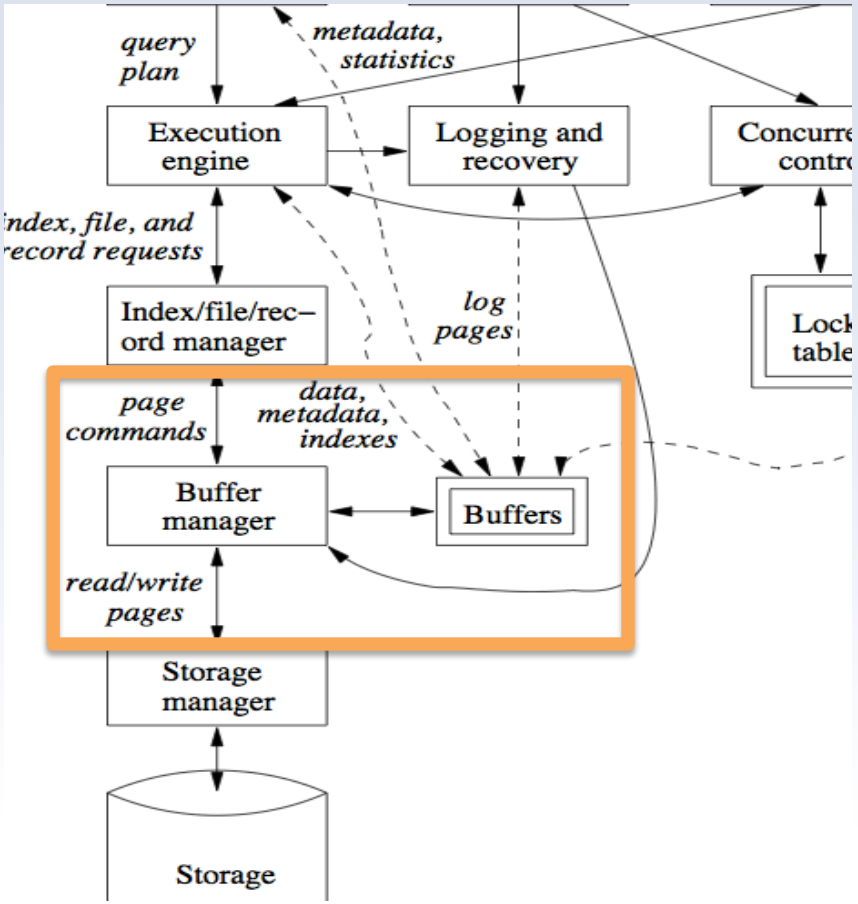

Transaction Management

- ACID Test
  - Atomicity - “all or nothing”
  - Isolation - “don’t interfere”
  - Consistency - “maintain constraints”
  - Durability - “don’t lose anything”

Transaction Management

- Transaction manager receives units of work called *transaction commands*
- It makes sure ACID test is satisfied for all transactions

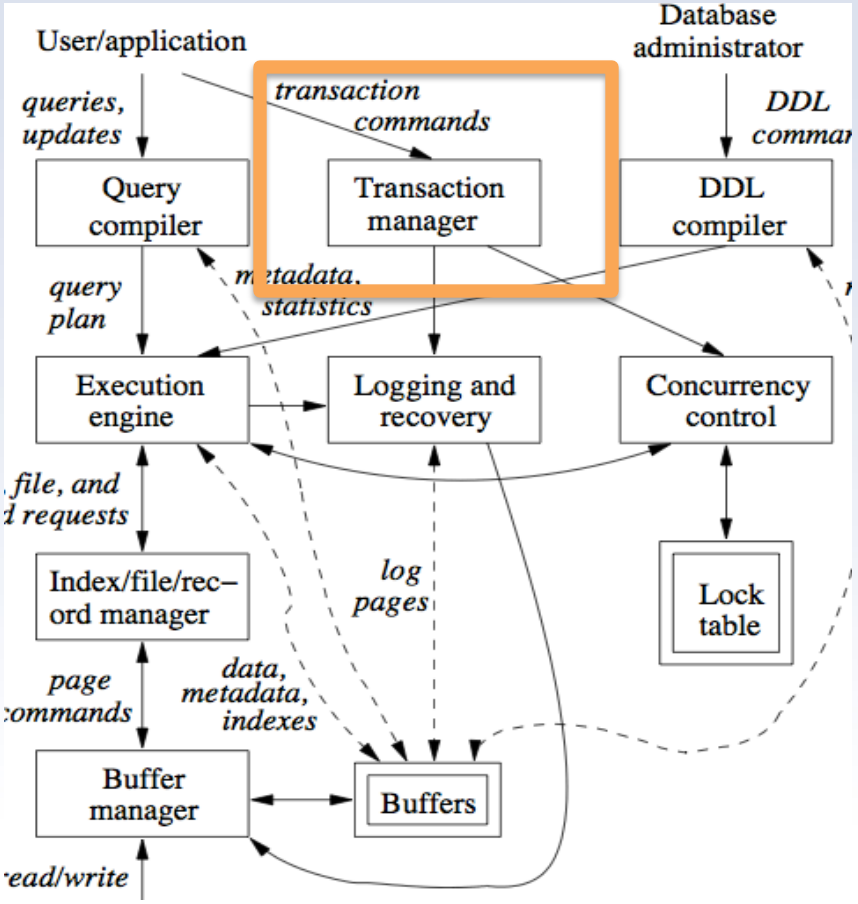

Transaction Management

- Logs execution of transactions so transactions that fail can be recovered

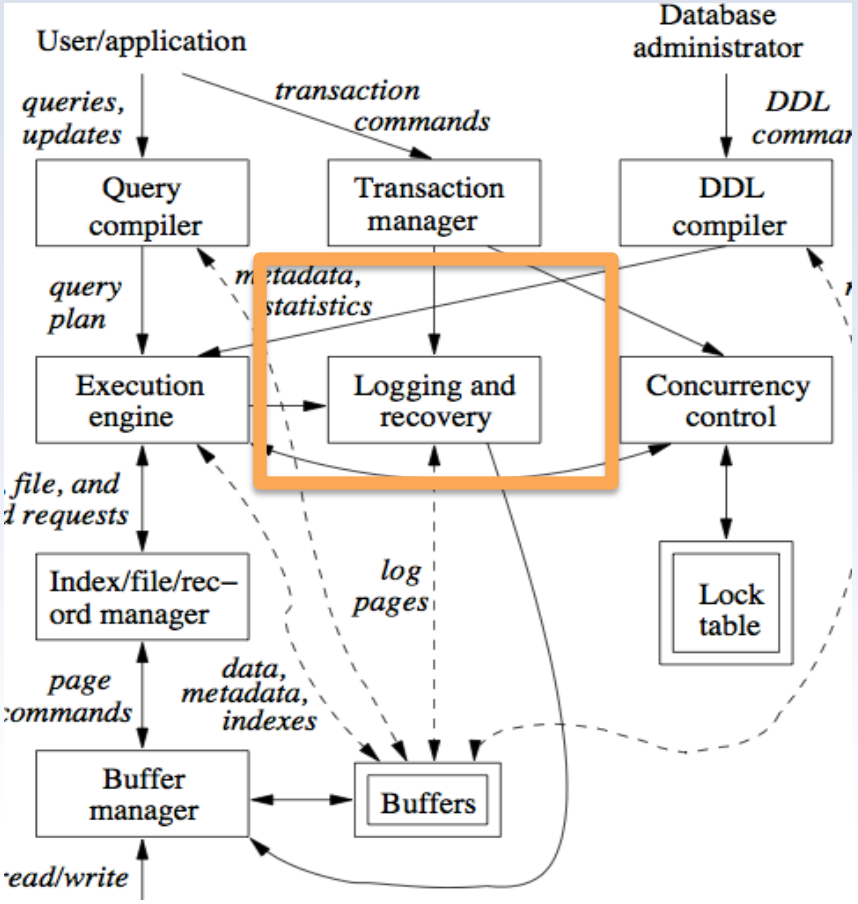

Transaction Management

- Tracks concurrently executing transaction commands
- Locks parts of the database as needed to ensure transactions don't interfere with each other

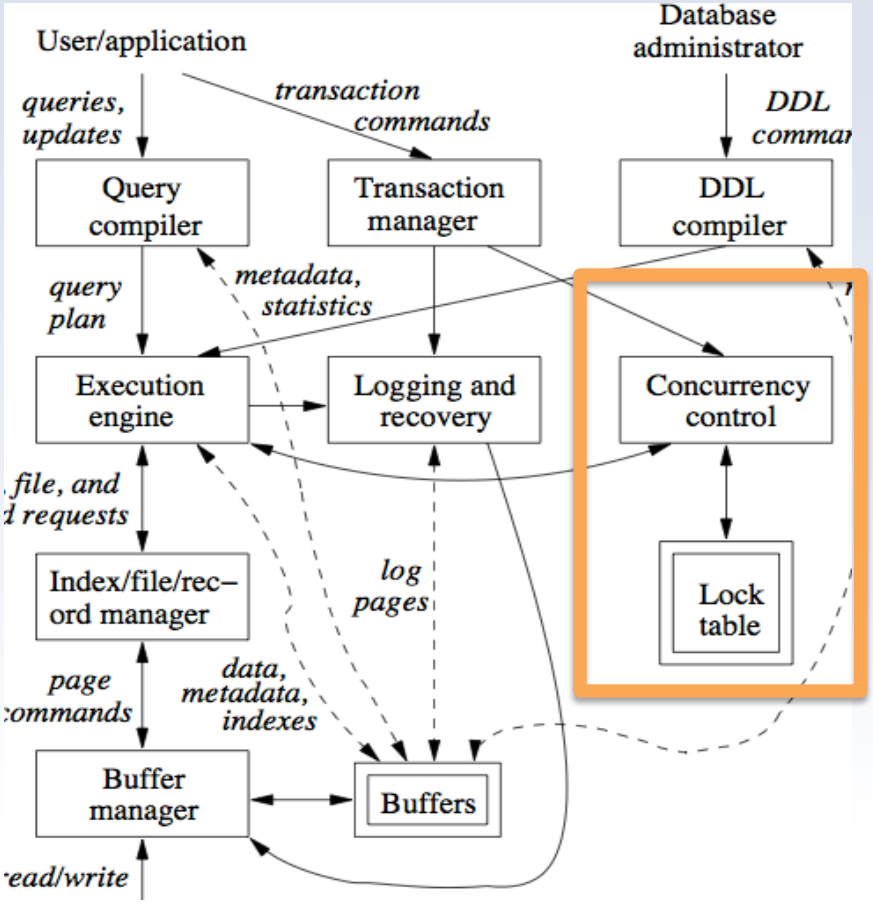

Transaction Management

- If multiple conflicting requests are waiting for the same data, must perform *deadlock resolution*

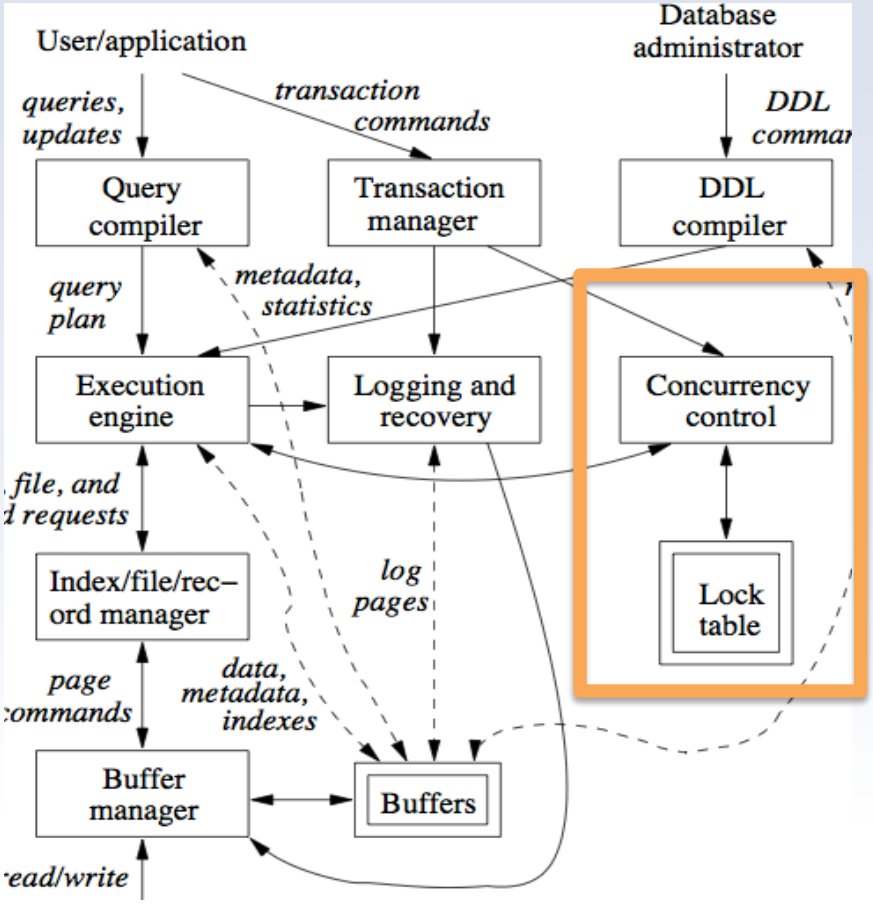

Course Overview

1. Relational Model: Query/DML
   - Theoretical and practical perspective
2. Relational Model: Design/DDL
   - Theoretical and practical perspective
   - Advanced Manipulation concepts
3. DBMS Implementation

Course Overview

4. Advanced Topics
   - Parallel/Distributed Databases
   - Information Integration
   - Data Mining/Information Retrieval

Next Lecture

- We'll start learning about Codd's relational model