

UNIVERSITY OF ILLINOIS
AT URBANA-CHAMPAIGN

CS411 - Query Parsing/Optimization



illinois.edu

Announcements

- HW3 is due Friday
- MP3 is coming
 - Related to query optimization
 - Won't involve programming
- Piazza poll for advanced topics



Review

- What is the “buffer pool”?
- What is a “table scan”?
- What is the main idea of the two pass sort based algorithms?



Review

- What is the main idea of the two pass hash based algorithms?
- How was “hybrid hash join” different?



One-Pass/NBLJ Summary

| Operator | M required | I/O Cost |
|----------------------------------|--------------------|---------------|
| σ, π | 1 | B |
| δ, γ | B | B |
| $\cup, \cap, -, \bowtie, \times$ | $\min(B(R), B(S))$ | $B(R) + B(S)$ |
| \bowtie | $M \geq 2$ | $B(R)B(S)/M$ |



Sorting-Based Summary

| Operator | M required | I/O Cost |
|------------------|---------------------------|----------------|
| δ, γ | \sqrt{B} | $3B$ |
| $\cup, \cap, -$ | $\sqrt{B(R)+B(S)}$ | $3(B(R)+B(S))$ |
| \bowtie | $\sqrt{\max(B(R), B(S))}$ | $5(B(R)+B(S))$ |
| \bowtie | $\sqrt{B(R)+B(S)}$ | $3(B(R)+B(S))$ |



2 Pass Hash-Based Summary


| Operator | M required | I/O Cost |
|------------------|---------------|---------------------------------|
| δ, γ | \sqrt{B} | $3B$ |
| $\cup, \cap, -$ | $\sqrt{B(S)}$ | $3(B(R)+B(S))$ |
| \bowtie | $\sqrt{B(S)}$ | $3(B(R)+B(S))$ |
| \bowtie | $\sqrt{B(S)}$ | $(3-2M/B(S))x$ $(B(R)+B(S))$ |

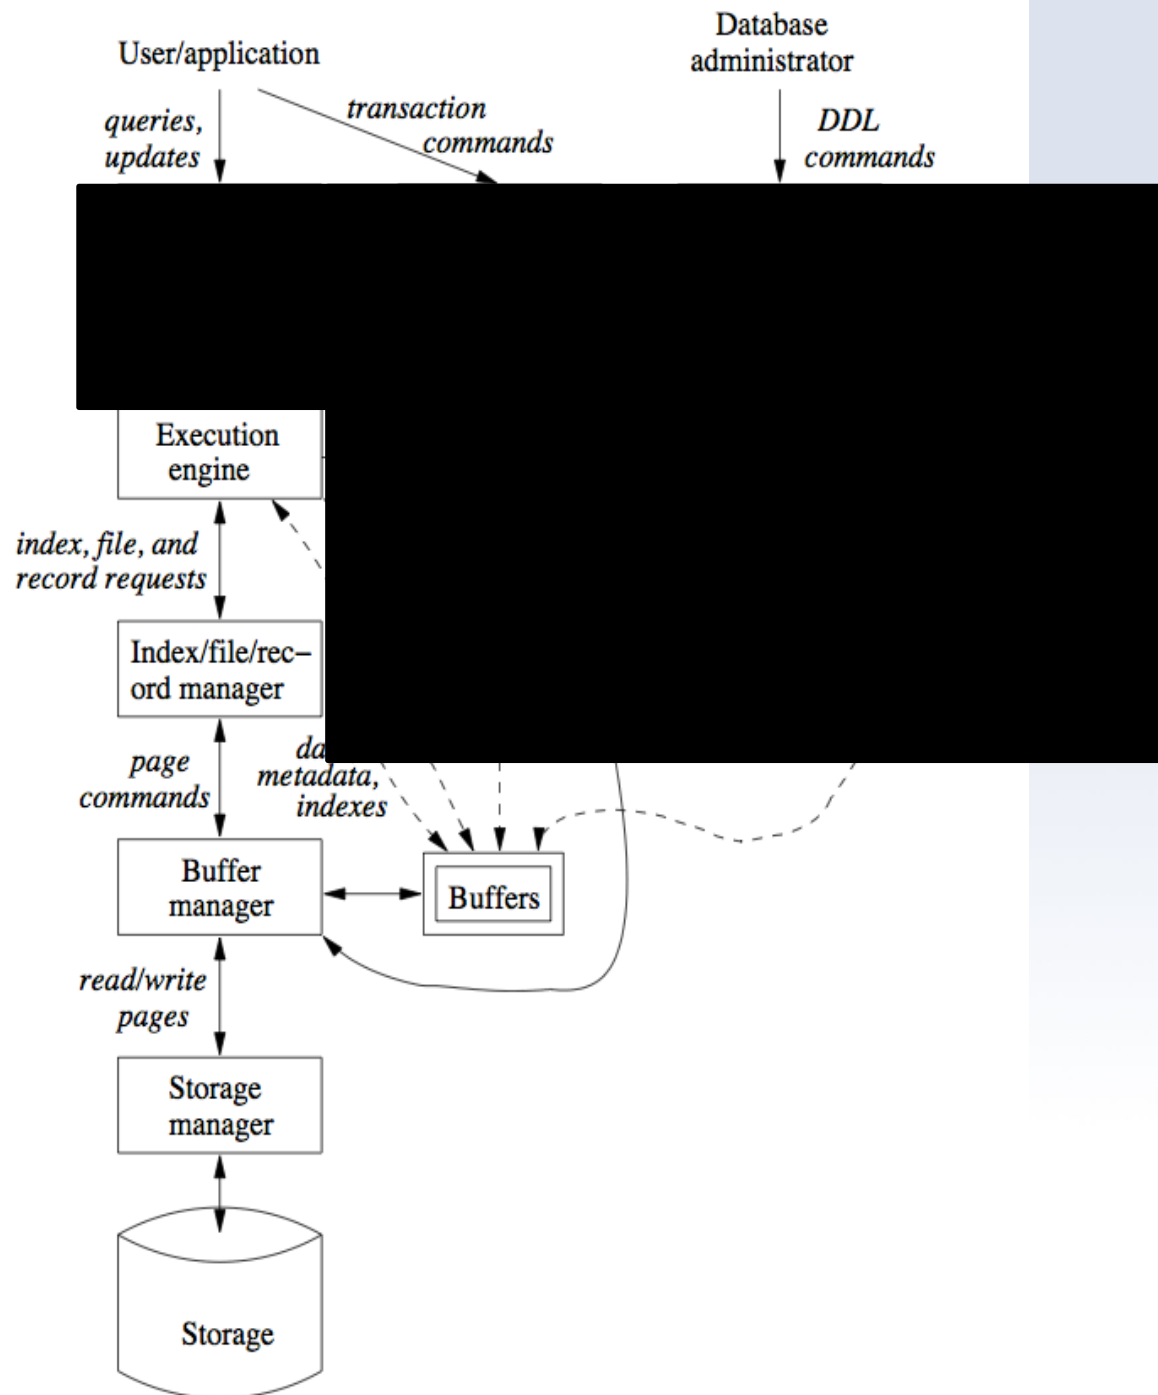


Practical Example

- On a system with $8\text{GB}=2^{33}$ of RAM
 - we can sort/join/deduplicate 2^{66}
 - bigger than an exabyte of data!
- On a system with $256\text{GB}=2^{38}$
 - we can sort/join/deduplicate 2^{76} bytes of data
 - almost a yottabyte!







User/application

*queries,
updates*

*transaction
commands*

Database
administrator

*DDL
commands*

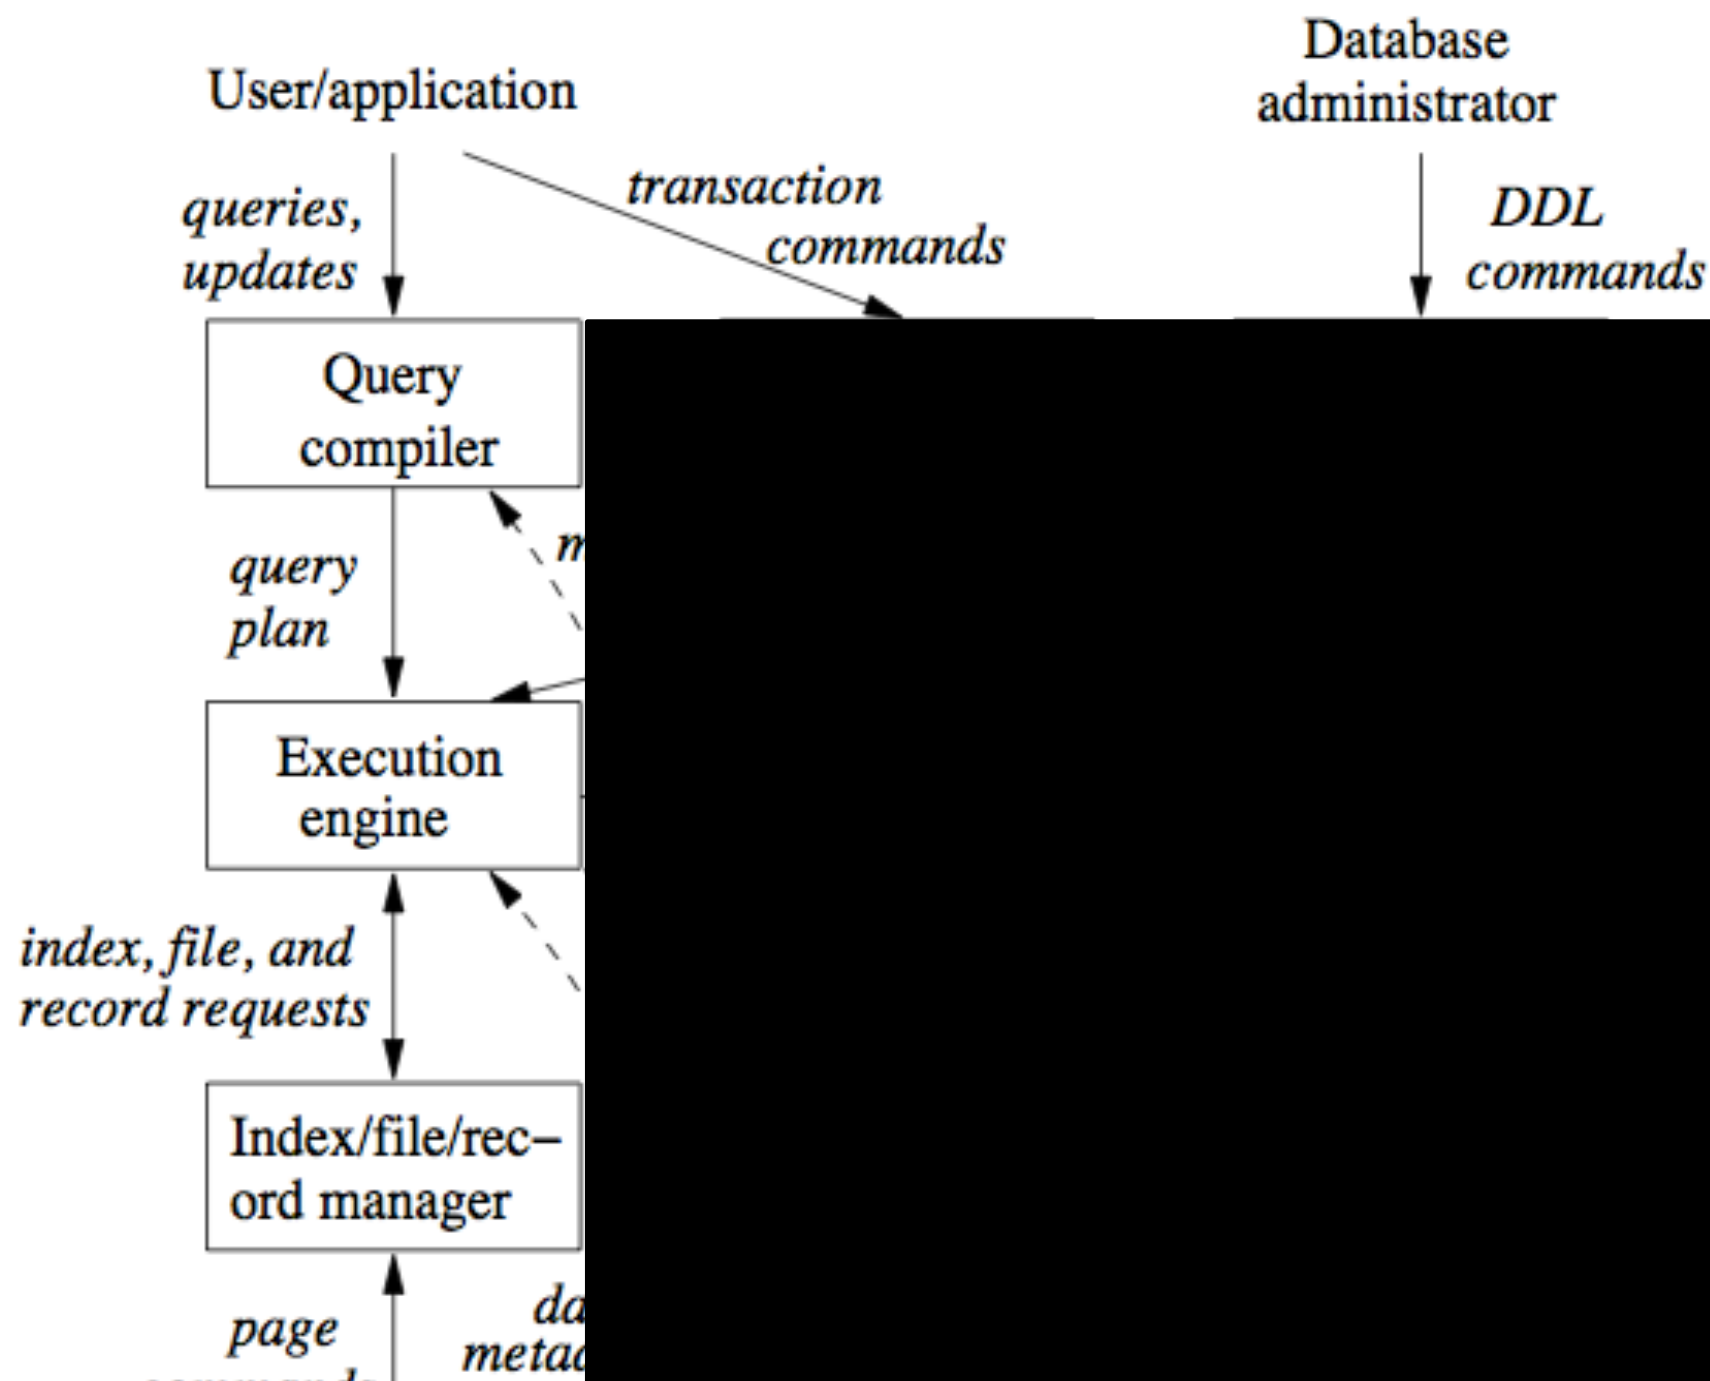
Execution
engine

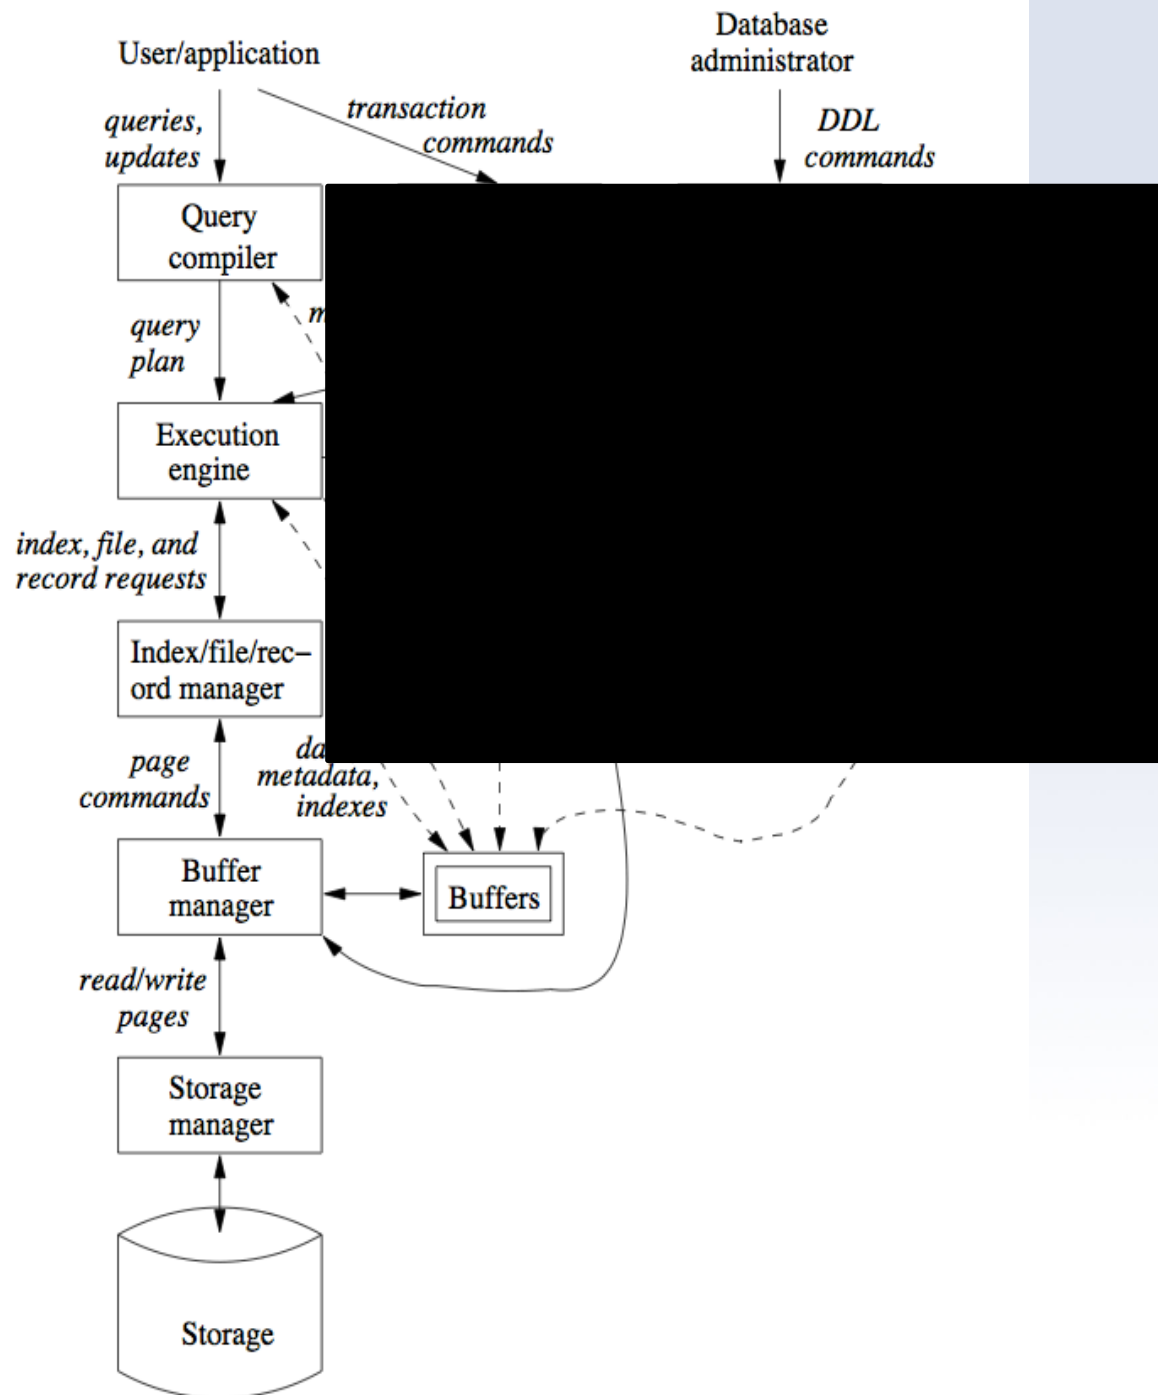
*index, file, and
record requests*

Index/file/rec-
ord manager

page

*data
metadata*





Query Processing

- Can be divided into three steps:
 1. Parse the query into a parse tree
 2. Transform parse tree into logical query plan
 3. Transform logical query plan into physical query plan



Parsing


- Result represented as a tree
- Nodes are either:
 - ***atoms***: keywords, names, attributes, constants, parentheses, operators
 - ***syntactic categories***: names of subparts of a query
 - represented like this: <Category>



Parsing

```
SELECT movieTitle
FROM StarsIn
WHERE starName IN(
    SELECT name
    FROM MovieStar
    WHERE birthdate LIKE '%1960'
);
```





Rules

- Rules:
 - consist of atoms and categories
 - allow us to substitute one thing for another
 - Example:

```
<Query> ::= SELECT <SelList> FROM  
<FromList> WHERE <Condition>
```



Partial Grammar for SQL

$\langle \text{Query} \rangle ::= \text{SELECT } \langle \text{SelList} \rangle \text{ FROM } \langle \text{FromList} \rangle \text{ WHERE } \langle \text{Condition} \rangle$

$\langle \text{SelList} \rangle ::= \langle \text{Attribute} \rangle, \langle \text{SelList} \rangle$

$\langle \text{SelList} \rangle ::= \langle \text{Attribute} \rangle$

$\langle \text{FromList} \rangle ::= \langle \text{Relation} \rangle, \langle \text{FromList} \rangle$

$\langle \text{FromList} \rangle ::= \langle \text{Relation} \rangle$



Partial Grammar for SQL

`<Query> ::= SELECT <SelList> FROM
 <FromList> WHERE <Condition>`

`<Condition> ::= <Attribute>=<Attribute>`

`<Condition> ::= <Attribute> LIKE <Pattern>`

`<Condition> ::= <Condition> AND <Condition>`

`<Condition> ::= <Attribute> IN (<Query>)`



Look familiar?

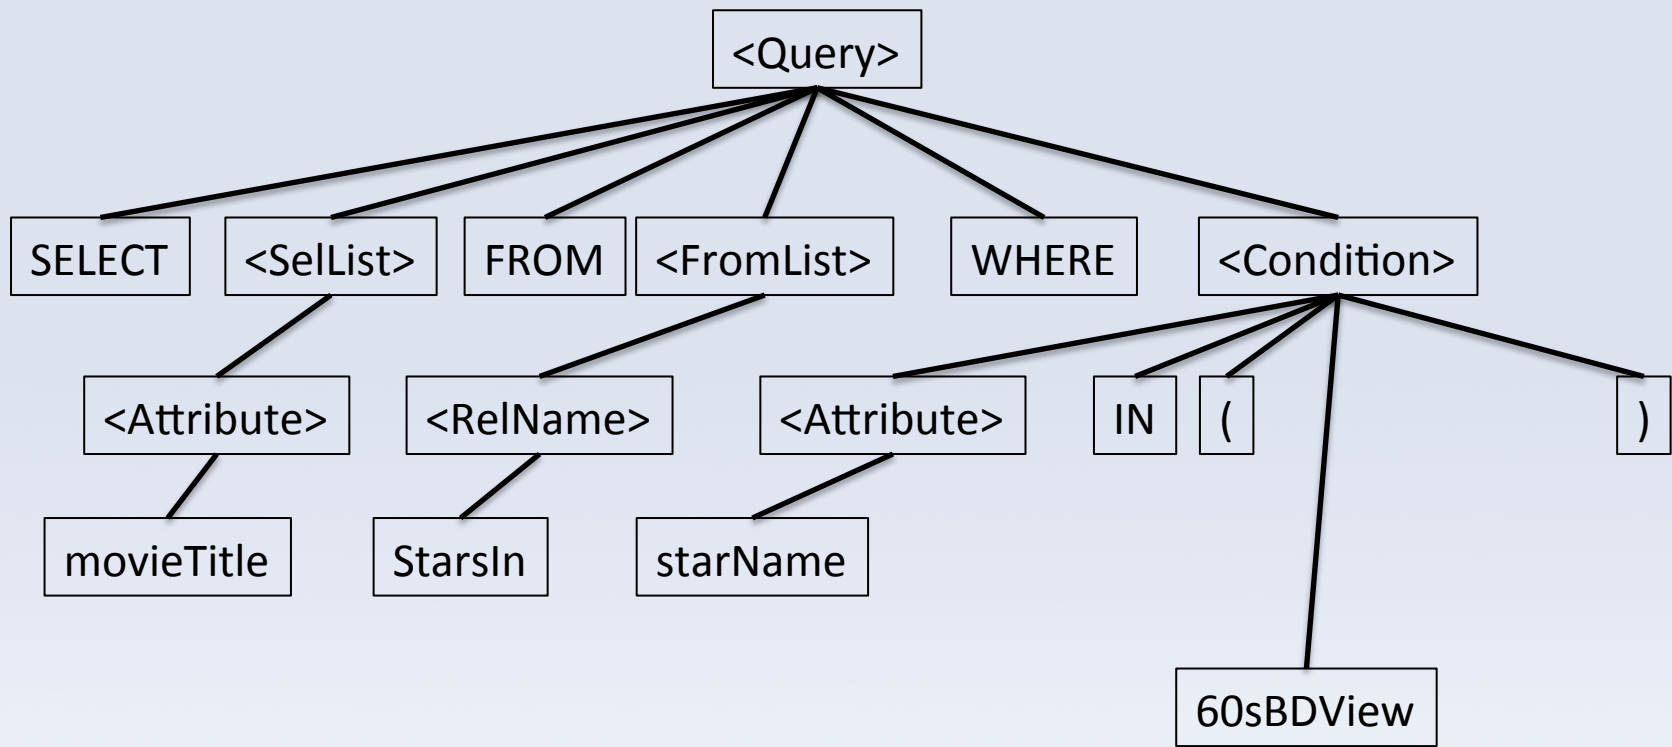
- CS373 or CS421?
- It's a CFG in BNF
- Entire grammar available here:
 - <http://savage.net.au/SQL/>
- We can find the parse tree using CYK, Earley parsing, etc.




Preprocessor

- Parser:
 - checks syntax while building tree
- Preprocessor checks ***semantics***
 - Are relations in our database schema?
 - Are the attributes correct for the schema?
 - Are the types correct?
- Also substitutes subqueries for views







Preprocessor

- You wrote a preprocessor in MP1
 - “is valid” methods checked semantics of the input query



Parse Tree to Logical Query Plan

- All of our algorithms for query execution were based on relational algebra
- Need to convert ***parse tree*** to relational algebra-based ***query plan***



Parse Tree to Logical Query Plan

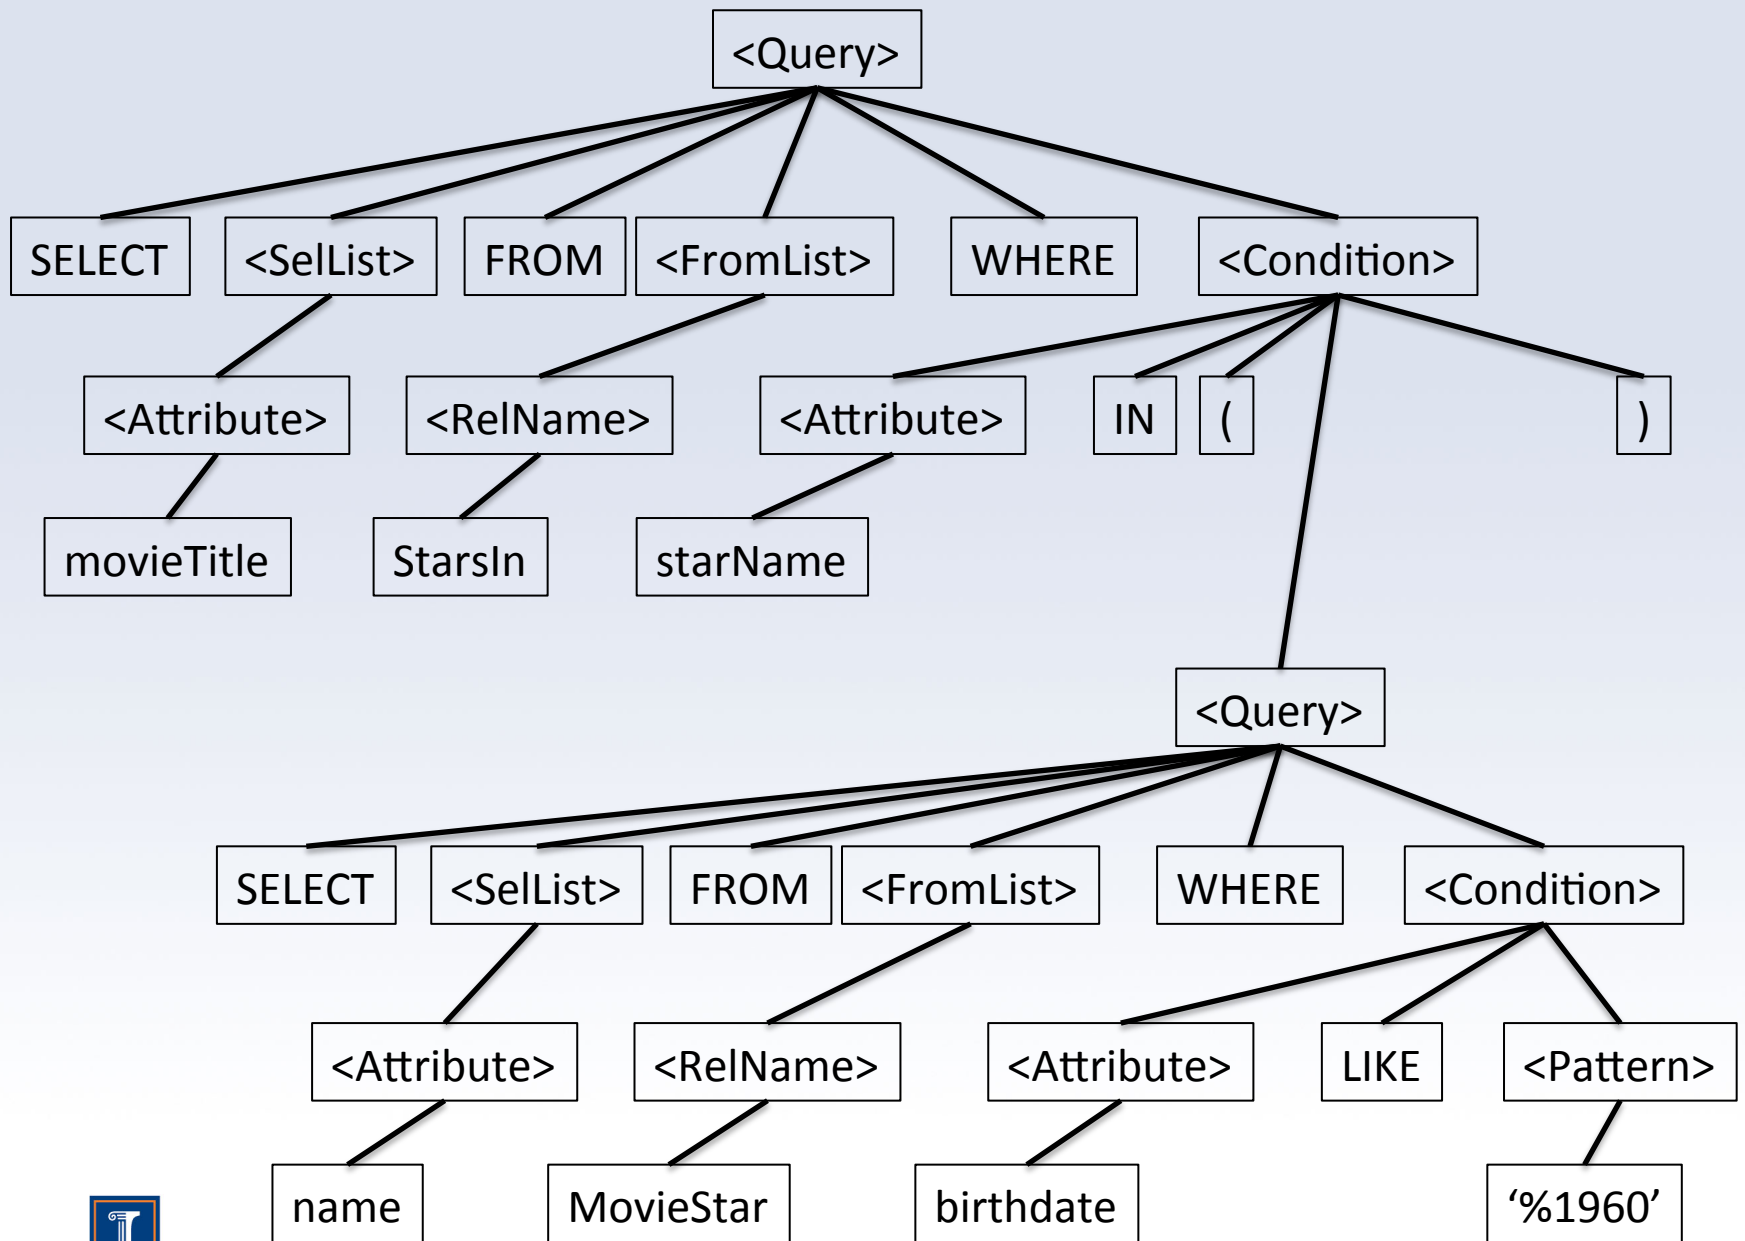
- Replacing $\langle \text{Query} \rangle$ (no subqueries)
 1. take product of all relations in $\langle \text{FromList} \rangle$
 2. σ_C where C is $\langle \text{Condition} \rangle$
 3. π_L where L is the attributes in $\langle \text{SelList} \rangle$

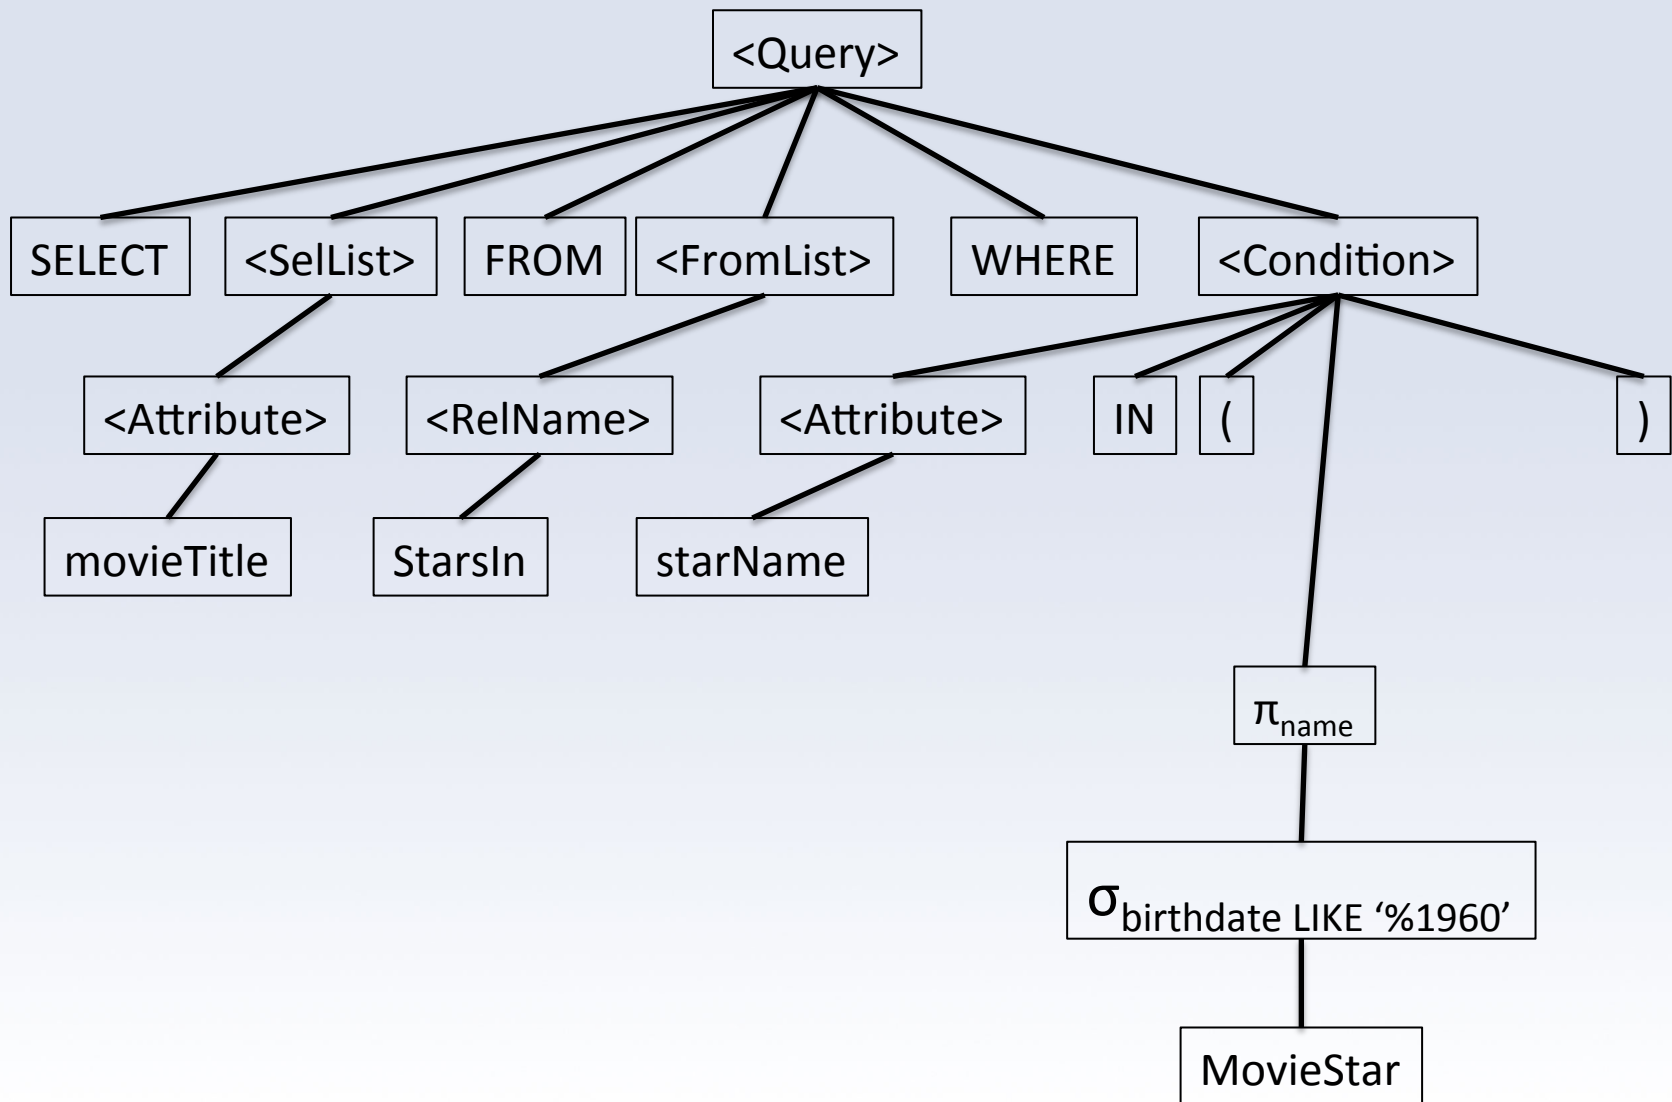


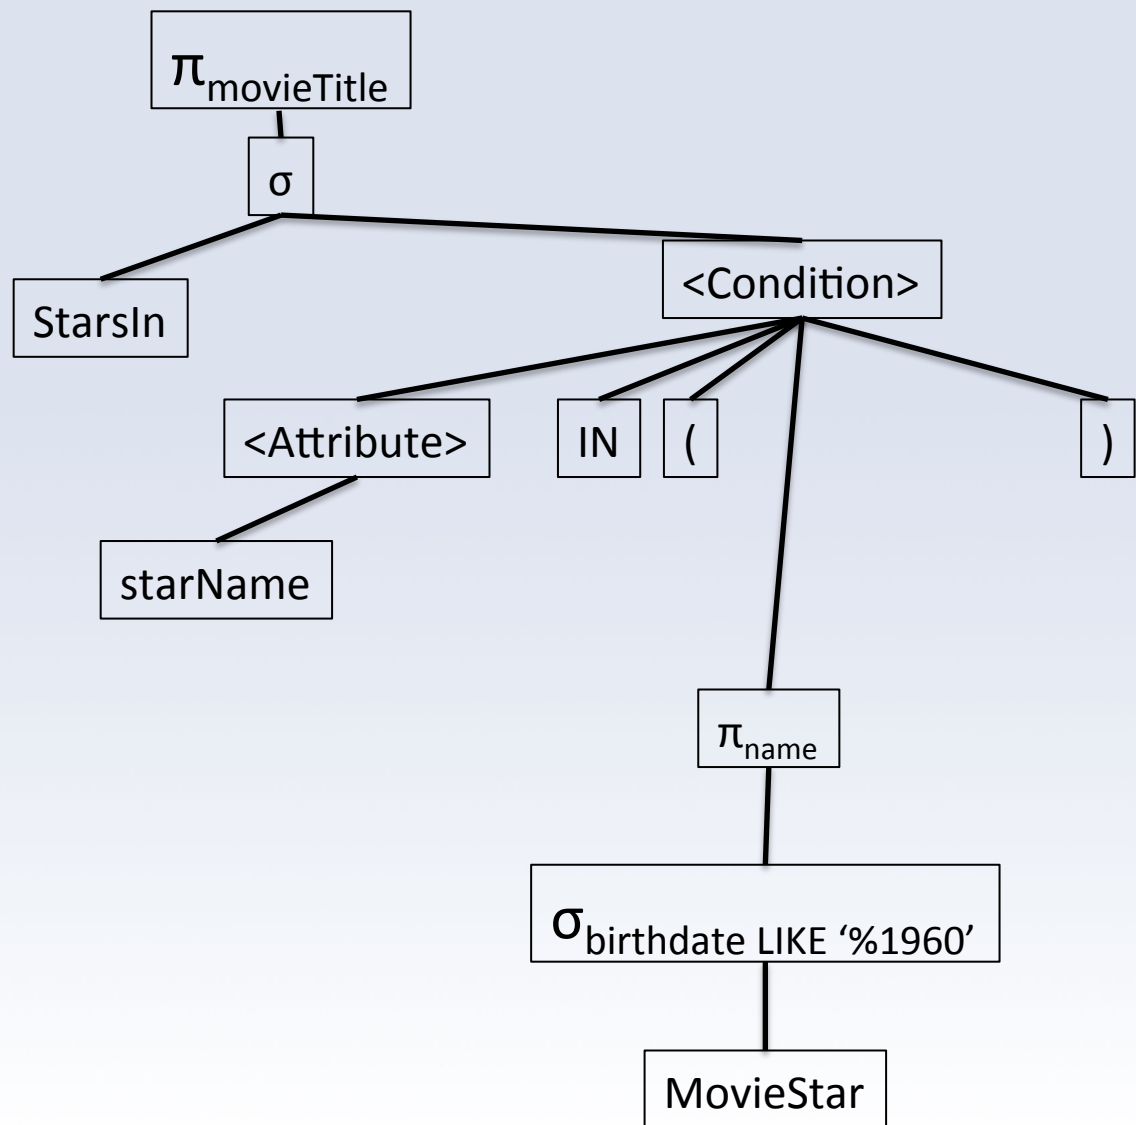
Example

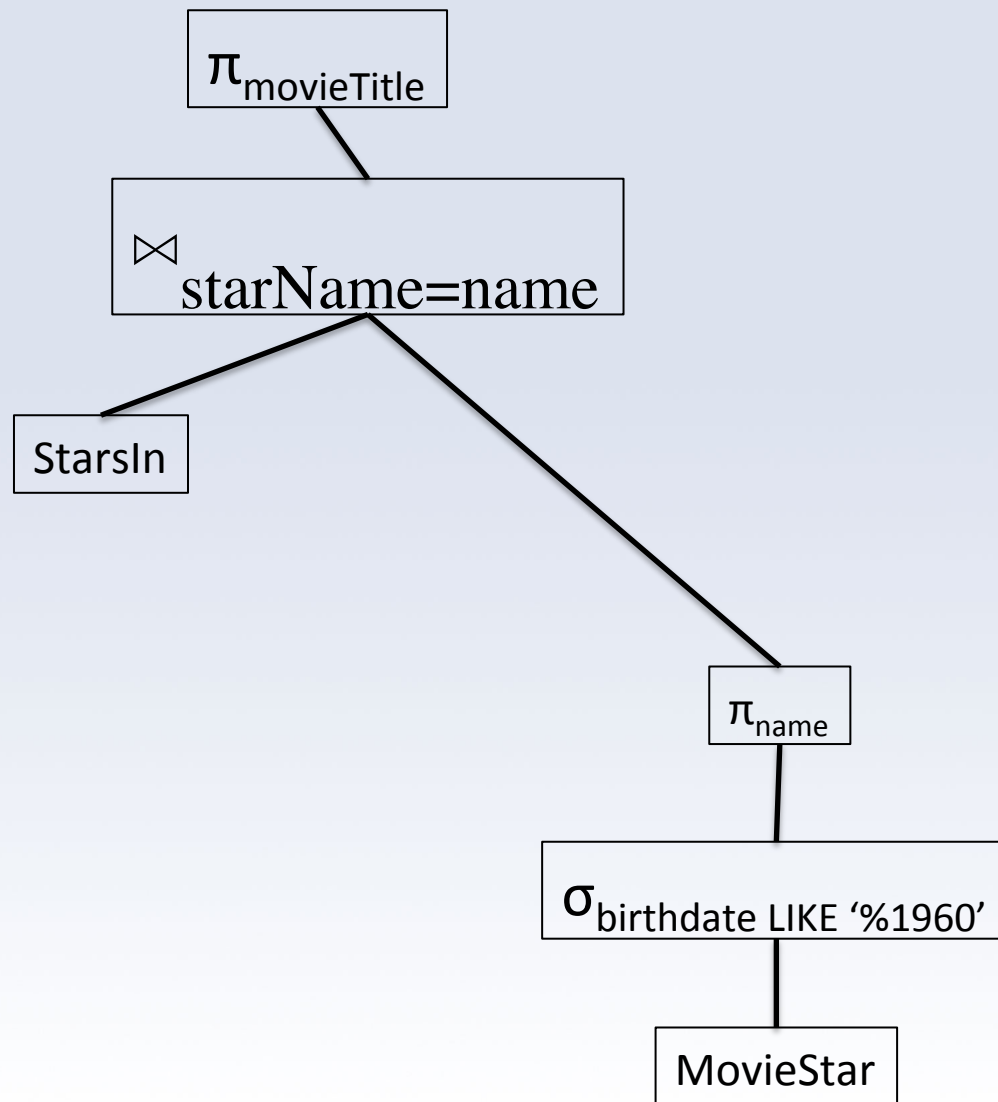
```
SELECT a,b,c  
FROM X,Y,Z  
WHERE d<e AND f=g
```

$$\pi_{a,b,c}(\sigma_{d<e \text{ AND } f=g}(X \times Y \times Z))$$







Parse Tree to Logical Query Plan

- Don't worry about algorithms for:
 - building parse trees
 - converting to query plans
- Subtle hint:
 - You should be able to convert SQL to relational algebra on the final
 - Think about methods for converting subqueries like IN, ALL, etc.



Refining Logical Query Plan

- One query can be expressed many ways
- Some more efficient than others
- Try to pick the best



Query Optimization

- We need three things:
 1. Rewrite rules
 2. An optimization algorithm
 3. A cost estimator



Rewrite Rules

- Algebraic laws for relational algebra
- Similar to $2(x+y)=2x+2y$



Commutative/Associative Laws

$$R \times S = S \times R, (R \times S) \times T = R \times (S \times T)$$

$$R \bowtie S = S \bowtie R, (R \bowtie S) \bowtie T = R \bowtie (S \bowtie T)$$

$$R \cup S = S \cup R, (R \cup S) \cup T = R \cup (S \cup T)$$

$$R \cap S = S \cap R, (R \cap S) \cap T = R \cap (S \cap T)$$

$$R \bowtie (S \cup T) = (R \bowtie S) \cup (R \bowtie T)$$



Selection Laws

$$\sigma_{C_1 \text{ AND } C_2}(R) = \sigma_{C_1}(\sigma_{C_2}(R)) = \sigma_{C_2}(\sigma_{C_1}(R))$$

$$\sigma_{C_1 \text{ OR } C_2}(R) = \sigma_{C_1}(R) \cup \sigma_{C_2}(R)$$

$$\sigma_C(R \times S) = \sigma_C(R) \times \sigma_C(S) = \sigma_C(R) \times S^*$$

$$\sigma_C(R \bowtie S) = \sigma_C(R) \bowtie \sigma_C(S) = \sigma_C(R) \bowtie S^*$$

$$\sigma_C(R \cap S) = \sigma_C(R) \cap \sigma_C(S) = \sigma_C(R) \cap S$$

$$\sigma_C(R \cup S) = \sigma_C(R) \cup \sigma_C(S)$$

$$\sigma_C(R - S) = \sigma_C(R) - \sigma_C(S) = \sigma_C(R) - S$$

*Valid only when C exclusively involves attributes of R



Example

$R(A, B, C, D), S(E, F, G)$

$\sigma_{F=3} (R \bowtie_{D=E} S)$



Example

$R(A, B, C, D), S(E, F, G)$

$\sigma_{A=5 \text{ AND } G=9} (R \bowtie_{D=E} S)$



Projection Laws

- Too many caveats to list
- Easier to summarize:
 - Can push a projection inside any operator that doesn't need the attributes projected out
- Example

$$\pi_{a+b \rightarrow x}(R \bowtie S) = \pi_{a+b \rightarrow x}(\pi_{a,c}(R) \bowtie \pi_{b,c}(S))$$



Join and Product Laws

$$R \bowtie_C S = \sigma_C (R \times S)$$
$$R \bowtie S = \pi_L (\sigma_C (R \times S))$$



Duplicate Elimination Laws

$$\delta(R \times S) = \delta(R) \times \delta(S)$$

$$\delta(R \bowtie S) = \delta(R) \bowtie \delta(S)$$

$$\delta(\sigma_c(R)) = \sigma_c(\delta(R))$$



Grouping and Aggregation

$$\delta(\gamma_L(R)) = \gamma_L(R)$$

$$\gamma_L(R) = \gamma_L(\pi_M(R))^*$$

$$\gamma_L(R) = \gamma_L(\delta(R))^*$$

*Valid only if $L \subseteq M$

*Valid only for MAX and MIN



Example

$\gamma_{movieYear, MAX(birthdate)}(\sigma_{name=starName}(MovieStar \times StarsIn))$



Query Optimization

- We need three things:
 1. ~~Rewrite rules~~
 2. An optimization algorithm
 3. A cost estimator



General Heuristics

- Some general heuristics for query optimization:
 1. Push down selections
 2. Push down/add projections
 3. Push/add duplicate eliminations when needed or remove when not needed
 4. Turn products into equijoins



Can we do better?

- Physical properties of data helpful
 - How big are the relations involved?
 - Are they sorted?
 - Are they indexed?
- Related to the physical plan, but we won't see physical plans until Friday



Cost Estimation

- We will find ways of estimating the **cost** of a query
- Has to be based on the properties of the underlying relation
- We can't know **exact** cost without executing the query, so only an estimation



Cost Estimation

- We start with metadata (data dictionary)
 - $B(R)$ number of blocks holding R
 - $T(R)$ number of tuples in R
 - $V(R,a)$ number of distinct values for attribute a in relation R



Projection Estimation

- Can be computed exactly
 - $B(S) = (T(R) * \text{size}(\text{newtuple})) / \text{blocksize}$
 - $T(S) = T(R)$
 - $V(S, a) = V(R, a)$ (if a in S)



Example

- Given $R(a,b,c)$ where $T(R)=10,000$
 - a,b are 4 byte integers, c is 100 byte string
 - record header=12 bytes
 - block header=24 bytes
 - block size=1024
- Total record size=120 bytes
- Total records per block=8
- $B(R)=1,250$



Example

- $S = \pi_{a+b \rightarrow x, c}(R)$
- Total record size = 116 bytes
- Total records per block = 8
- $T(S) = 10,000$
- $B(S) = 1,250$



Selection Estimation

- For $S = \sigma_{a=c}(R)$
 - $T(S) \approx T(R)/V(R,A)$
- For $S = \sigma_{a < c}(R)$
 - $T(S) \approx T(R)/3$
- Treat AND as multiple nested selections
- Treat OR as independent (more on this later)
 - $T(S) \approx n(1 - (1 - m_1/n)(1 - m_2/n))$



Example

- Given $R(a,b,c)$ with:
 - $T(R)=10,000$
 - $V(R,a)=50, V(R,b)=20, V(R,c)=100$
- Let $S=\sigma_{a=5 \text{ AND } b<100} (R)$
 - $T(S)\approx(10,000/50)/3=67$



Selection with OR

- $S = \sigma_{c_1 \text{ OR } c_2}(R)$
 - Assume m_1 tuples satisfy c_1 , m_2 satisfy c_2
 - Assume conditions are independent
 - $(1 - m_1/n)$ fraction of tuples that don't satisfy c_1
 - $(1 - m_2/n)$ fraction of tuples that don't satisfy c_2
 - $(1 - m_1/n)(1 - m_2/n)$ fraction that satisfy neither
 - $(1 - m_1/n)(1 - m_2/n)$ fraction that satisfy both
 - $T(S) \approx n(1 - (1 - m_1/n)(1 - m_2/n))$



Example

- Given $R(a,b,c)$ with:
 - $T(R)=10,000$
 - $V(R,a)=50, V(R,b)=20, V(R,c)=100$
- Let $S=\sigma_{a=5 \text{ OR } b<100} (R)$
 - Tuples satisfying $a=5 \approx 10k/50=200=m_1$
 - Tuples satisfying $b<100 \approx 3333=m_2$
 - $T(S) \approx (1-(1-200/10k)(1-3333/10k))=3466$



Join Estimation

- Arguably the most important cost
- We need to predict how many tuples will relate between R and S
 - Could be anywhere between 0 and $T(R)T(S)$!
- We'll first work with natural join
 - $U = R(X, Y) \bowtie S(Y, Z)$
- Can easily generalize to other joins



Join Estimation

- Simplifying assumptions:
 1. If join attribute has more values in R than in S, **all** values in S occur in R
 - $V(R,Y) \geq V(S,Y)$
 - in other words, assume every tuple in S joins
 2. Non-join attributes do not lose values
 - $V(R \bowtie S, A) \approx V(R, A)$ for A not in Y
 - $V(R \bowtie S, B) \approx V(S, B)$ for B not in Y



Join Estimation

- Given a pair of tuples, probability they join is $1/\max(V(R,Y),V(S,Y))$
- Total number of pairs of tuples is $T(R)T(S)$
- $T(R \bowtie S) \approx T(R)T(S)/\max(V(R,Y),V(S,Y))$



Example

- $R(a,b,c) \bowtie S(b,c,d)$
 - $T(R)=1000$, $V(R,a)=100$, $V(R,b)=20$, $V(R,c)=200$
 - $T(S)=2000$, $V(S,b)=50$, $V(S,c)=100$, $V(S,d)=400$
- $T(R \bowtie S) \approx 1000 * 2000 / 50 = 40,000$
- $V(R \bowtie S, a) \approx 100$, $V(R \bowtie S, b) \approx 20$, $V(R \bowtie S, c) \approx 200$,
 $V(R \bowtie S, d) \approx 400$



Other Operators

- Assume $T(S) < T(R)$
 - $T(R \cup S) \approx T(R) + T(S)/2$
 - $T(R \cap S) \approx T(S)/2$
 - $T(R - S) \approx T(R) - T(S)/2$
 - $T(\delta(R)) \approx \min(T(R)/2, V(R,a)V(R,b), \dots V(R,z))$
 - $T(\gamma(R)) \approx \min(T(R)/2, V(R,a)V(R,b), \dots V(R,z))$



Example

Given $R(a,b)$ and $S(b,c)$

$$T(R)=5000, V(R,a)=50, V(R,b)=100$$

$$T(S)=2000, V(S,b)=200, V(S,c)=100$$

$$U=\delta(\sigma_{a=10}(R \bowtie S))$$

$$T(R \bowtie S)=5000*2000/100=10000$$

$$T(\sigma_{a=10}(R \bowtie S))=10000/50=200$$

$$T(U)=\min(100,5000)=100$$



Example

Given $R(a,b)$ and $S(b,c)$

$$T(R)=5000, V(R,a)=50, V(R,b)=100$$

$$T(S)=2000, V(S,b)=200, V(S,c)=100$$

$$U = \delta(\sigma_{a=10}(R)) \bowtie \delta(S)$$

$$T(\sigma_{a=10}(R))=100, T(\delta(\sigma_{a=10}(R)))=50$$

$$T(\delta(S))=1000$$

$$T(U)=1000*50/200=250$$



Query Optimization

- We need three things:
 1. ~~Rewrite rules~~
 2. An optimization algorithm
 3. ~~A cost estimator~~



Next Time...

- We'll see:
 - how we use these estimates to find a good plan
 - order of joins is especially important
 - how to represent and choose a physical plan
 - other factors to consider (e.g. indexes, sorts)
 - pipelining vs. materialization

