

UNIVERSITY OF ILLINOIS
AT URBANA-CHAMPAIGN

CS411 - Transaction Management



illinois.edu

Review

- What is “non-quiescent checkpointing”?
- What is an advantage of undo/redo logging?
- Why do we care about scheduling?
- What is a “serial schedule”?
- What is a “serializable schedule”?



Example

$\langle \text{START } T_1 \rangle$

$\langle T_1, A, 4, 5 \rangle$

$\langle \text{START } T_2 \rangle$

$\langle \text{START } T_3 \rangle$

$\langle \text{COMMIT } T_1 \rangle$

$\langle T_2, B, 9, 10 \rangle$

$\langle T_2, C, 14, 15 \rangle$

$\langle \text{START } T_4 \rangle$

$\langle T_3, D, 19, 20 \rangle$

$\langle T_4, E, 30, 31 \rangle$

$\langle \text{COMMIT } T_3 \rangle$

$\langle T_4, F, 43, 44 \rangle$



Example

<START T_1 >

< T_1 ,A,4,5>

<START T_2 >

<COMMIT T_1 >

< T_2 ,B,9,10>

<START CKPT (T_2)>

< T_2 ,C,14,15>

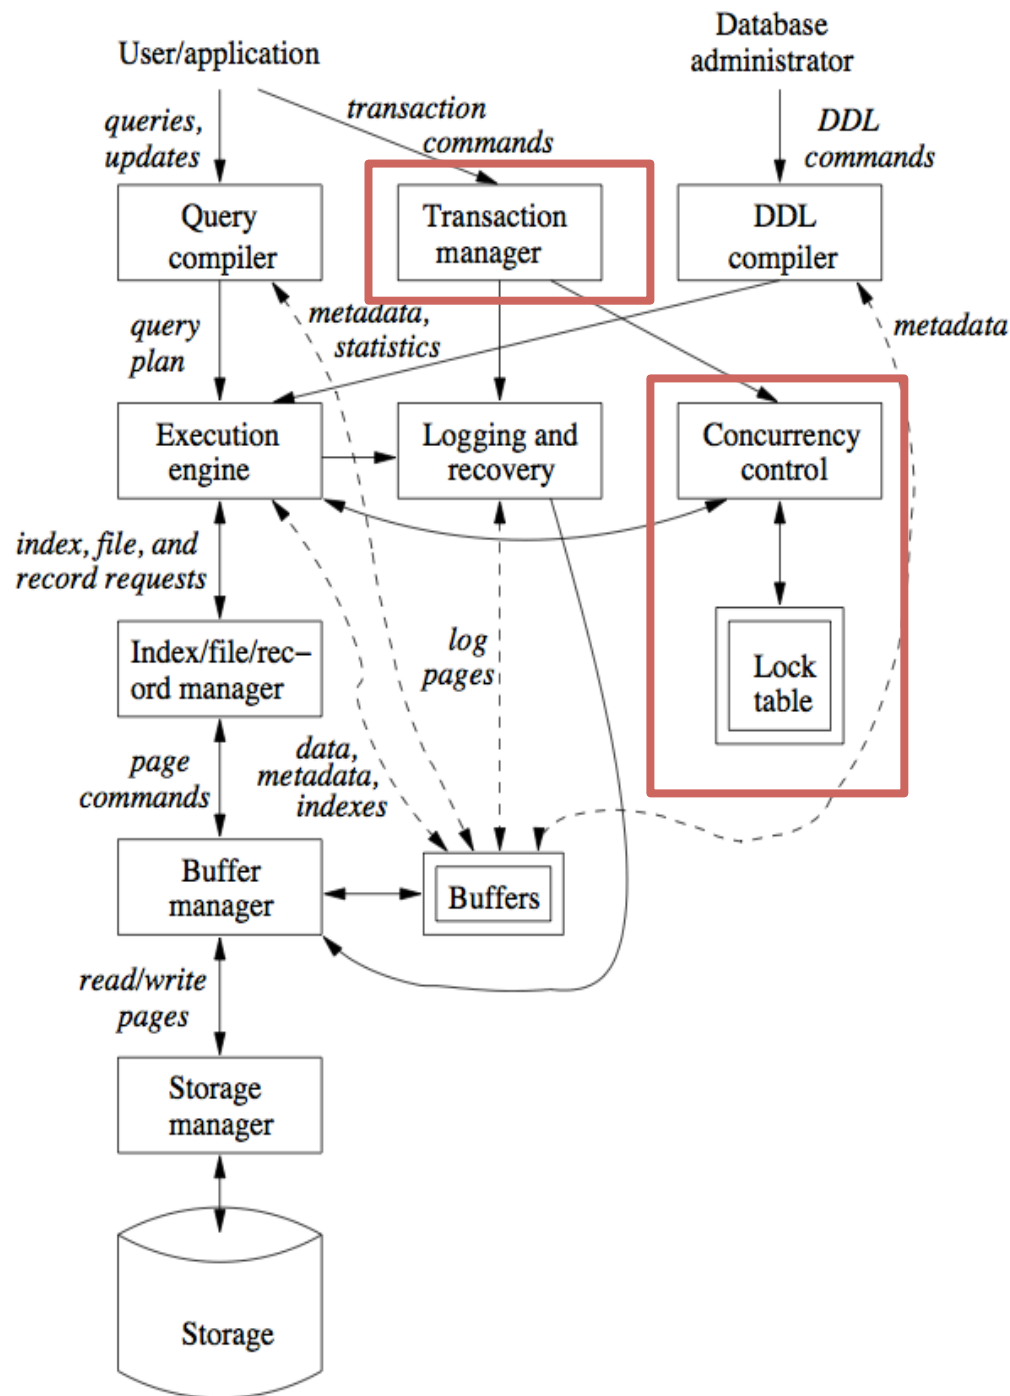
<START T_3 >

< T_3 ,D,19,20>

<END CKPT>

<COMMIT T_2 >





Serializable Schedule

- A schedule is *serializable* if it has the same effect as some serial schedule



Notation

- We don't care about values
- Just ***what*** is accessed, ***who*** accessed it, and ***how*** it is accessed
 - Read: $r_i(X)$ transaction i reads element X
 - Write: $w_i(X)$ transaction i writes element X



Example

T1: $r_1(A)$; $w_1(A)$; $r_1(B)$; $w_1(B)$;

T2: $r_2(A)$; $w_2(A)$; $r_2(B)$; $w_2(B)$;

S: $r_1(A)$; $w_1(A)$; $r_2(A)$; $w_2(A)$; $r_1(B)$; $w_1(B)$; $r_2(B)$; $w_2(B)$;



Conflicting swaps

- A pair of consecutive actions conflict if changing order changes the behavior of a transaction:
 1. two actions from the same transaction
 2. $r_i(X), w_j(X)$ or $w_i(X), r_j(X)$
 3. $w_i(X), w_j(X)$



Conflict-Serializable

- Making a series of non-conflicting swaps to a schedule, we can produce a serial schedule
- Non-conflicting swaps don't change behavior, so conflict-serializable \Rightarrow serializable



Example

r1(A); w1(A); r2(A); w2(A); r1(B); w1(B); r2(B); w2(B);
r1(A); w1(A); r2(A); r1(B); w2(A); w1(B); r2(B); w2(B);
r1(A); w1(A); r1(B); r2(A); w2(A); w1(B); r2(B); w2(B);
r1(A); w1(A); r1(B); r2(A); w1(B); w2(A); r2(B); w2(B);
r1(A); w1(A); r1(B); w1(B); r2(A); w2(A); r2(B); w2(B);



Conflict-Serializable

- Conflict-Serializable schedules are a subset of serializable schedules
 - Conflict-serializable is a “stronger” definition
- Example:
 - $w_1(A); w_2(A); w_2(B); w_1(B); w_3(B);$
 - Executing $T_1, T_2, T_3 \Rightarrow$
 - A gets value from T_2
 - B gets value from T_3



Precedence Graphs

- Checking a schedule to see if it is conflict-serializable
 - Graph based approach
 - A node for each transaction
 - A directed edge from T_i to T_j if $T_i <_S T_j$
 - T_i takes precedence over T_j in schedule S






Precedence Graphs

- $T_i <_S T_j$ if there are actions A_i and A_j :
 - A_i is in T_i and A_j is in T_j
 - A_i happens before A_j
 - A_i and A_j both involve the same database element
 - Either A_i or A_j is a write action
- In other words, T_i and T_j have some conflicting action



Example

$r_2(A); r_1(B); w_2(A); r_3(A); w_1(B); w_3(A); r_2(B); w_2(B);$






Precedence Graphs

- S is conflict-serializable iff its precedence graph has no cycles
 - Intuitively: a loop means the transactions in the loop share actions that can't be swapped to create a serial schedule
 - Formal proof is in the book



Example

$r_2(A); r_1(B); w_2(A); r_2(B); r_3(A); w_1(B); w_3(A); w_2(B);$



Precedence Graphs

- Useful for humans to prove things about schedules
- Not very helpful for a DBMS, though
 - examine all possible schedules?
 - search for schedule with no cycles?



Scheduler

- Assures serializability in real time by:
 - Viewing actions of incoming transactions
 - “Locking” database elements to indicate they are in use: $l_1(A)$
 - “Unlocking” them when transaction finishes with them: $u_1(A)$
 - Delaying actions requiring locked elements



Two-Phase Locking

- ***All*** locks for a transaction ***must*** come before ***all*** unlocks
 - Scheduler inserts locks before all reads and rights
 - Scheduler inserts unlocks after transaction commits or aborts



Example

$l_1(A); w_1(A); l_2(B); r_2(B); w_1(A); u_1(A); l_2(A); w_2(A);$
 $u_2(A); u_2(B);$



Serializability

- A two-phase locking scheduler ***guarantees*** conflict-serializability
 - Intuitively:
 - all actions for a transaction could be moved just before unlock phase
 - unlocks happen serially, so schedule is serializable
 - Again, we won't prove this



Deadlock

- Example:
 - T_1 has a lock on A, waiting for B
 - T_2 has a lock on B, waiting for A
- Two-phase lock scheduler can't stop this
 - Deadlock resolution is covered enough in other classes
 - We won't cover it



Lock Modes

- Practical schemes have multiple modes
 - Shared lock: $sl_1(A)$
 - **Multiple transactions** can have a shared lock on one element
 - Can only **read** the element, though
 - Exclusive lock: $xl_1(A)$
 - Only **one transaction** can have an exclusive lock
 - No shared locks allowed on the element
 - Can **read or write** the element



Compatibility Matrix

		Lock Requested	
Lock Held		S	X
	S	Yes	No
	X	No	No



Example

$sl_1(A); w_1(A);$

$sl_2(A); r_2(A);$

$sl_2(B); r_2(B);$

$xl_1(B); \mathbf{DENIED}$

$u_2(A); u_2(B)$

$xl_1(A); r_1(B); w_1(B)$

$u_1(A); u_2(B)$



Scheduler Architecture

- How is this scheme implemented in a DBMS?
 - need to track locking and unlocking of elements
 - need to decide when to grant or deny a lock
 - need to delay transactions whose locks are denied



Scheduler Architecture

- Locks are entered in a ***lock table***
 - A relation of database elements and lock information

A	
B	
C	
D	
E	
F	
G	

Group Mode: S

Waiting: Yes

List:

Tran	Mode	Wait?	Tnext	Next
T ₁	S	no	NULL	

Tran	Mode	Wait?	Tnext	Next
T ₂	S	no	NULL	

Tran	Mode	Wait?	Tnext	Next
T ₃	X	yes	NULL	NULL



Scheduler Architecture

- Group Mode: current lock mode for B
- Waiting: is anyone waiting to lock B
- List: (next slide)

A	
B	
C	
D	
E	
F	
G	

Group Mode: U

Waiting: Yes

List:

Tran	Mode	Wait?	Tnext	Next
T ₁	S	no	NULL	

Tran	Mode	Wait?	Tnext	Next
T ₂	S	no	NULL	

Tran	Mode	Wait?	Tnext	Next
T ₃	X	yes	NULL	NULL



Scheduler Architecture

- Tran: transaction for this entry
- Mode: lock mode for transaction
- Waiting: status of wait

A	
B	
C	
D	
E	
F	
G	

Group Mode: S
Waiting: Yes
List:

Tran	Mode	Wait?	Tnext	Next
T ₁	S	no	NULL	

Tran	Mode	Wait?	Tnext	Next
T ₂	S	no	NULL	

Tran	Mode	Wait?	Tnext	Next
T ₃	X	yes	NULL	NULL



Handling locks

- If transaction T needs a lock on A
 - No lock-table entry: grant lock and create entry
 - Lock-table entry: check group mode and grant/deny as appropriate
 - Add to list and indicate waiting or not waiting



Handling unlocks

- If transaction T unlocks A
 - Delete T from list
 - Update group mode by examining remaining entries
 - If waiting is “yes”, search list to see if any waiting locks can be granted




Locking Elements

- “Elements” is vague
- What are we locking?
 - Tables?
 - Blocks?
 - Tuples?
- Yes.



Hierarchical Locking



Warning Locks

- To lock an element with S or X, work down the hierarchy tree
- If element we want to lock is down the tree, request IS or IX lock on this node
 - If granted, go lower. If not, wait.
- When we get to element we want, request S or X lock



Compatibility Matrix

		Lock Requested			
Lock Held		IS	IX	S	X
	IS	Yes	Yes	Yes	No
	IX	Yes	Yes	No	No
	S	Yes	No	Yes	No
	X	No	No	No	No



Example

- Page 925



Read phenomena

- Three types:
 1. Phantom reads
 2. Non-repeatable reads
 3. Dirty reads



Phantom Reads

- Can't lock tuples that haven't been inserted yet
 - If another transaction inserts a new tuple, it won't know about our read locks
 - We might see part of the insert
- To avoid this, we can lock the entire table



Non-repeatable Reads

- Release read locks before committing
 - Other transactions might modify the data before our transaction commits
 - But we won't see data that hasn't been committed



Dirty reads

- Don't bother getting read locks
 - Just get whatever the current value is
 - Other transactions could modify the data
 - Those transactions might even be rolled back



Further reading

- Chapter 19 covers:
 - interaction between logging and transaction management
 - how to deal with deadlocks
- [http://msdn.microsoft.com/en-us/library/aa213039\(v=sql.80\).aspx](http://msdn.microsoft.com/en-us/library/aa213039(v=sql.80).aspx)
- <http://dev.mysql.com/doc/refman/5.0/en/internal-locking.html>

