

Notes 7: Polynomials and Interactions

Nathaniel E. Helwig

Department of Statistics
University of Illinois at Urbana-Champaign

Stat 420: Methods of Applied Statistics
Section N1U/N1G – Spring 2014

Outline of Notes

1) Polynomial Regression:

- Polynomials review
- Model form
- Model assumptions
- OLS estimation
- Orthogonal polynomials
- Example: MPG vs HP

2) Interactions in Regression:

- Overview
- Nominal*Continuous
- Example #1: Real Estate
- Example #2: Depression
- Continuous*Continuous
- Example #3: Oceanography

Polynomial Function: Definition

Reminder: a polynomial function has the form

$$\begin{aligned} f(x) &= a_0 + a_1x + a_2x^2 + a_3x^3 + \cdots + a_nx^n \\ &= \sum_{j=0}^n a_jx^j \end{aligned}$$


where $a_j \in \mathbb{R}$ are the coefficients and x is the indeterminate (variable).

Note: x^j is the j -th order polynomial term


- x is first order term, x^2 is second order term, etc.
- The *degree* of a polynomial is the highest order term

Polynomial Function: Simple Regression

```
> x=seq(-1,1,l=50)
> y=2+2*(x^2)
> plot(x,y,main="Quadratic")
> qmod=lm(y~x)
> abline(qmod)
```



```
> x=seq(-1,1,l=50)
> y=2+2*(x^3)
> plot(x,y,main="Cubic")
> cmod=lm(y~x)
> abline(cmod)
```



Model Form (scalar)

The polynomial regression model has the form

$$y_i = b_0 + \sum_{j=1}^p b_j x_i^j + e_i$$

for $i \in \{1, \dots, n\}$ where

- $y_i \in \mathbb{R}$ is the real-valued response for the i -th observation
- $b_0 \in \mathbb{R}$ is the regression intercept
- $b_j \in \mathbb{R}$ is the regression slope for the j -th degree polynomial
- $x_i \in \mathbb{R}$ is the predictor for the i -th observation
- $e_i \stackrel{\text{iid}}{\sim} \mathcal{N}(0, \sigma^2)$ is Gaussian measurement error

Model Form (matrix)

The polynomial regression model has the form

$$\mathbf{y} = \mathbf{X}\mathbf{b} + \mathbf{e}$$

or

$$\begin{pmatrix} y_1 \\ y_2 \\ y_3 \\ \vdots \\ y_n \end{pmatrix} = \begin{pmatrix} 1 & x_1 & x_1^2 & \cdots & x_1^p \\ 1 & x_2 & x_2^2 & \cdots & x_2^p \\ 1 & x_3 & x_3^2 & \cdots & x_3^p \\ \vdots & \vdots & \vdots & \ddots & \vdots \\ 1 & x_n & x_n^2 & \cdots & x_n^p \end{pmatrix} \begin{pmatrix} b_0 \\ b_1 \\ b_2 \\ \vdots \\ b_p \end{pmatrix} + \begin{pmatrix} e_1 \\ e_2 \\ e_3 \\ \vdots \\ e_n \end{pmatrix}$$

Note that this is still a linear model, even though we have polynomial terms in the design matrix.

PR Model Assumptions (scalar)

The fundamental assumptions of the PR model are:

- 1 Relationship between x and y is polynomial
- 2 x_i and y_i are observed random variables (constants)
- 3 $e_i \stackrel{\text{iid}}{\sim} N(0, \sigma^2)$ is an unobserved random variable
- 4 b_0, b_1, \dots, b_p are unknown constants
- 5 $(y_i|x_i) \stackrel{\text{ind}}{\sim} N(b_0 + \sum_{j=1}^p b_j x_i^j, \sigma^2)$
note: homogeneity of variance

Note: focus is estimation of the polynomial curve.

PR Model: Assumptions (matrix)

In matrix terms, the error vector is multivariate normal:

$$\mathbf{e} \sim N(\mathbf{0}_n, \sigma^2 \mathbf{I}_n)$$

In matrix terms, the response vector is multivariate normal given \mathbf{X} :

$$(\mathbf{y}|\mathbf{X}) \sim N(\mathbf{X}\mathbf{b}, \sigma^2 \mathbf{I}_n)$$

Polynomial Regression: Properties

Some important properties of the PR model include:

- 1 Need $n > p$ to fit the polynomial regression model
- 2 Setting $p = 1$ produces simple linear regression
- 3 Setting $p = 2$ is quadratic polynomial regression
- 4 Setting $p = 3$ is cubic polynomial regression
- 5 Rarely set $p > 3$; use cubic spline instead

Polynomial Regression: OLS Estimation

The ordinary least squares (OLS) problem is

$$\min_{\mathbf{b} \in \mathbb{R}^{p+1}} \|\mathbf{y} - \mathbf{X}\mathbf{b}\|^2$$

where $\|\cdot\|$ denotes the Frobenius norm.

The OLS solution has the form

$$\hat{\mathbf{b}} = (\mathbf{X}'\mathbf{X})^{-1}\mathbf{X}'\mathbf{y}$$

which is the same formula from SLR and MLR!

Fitted Values and Residuals

SCALAR FORM:

Fitted values are given by

$$\hat{y}_i = \hat{b}_0 + \sum_{j=1}^p \hat{b}_j x_i^j$$

and *residuals* are given by

$$\hat{e}_i = y_i - \hat{y}_i$$

MATRIX FORM:

Fitted values are given by

$$\hat{\mathbf{y}} = \mathbf{X}\hat{\mathbf{b}} = \mathbf{H}\mathbf{y}$$

and *residuals* are given by

$$\hat{\mathbf{e}} = \mathbf{y} - \hat{\mathbf{y}} = (\mathbf{I}_n - \mathbf{H})\mathbf{y}$$

Estimated Error Variance (Mean Squared Error)

The estimated error variance is

$$\begin{aligned}\hat{\sigma}^2 &= SSE/(n - p - 1) \\ &= \sum_{i=1}^n (y_i - \hat{y}_i)^2 / (n - p - 1) \\ &= \|(\mathbf{I}_n - \mathbf{H})\mathbf{y}\|^2 / (n - p - 1)\end{aligned}$$

which is an unbiased estimate of error variance σ^2 .

The estimate $\hat{\sigma}^2$ is the *mean squared error (MSE)* of the model.

Distribution of Estimator, Fitted Values, and Residuals

Just like in SLR and MLR, the PR assumptions imply that

$$\hat{\mathbf{b}} \sim N(\mathbf{b}, \sigma^2(\mathbf{X}'\mathbf{X})^{-1})$$

$$\hat{\mathbf{y}} \sim N(\mathbf{X}\mathbf{b}, \sigma^2\mathbf{H})$$

$$\hat{\mathbf{e}} \sim N(\mathbf{0}, \sigma^2(\mathbf{I}_n - \mathbf{H}))$$

Typically σ^2 is unknown, so we use the MSE $\hat{\sigma}^2$ in practice.

Multicollinearity: Problem

Note that x_i , x_i^2 , x_i^3 , etc. can be highly correlated with one another, which introduces multicollinearity problem.

```
> set.seed(123)
> x=runif(100)*2
> X=cbind(x,xsq=x^2,xcu=x^3)
> cor(X)
```

	x	xsq	xcu
x	1.0000000	0.9703084	0.9210726
xsq	0.9703084	1.0000000	0.9866033
xcu	0.9210726	0.9866033	1.0000000

Multicollinearity: Partial Solution

You could mean-center the x_i terms to reduce multicollinearity.

```
> set.seed(123)
> x=runif(100)*2
> x=x-mean(x)
> X=cbind(x,xsq=x^2,xcu=x^3)
> cor(X)
```

	x	xsq	xcu
x	1.00000000	0.03854803	0.91479660
xsq	0.03854803	1.00000000	0.04400704
xcu	0.91479660	0.04400704	1.00000000

But this doesn't fully solve our problem...

Orthogonal Polynomials: Definition

To deal with multicollinearity, define the set of variables

$$z_0 = a_0$$

$$z_1 = a_1 + b_1 x$$

$$z_2 = a_2 + b_2 x + c_2 x^2$$

$$z_3 = a_3 + b_3 x + c_3 x^2 + d_3 x^3$$

where the coefficients are chosen so that $z_j' z_k = 0$ for all $j \neq k$.


The transformed z_j variables are called *orthogonal polynomials*.

Orthogonal Polynomials: Orthogonal Projection

The orthogonal projection of a vector $\mathbf{v} = \{v_i\}_{n \times 1}$ on to the line spanned by the vector $\mathbf{u} = \{u_i\}_{n \times 1}$ is

$$\text{proj}_{\mathbf{u}}(\mathbf{v}) = \frac{\langle \mathbf{v}, \mathbf{u} \rangle}{\langle \mathbf{u}, \mathbf{u} \rangle} \mathbf{u}$$

where $\langle \mathbf{u}, \mathbf{v} \rangle = \mathbf{u}'\mathbf{v}$ and $\langle \mathbf{u}, \mathbf{u} \rangle = \mathbf{u}'\mathbf{u}$ denote the inner products.



<http://thejuniverse.org/PUBLIC/LinearAlgebra/LOLA/dotProd/proj.html>

Orthogonal Polynomials: Gram-Schmidt

Can use the Gram-Schmidt process to form orthogonal polynomials.

Start with a linearly independent design matrix $\mathbf{X} = [\mathbf{x}_0, \mathbf{x}_1, \mathbf{x}_3, \mathbf{x}_4]$ where $\mathbf{x}_j = (x_1^j, \dots, x_n^j)'$ is the j -th order polynomial vector.

Gram-Schmidt algorithm to form columnwise orthogonal matrix \mathbf{Z} that spans the same column space as \mathbf{X} :

$$\mathbf{z}_0 = \mathbf{x}_0$$

$$\mathbf{z}_1 = \mathbf{x}_1 - \text{proj}_{\mathbf{z}_0}(\mathbf{x}_1)$$

$$\mathbf{z}_2 = \mathbf{x}_2 - \text{proj}_{\mathbf{z}_0}(\mathbf{x}_2) - \text{proj}_{\mathbf{z}_1}(\mathbf{x}_2)$$

$$\mathbf{z}_3 = \mathbf{x}_3 - \text{proj}_{\mathbf{z}_0}(\mathbf{x}_3) - \text{proj}_{\mathbf{z}_1}(\mathbf{x}_3) - \text{proj}_{\mathbf{z}_2}(\mathbf{x}_3)$$

Orthogonal Polynomials: R Functions

Simple R function to orthogonalize an input matrix:

```
orthog<-function(X,normalize=FALSE) {  
  np=dim(X)  
  Z=matrix(0,np[1],np[2])  
  Z[,1]=X[,1]  
  for(k in 2:np[2]){  
    Z[,k]=X[,k]  
    for(j in 1:(k-1)){  
      Z[,k]=Z[,k]-Z[,j]*sum(Z[,k]*Z[,j])/sum(Z[,j]^2)  
    }  
  }  
  if(normalize){Z=Z%*%diag(colSums(Z^2)^-0.5)}  
  Z  
}
```

Orthogonal Polynomials: R Functions (continued)

```

> set.seed(123)
> X=cbind(1,runif(10),runif(10))
> crossprod(X)
      [,1]      [,2]      [,3]
[1,] 10.000000  5.782475  5.233693
[2,]  5.782475  4.125547  2.337238
[3,]  5.233693  2.337238  3.809269
> Z=orthog(X)
> crossprod(Z)
      [,1]      [,2]      [,3]
[1,] 1.000000e+01 -4.440892e-16 -4.440892e-16
[2,] -4.440892e-16  7.818448e-01 -1.387779e-17
[3,] -4.440892e-16 -1.387779e-17  4.627017e-01
> Z=orthog(X,norm=TRUE)
> crossprod(Z)
      [,1]      [,2]      [,3]
[1,] 1.000000e+00 -1.942890e-16 -2.220446e-16
[2,] -1.942890e-16  1.000000e+00  1.387779e-17
[3,] -2.220446e-16  1.387779e-17  1.000000e+00

```

Orthogonal Polynomials: R Functions (continued)

Can also use the default `poly` function in R.

```
> set.seed(123)
> x=runif(10)
> X=cbind(1,x,xsq=x^2,xcu=x^3)
> Z=orthog(X,norm=TRUE)
> z=poly(x,degree=3)
> Z[,2:4]=Z[,2:4]%*%diag(colSums(z^2)^0.5)
> Z[1:3,]
```

	[,1]	[,2]	[,3]	[,4]
[1,]	0.3162278	-0.3287304	-0.07537277	0.5363745
[2,]	0.3162278	0.2375627	-0.06651752	-0.5097714
[3,]	0.3162278	-0.1914349	-0.26206273	0.2473705

```
> cbind(Z[1:3,1],z[1:3,])
```

	1	2	3	
[1,]	0.3162278	-0.3287304	-0.07537277	0.5363745
[2,]	0.3162278	0.2375627	-0.06651752	-0.5097714
[3,]	0.3162278	-0.1914349	-0.26206273	0.2473705

Real Polynomial Data

Auto-MPG data from the UCI Machine Learning repository:

<http://archive.ics.uci.edu/ml/datasets/Auto+MPG>


Have variables collected from $n = 398$ cars from years 1970–1982.

mpg	miles per gallon
cylinder	number of cylinders
disp	displacement
hp	horsepower
weight	weight
accel	acceleration
year	model year
origin	origin
name	make and model

Best Linear Relationship

Best linear relationship predicting `mpg` from `hp`.


```
> plot(hp, mpg)
> linmod=lm(mpg~hp)
> abline(linmod)
```



Best Quadratic Relationship

Best quadratic relationship predicting mpg from hp .

```
> quadmod=lm(mpg~hp+I(hp^2))  
> hpseq=seq(50,250,by=5)  
> Xmat=cbind(1,hpseq,hpseq^2)  
> hphat=Xmat%*%quadmod$coef  
> plot(hp,mpg)  
> lines(hpseq,hphat)
```



Best Cubic Relationship

Check for possible cubic relationship:

```
> cubmod=lm(mpg~hp+I(hp^2)+I(hp^3))
> summary(cubmod)
```

```
Call:
lm(formula = mpg ~ hp + I(hp^2) + I(hp^3))
```

Residuals:

Min	1Q	Median	3Q	Max
-14.7039	-2.4491	-0.1519	2.2035	15.8159

Coefficients:

	Estimate	Std. Error	t value	Pr(> t)
(Intercept)	6.068e+01	4.563e+00	13.298	< 2e-16 ***
hp	-5.689e-01	1.179e-01	-4.824	2.03e-06 ***
I(hp^2)	2.079e-03	9.479e-04	2.193	0.0289 *
I(hp^3)	-2.147e-06	2.378e-06	-0.903	0.3673

```
---
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```

```
Residual standard error: 4.375 on 388 degrees of freedom
(6 observations deleted due to missingness)
```

```
Multiple R-squared:  0.6882, Adjusted R-squared:  0.6858
F-statistic: 285.5 on 3 and 388 DF,  p-value: < 2.2e-16
```

Orthogonal versus Raw Polynomials

Compare orthogonal and raw polynomials:

```
> quadomod=lm(mpg~poly(hp,degree=2))
> summary(quadomod)$coef
```

	Estimate	Std. Error	t value	Pr(> t)
(Intercept)	23.44592	0.2209163	106.13030	2.752212e-289
poly(hp, degree = 2)1	-120.13774	4.3739206	-27.46683	4.169400e-93
poly(hp, degree = 2)2	44.08953	4.3739206	10.08009	2.196340e-21

```
> summary(quadmod)$coef
```

	Estimate	Std. Error	t value	Pr(> t)
(Intercept)	56.900099702	1.8004268063	31.60367	1.740911e-109
hp	-0.466189630	0.0311246171	-14.97816	2.289429e-40
I(hp^2)	0.001230536	0.0001220759	10.08009	2.196340e-21

Orthogonal Polynomials from Scratch

We can reproduce the same significance test results using `orthog`:

```
> widx=which(is.na(hp)==FALSE)
> hp=hp[widx]
> mpg=mpg[widx]
> X=orthog(cbind(1, hp, hp^2))
> quadomod=lm(mpg~X[,2]+X[,3])
> summary(quadomod)$coef
```

	Estimate	Std. Error	t value	Pr(> t)
(Intercept)	23.445918367	0.2209163488	106.13030	2.752212e-289
X[, 2]	-0.157844733	0.0057467395	-27.46683	4.169400e-93
X[, 3]	0.001230536	0.0001220759	10.08009	2.196340e-21

Interaction Term: Definition

MLR model with two predictors and an interaction

$$y_i = b_0 + b_1 x_{i1} + b_2 x_{i2} + b_3 x_{i1} x_{i2} + e_i$$

where

- $y_i \in \mathbb{R}$ is the real-valued response for the i -th observation
- $b_0 \in \mathbb{R}$ is the regression intercept
- $b_1 \in \mathbb{R}$ is the first predictor's main effect slope
- $b_2 \in \mathbb{R}$ is the second predictor's main effect slope
- $b_3 \in \mathbb{R}$ is the interaction effect slope
- $e_i \stackrel{\text{iid}}{\sim} \mathcal{N}(0, \sigma^2)$ is Gaussian measurement error

Interaction Term: Interpretation

An interaction between X_1 and X_2 means that the relationship between X_1 and Y differs depending on the value of X_2 (and vice versa).

Pro: model is more flexible (i.e., we've added a parameter)

Con: model is (sometimes) more difficult to interpret.

Nominal Variables

Suppose that $X \in \{x_1, \dots, x_g\}$ is a nominal variable with g levels.

- Nominal variables are also called categorical variables
- Example: $\text{sex} \in \{\text{female}, \text{male}\}$ has two levels
- Example: $\text{drug} \in \{A, B, C\}$ has three levels

To code a nominal variable (with g levels) in a regression model, we need to include $g - 1$ different variables in the model.

- Use “dummy” coding to absorb g -th level into intercept
- $$x_{ij} = \begin{cases} 1 & \text{if } i\text{-th observation is in } j\text{-th level} \\ 0 & \text{otherwise} \end{cases}$$
for $j \in \{1, \dots, g - 1\}$

Nominal Interaction with Two Levels

Revisit the MLR model with two predictors and an interaction

$$y_i = b_0 + b_1 x_{i1} + b_2 x_{i2} + b_3 x_{i1} x_{i2} + e_i$$

and suppose that $x_{i2} \in \{0, 1\}$ is a nominal predictor.

In this case, the model terms can be interpreted as:

- b_0 is expected value of Y when $x_{i1} = x_{i2} = 0$
- b_1 is expected change in Y corresponding to 1-unit change in x_{i1} for observations with $x_{i2} = 0$
- b_2 is expected change in Y corresponding to $x_{i2} = 1$ level
- $(b_1 + b_3)$ is expected change in Y corresponding to 1-unit change in x_{i1} for observations with $x_{i2} = 1$

Real Estate Data Description

Using house price data from Kutner, Nachtsheim, Neter, and Li (2005).

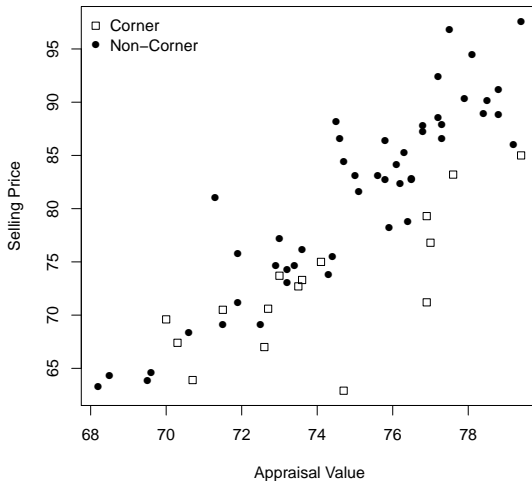
Have three variables in the data set:

- `price`: price house sells for (thousands of dollars)
- `value`: house value before selling (thousands of dollars)
- `corner`: indicator variable (=1 if house is on corner of block)

Total of $n = 64$ independent observations (16 corners).

Want to predict selling `price` from appraisal `value`, and determine if the relationship depends on the `corner` status of the house.

Real Estate Data Visualization



Real Estate Data Visualization (R Code)

```
> house=read.table("/Users/Nate/Desktop/houseprice.txt",header=TRUE)
> house[1:3,]
  price value corner
1  78.8  76.4      0
2  73.8  74.3      0
3  64.6  69.6      0
4  76.2  73.6      0
5  87.2  76.8      0
6  70.6  72.7      1
7  86.0  79.2      0
8  83.1  75.6      0
> plot(house$value,house$price,pch=ifelse(house$corner==1,0,16),
+       xlab="Appraisal Value",ylab="Selling Price")
> legend("topleft",c("Corner","Non-Corner"),pch=c(0,16),bty="n")
```

Real Estate Regression: Fit Model

Fit model with interaction between value and corner

```
> hmod=lm(price~value*corner,data=house)
> hmod
```

Call:

```
lm(formula = price ~ value * corner, data = house)
```

Coefficients:

(Intercept)	value	corner	value:corner
-126.905	2.776	76.022	-1.107

Real Estate Regression: Significance of Terms

```
> summary(hmod)
```

Call:

```
lm(formula = price ~ value * corner, data = house)
```

Residuals:

Min	1Q	Median	3Q	Max
-10.8470	-2.1639	0.0913	1.9348	9.9836

Coefficients:

	Estimate	Std. Error	t value	Pr(> t)	
(Intercept)	-126.9052	14.7225	-8.620	4.33e-12	***
value	2.7759	0.1963	14.142	< 2e-16	***
corner	76.0215	30.1314	2.523	0.01430	*
value:corner	-1.1075	0.4055	-2.731	0.00828	**

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 3.893 on 60 degrees of freedom

Multiple R-squared: 0.8233, Adjusted R-squared: 0.8145

F-statistic: 93.21 on 3 and 60 DF, p-value: < 2.2e-16

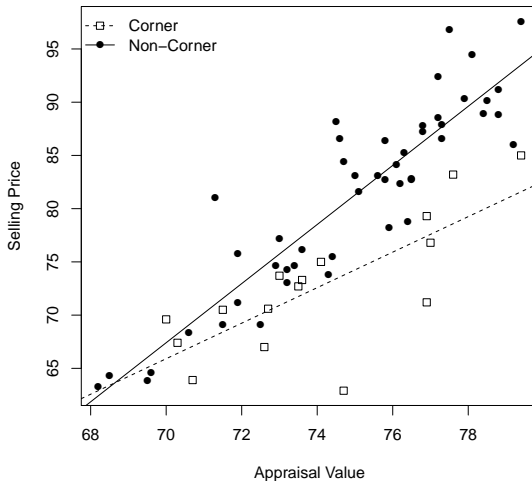
Real Estate Regression: Interpreting Results

```
> hmod$coef
```

(Intercept)	value	corner	value:corner
-126.905171	2.775898	76.021532	-1.107482

- $\hat{b}_0 = -126.90$ is expected selling price (in thousands of dollars) for non-corner houses that were valued at \$0.
- $\hat{b}_0 + \hat{b}_2 = -126.90 + 76.022 = -50.878$ is expected selling price (in thousands of dollars) for corner houses that were valued at \$0.
- $\hat{b}_1 = 2.776$ is the expected increase in selling price (in thousands of dollars) corresponding to a 1-unit (\$1,000) increase in appraisal value for non-corner houses
- $\hat{b}_1 + \hat{b}_3 = 2.775 - 1.107 = 1.668$ is the expected increase in selling price (in thousands of dollars) corresponding to a 1-unit (\$1,000) increase in appraisal value for corner houses

Real Estate Regression: Plotting Results



Real Estate Regression: Plotting Results (R Code)

```
> plot(house$value,house$price,pch=ifelse(house$corner==1,0,16),  
+      xlab="Appraisal Value",ylab="Selling Price")  
> abline(hmod$coef[1],hmod$coef[2])  
> abline(hmod$coef[1]+hmod$coef[3],hmod$coef[2]+hmod$coef[4],lty=2)  
> legend("topleft",c("Corner","Non-Corner"),lty=2:1,pch=c(0,16),bty="n")
```

Note that if you input the (appropriate) coefficients, you can still use the `abline` function to draw the regression lines.

Depression Data Description

Using depression data from Daniel (1999) *Biostatistics: A Foundation for Analysis in the Health Sciences*.


Total of $n = 36$ subjects participated in a depression study.

Have three variables in the data set:

- `effect`: effectiveness of depression treatment (high=effective)
- `age`: age of the participant (in years)
- `method`: method of treatment (3 levels: A, B, C)

Predict effectiveness from participant's age and treatment method.

Depression Data Visualization



Depression Data Visualization (R Code)

```
> depression=read.table("/Users/Nate/Desktop/depression.txt",header=TRUE)
> depression[1:8,]
  effect age method
1     56  21      A
2     41  23      B
3     40  30      B
4     28  19      C
5     55  28      A
6     25  23      C
7     46  33      B
8     71  67      C
> plot(depression$age,depression$effect,xlab="Age",ylab="Effect",type="n")
> text(depression$age,depression$effect,depression$method)
```

Depression Regression: Fit Model

Fit model with interaction between age and method

```
> dmod=lm(effect~age*method,data=depression)
```

```
> dmod$coef
```

(Intercept)	age	methodB	methodC	age:methodB	age:methodC
47.5155913	0.3305073	-18.5973852	-41.3042101	0.1931769	0.7028836

Note that R creates two indicator variables for `method`:

- $x_{iB} = \begin{cases} 1 & \text{if } i\text{-th observation is in treatment method B} \\ 0 & \text{otherwise} \end{cases}$
- $x_{iC} = \begin{cases} 1 & \text{if } i\text{-th observation is in treatment method C} \\ 0 & \text{otherwise} \end{cases}$

Depression Regression: Significance of Terms

```
> summary(dmod)
```

Call:

```
lm(formula = effect ~ age * method, data = depression)
```

Residuals:

	Min	1Q	Median	3Q	Max
	-6.4366	-2.7637	0.1887	2.9075	6.5634

Coefficients:

	Estimate	Std. Error	t value	Pr(> t)	
(Intercept)	47.51559	3.82523	12.422	2.34e-13	***
age	0.33051	0.08149	4.056	0.000328	***
methodB	-18.59739	5.41573	-3.434	0.001759	**
methodC	-41.30421	5.08453	-8.124	4.56e-09	***
age:methodB	0.19318	0.11660	1.657	0.108001	
age:methodC	0.70288	0.10896	6.451	3.98e-07	***

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 3.925 on 30 degrees of freedom

Multiple R-squared: 0.9143, Adjusted R-squared: 0.9001

F-statistic: 64.04 on 5 and 30 DF, p-value: 4.264e-15


Depression Regression: Interpreting Results

```
> dmod$coef
```

(Intercept)	age	methodB	methodC	age:methodB	age:methodC
47.5155913	0.3305073	-18.5973852	-41.3042101	0.1931769	0.7028836

- $\hat{b}_0 = 47.516$ is expected treatment effectiveness for subjects in method A who are 0 y/o.
- $\hat{b}_0 + \hat{b}_2 = 47.516 - 18.598 = 28.918$ is expected treatment effectiveness for subjects in method B who are 0 years old.
- $\hat{b}_0 + \hat{b}_3 = 47.516 - 41.304 = 6.212$ is expected treatment effectiveness for subjects in method C who are 0 years old.
- $\hat{b}_1 = 0.331$ is the expected increase in treatment effectiveness corresponding to a 1-unit (1 year) increase in age for treatment method A
- $\hat{b}_1 + \hat{b}_4 = 0.331 + 0.193 = 0.524$ is the expected increase in treatment effectiveness corresponding to a 1-unit (1 year) increase in age for treatment method B
- $\hat{b}_1 + \hat{b}_5 = 0.331 + 0.703 = 1.033$ is the expected increase in treatment effectiveness corresponding to a 1-unit (1 year) increase in age for treatment method C

Depression Regression: Plotting Results



Depression Regression: Plotting Results

```
> plot(depression$age, depression$effect, xlab="Age", ylab="Effect", type="n")
> text(depression$age, depression$effect, depression$method)
> abline(dmod$coef[1], dmod$coef[2])
> abline(dmod$coef[1]+dmod$coef[3], dmod$coef[2]+dmod$coef[5], lty=2)
> abline(dmod$coef[1]+dmod$coef[4], dmod$coef[2]+dmod$coef[6], lty=3)
> legend("bottomright", c("A", "B", "C"), lty=1:3, bty="n")
```

Interactions between Continuous Variables

Revisit the MLR model with two predictors and an interaction

$$y_i = b_0 + b_1 x_{i1} + b_2 x_{i2} + b_3 x_{i1} x_{i2} + e_i$$


and suppose that $x_{i1}, x_{i2} \in \mathbb{R}$ are both continuous predictors.

In this case, we model terms can be interpreted as:


- b_0 is expected value of Y when $x_{i1} = x_{i2} = 0$
- $b_1 + b_3 x_{i2}$ is expected change in Y corresponding to 1-unit change in x_{i1} holding x_{i2} fixed (i.e., conditioning on x_{i2})
- $b_2 + b_3 x_{i1}$ is expected change in Y corresponding to 1-unit change in x_{i2} holding x_{i1} fixed (i.e., conditioning on x_{i1})

Visualizing Continuous*Continuous Interactions


Simple regression



Multiple regression (additive)




Multiple regression (interaction)



Oceanography Data Description

Data from UCI Machine Learning: <http://archive.ics.uci.edu/ml/>

- Data originally from TAO project: <http://www.pmel.noaa.gov/tao/>
- Note that I have preprocessed the data a bit before analysis.



Oceanography Data Description (continued)

Buoys collect lots of different data:

```
> elnino[1:4,]
      obs year month day   date latitude longitude zon.winds mer.winds humidity air.temp ss.temp
4297 4297   94     1   1 940101    -0.01    250.01     -4.3       2.6      89.9    23.21    23.37
4298 4298   94     1   2 940102    -0.01    250.01     -4.1       1       90     23.16    23.45
4299 4299   94     1   3 940103    -0.01    250.01      -3       1.6     87.7    23.14    23.71
4300 4300   94     1   4 940104     0.00    250.00      -3       2.9     85.8    23.39    24.29
```

We will focus on predicting the sea surface temperatures (`ss.temp`) from the `latitude` and `longitude` locations of the buoys.

Oceanography Regression: Fit Additive Model

Fit additive model of latitude and longitude

```
> eladd=lm(ss.temp~latitude+longitude,data=elnino)
> summary(eladd)
```

Call:

```
lm(formula = ss.temp ~ latitude + longitude, data = elnino)
```

Residuals:

	Min	1Q	Median	3Q	Max
	-7.6055	-0.7229	0.1261	0.9039	5.0987

Coefficients:

	Estimate	Std. Error	t value	Pr(> t)
(Intercept)	35.2636388	0.0305722	1153.45	<2e-16 ***
latitude	0.0257867	0.0010006	25.77	<2e-16 ***
longitude	-0.0357496	0.0001445	-247.33	<2e-16 ***

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 1.48 on 86498 degrees of freedom

Multiple R-squared: 0.4184, Adjusted R-squared: 0.4184

F-statistic: 3.112e+04 on 2 and 86498 DF, p-value: < 2.2e-16

Oceanography Regression: Fit Interaction Model

Fit model with interaction between latitude and longitude

```
> elint=lm(ss.temp~latitude*longitude,data=elnino)
> summary(elint)
```

Call:

```
lm(formula = ss.temp ~ latitude * longitude, data = elnino)
```

Residuals:

Min	1Q	Median	3Q	Max
-7.5867	-0.6496	0.1000	0.8127	5.0922

Coefficients:

	Estimate	Std. Error	t value	Pr(> t)	
(Intercept)	3.541e+01	2.913e-02	1215.61	<2e-16	***
latitude	-5.245e-01	5.862e-03	-89.47	<2e-16	***
longitude	-3.638e-02	1.377e-04	-264.22	<2e-16	***
latitude:longitude	2.618e-03	2.752e-05	95.13	<2e-16	***

 Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1


Residual standard error: 1.408 on 86497 degrees of freedom

Multiple R-squared: 0.4735, Adjusted R-squared: 0.4735


F-statistic: 2.593e+04 on 3 and 86497 DF, p-value: < 2.2e-16

Oceanography Regression: Visualize Results

Additive Prediction



Interaction Prediction




Oceanography Regression: Visualize (R Code)


```
> newdata=expand.grid(longitude=seq(min(elnino$longitude),max(elnino$longitude),l=50),
+                      latitude=seq(min(elnino$latitude),max(elnino$latitude),l=50))
> yadd=predict(eladd,newdata)
> image(seq(min(elnino$longitude),max(elnino$longitude),l=50),
+        seq(min(elnino$latitude),max(elnino$latitude),l=50),
+        matrix(yadd,50,50),col=rev(rainbow(100,end=3/4)),
+        xlab="Longitude",ylab="Latitude",main="Additive Prediction")
> yint=predict(elint,newdata)
> image(seq(min(elnino$longitude),max(elnino$longitude),l=50),
+        seq(min(elnino$latitude),max(elnino$latitude),l=50),
+        matrix(yint,50,50),col=rev(rainbow(100,end=3/4)),
+        xlab="Longitude",ylab="Latitude",main="Interaction Prediction")
```

Oceanography Regression: Smoothing Spline Solution

a) Buoy Locations



b) Main Effect: Temporal



c) Main Effect: Spatial

