

Chapter 5: SQL Joins



5.1: Introduction to SQL Joins

5.2: Complex SQL Joins

Chapter 5: SQL Joins

5.1: Introduction to SQL Joins

5.2: Complex SQL Joins

Objectives

- Horizontally combine data from multiple tables.
- Distinguish between inner and outer SQL joins.
- Compare SQL joins to DATA step merges.

Combining Data from Multiple Tables

SQL uses set operators to combine tables vertically.

| |
|---------|
| Table A |
| Table B |

This produces results that can be compared to a DATA step concatenation.

Combining Data from Multiple Tables

SQL uses joins to combine tables horizontally.



This produces results that can be compared to a DATA step merge.

Poll 

Quiz

5.01 Multiple Choice Poll

Which of these DATA step statements is used to combine tables horizontally?

- a. SET
- b. APPEND
- c. MERGE
- d. INPUT
- e. INFILE

5.01 Multiple Choice Poll – Correct Answer

Which of these DATA step statements is used to combine tables horizontally?

- a. SET
- b. APPEND
- ☒ c. MERGE
- d. INPUT
- e. INFILE

Types of Joins


PROC SQL supports two types of joins:

- inner joins
- outer joins

Types of Joins

Inner joins


- return only matching rows
- enable a maximum of 256 tables to be joined at the same time.




Types of Joins

Outer joins


- return all matching rows, plus nonmatching rows from one or both tables
- can be performed on only two tables or views at a time.



Left



Full



Right

Cartesian Product

To understand how SQL processes a join, it is important to understand the concept of the Cartesian product.

A query that lists multiple tables in the FROM clause without a WHERE clause produces all possible combinations of rows from all tables. This result is called the *Cartesian product*.

```
select *  
  from one, two;
```

Cartesian Product

Table One

| X | A |
|---|---|
| 1 | a |
| 4 | d |
| 2 | b |

Table Two

| X | B |
|---|---|
| 2 | x |
| 3 | y |
| 5 | v |

Cartesian Product

Table One

| X | A |
|---|---|
| 1 | a |
| 4 | d |
| 2 | b |

Table Two

| X | B |
|---|---|
| 2 | x |
| 3 | y |
| 5 | v |

Result Set

| X | A | X | B |
|---|---|---|---|
| 1 | a | 2 | x |

Cartesian Product

Table One

| X | A |
|---|---|
| 1 | a |
| 4 | d |
| 2 | b |

Table Two

| X | B |
|---|---|
| 2 | x |
| 3 | y |
| 5 | v |

Result Set

| X | A | X | B |
|---|---|---|---|
| 1 | a | 2 | x |
| 1 | a | 3 | y |

Cartesian Product

Table One

| X | A |
|---|---|
| 1 | a |
| 4 | d |
| 2 | b |

Table Two

| X | B |
|---|---|
| 2 | x |
| 3 | y |
| 5 | v |

Result Set

| X | A | X | B |
|---|---|---|---|
| 1 | a | 2 | x |
| 1 | a | 3 | y |
| 1 | a | 5 | v |

Cartesian Product

Table One

| X | A |
|---|---|
| 1 | a |
| 4 | d |
| 2 | b |

Table Two

| X | B |
|---|---|
| 2 | x |
| 3 | y |
| 5 | v |

Result Set

| X | A | X | B |
|---|---|---|---|
| 1 | a | 2 | x |
| 1 | a | 3 | y |
| 1 | a | 5 | v |
| 4 | d | 2 | x |

Cartesian Product

Table One

| X | A |
|---|---|
| 1 | a |
| 4 | d |
| 2 | b |

Table Two

| X | B |
|---|---|
| 2 | x |
| 3 | y |
| 5 | v |

Result Set

| X | A | X | B |
|---|---|---|---|
| 1 | a | 2 | x |
| 1 | a | 3 | y |
| 1 | a | 5 | v |
| 4 | d | 2 | x |
| 4 | d | 3 | y |

Cartesian Product

Table One

| X | A |
|---|---|
| 1 | a |
| 4 | d |
| 2 | b |

Table Two

| X | B |
|---|---|
| 2 | x |
| 3 | y |
| 5 | v |

Result Set

| X | A | X | B |
|---|---|---|---|
| 1 | a | 2 | x |
| 1 | a | 3 | y |
| 1 | a | 5 | v |
| 4 | d | 2 | x |
| 4 | d | 3 | y |
| 4 | d | 5 | v |

Cartesian Product

Table One

| X | A |
|---|---|
| 1 | a |
| 4 | d |
| 2 | b |

Table Two

| X | B |
|---|---|
| 2 | x |
| 3 | y |
| 5 | v |

Result Set

| X | A | X | B |
|---|---|---|---|
| 1 | a | 2 | x |
| 1 | a | 3 | y |
| 1 | a | 5 | v |
| 4 | d | 2 | x |
| 4 | d | 3 | y |
| 4 | d | 5 | v |
| 2 | b | 2 | x |

Cartesian Product

Table One

| X | A |
|---|---|
| 1 | a |
| 4 | d |
| 2 | b |

Table Two

| X | B |
|---|---|
| 2 | x |
| 3 | y |
| 5 | v |

Result Set

| X | A | X | B |
|---|---|---|---|
| 1 | a | 2 | x |
| 1 | a | 3 | y |
| 1 | a | 5 | v |
| 4 | d | 2 | x |
| 4 | d | 3 | y |
| 4 | d | 5 | v |
| 2 | b | 2 | x |
| 2 | b | 3 | y |

Cartesian Product

Table One

| X | A |
|---|---|
| 1 | a |
| 4 | d |
| 2 | b |

Table Two

| X | B |
|---|---|
| 2 | x |
| 3 | y |
| 5 | v |

Result Set

| X | A | X | B |
|---|---|---|---|
| 1 | a | 2 | x |
| 1 | a | 3 | y |
| 1 | a | 5 | v |
| 4 | d | 2 | x |
| 4 | d | 3 | y |
| 4 | d | 5 | v |
| 2 | b | 2 | x |
| 2 | b | 3 | y |
| 2 | b | 5 | v |


Cartesian Product

Table One

| X | A |
|---|---|
| 1 | a |
| 4 | d |
| 2 | b |

Table Two

| X | B |
|---|---|
| 2 | x |
| 3 | y |
| 5 | v |



Result Set

| X | A | X | B |
|---|---|---|---|
| 1 | a | 2 | x |
| 1 | a | 3 | y |
| 1 | a | 5 | v |
| 4 | d | 2 | x |
| 4 | d | 3 | y |
| 4 | d | 5 | v |
| 2 | b | 2 | x |
| 2 | b | 3 | y |
| 2 | b | 5 | v |

Cartesian Product

Table One

| X | A |
|---|---|
| 1 | a |
| 4 | d |
| 2 | b |

} 3 rows

Table Two

| X | B |
|---|---|
| 2 | x |
| 3 | y |
| 5 | v |

Result Set

| X | A | X | B |
|---|---|---|---|
| 1 | a | 2 | x |
| 1 | a | 3 | y |
| 1 | a | 5 | v |
| 4 | d | 2 | x |
| 4 | d | 3 | y |
| 4 | d | 5 | v |
| 2 | b | 2 | x |
| 2 | b | 3 | y |
| 2 | b | 5 | v |

Cartesian Product

Table One

| X | A |
|---|---|
| 1 | a |
| 4 | d |
| 2 | b |

3 rows

Table Two

| X | B |
|---|---|
| 2 | x |
| 3 | y |
| 5 | v |

3 rows

Result Set

| X | A | X | B |
|---|---|---|---|
| 1 | a | 2 | x |
| 1 | a | 3 | y |
| 1 | a | 5 | v |
| 4 | d | 2 | x |
| 4 | d | 3 | y |
| 4 | d | 5 | v |
| 2 | b | 2 | x |
| 2 | b | 3 | y |
| 2 | b | 5 | v |

Cartesian Product

Table One

| X | A |
|---|---|
| 1 | a |
| 4 | d |
| 2 | b |

3 rows

X

3 rows

Table Two

| X | B |
|---|---|
| 2 | x |
| 3 | y |
| 5 | v |

Result Set

| X | A | X | B |
|---|---|---|---|
| 1 | a | 2 | x |
| 1 | a | 3 | y |
| 1 | a | 5 | v |
| 4 | d | 2 | x |
| 4 | d | 3 | y |
| 4 | d | 5 | v |
| 2 | b | 2 | x |
| 2 | b | 3 | y |
| 2 | b | 5 | v |

9 rows

Cartesian Product

The number of rows in a Cartesian product is the product of the number of rows in the contributing tables.

$$3 \times 3 = 9$$

$$1,000 \times 1,000 = 1,000,000$$

$$100,000 \times 100,000 = 10,000,000,000$$

A Cartesian product is rarely the **desired** result of a query.

Poll

Quiz



5.02 Quiz

How many rows are returned from this query?

```
select *  
  from three, four;
```

Table Three

| X | A |
|---|----|
| 1 | a1 |
| 1 | a2 |
| 2 | b1 |
| 2 | b2 |
| 4 | d |

Table Four

| X | B |
|---|----|
| 2 | x1 |
| 2 | x2 |
| 3 | y |
| 5 | v |

5.02 Quiz – Correct Answer

How many rows are returned from this query?

The query produces 20 rows.

```
select *  
  from three, four;
```

Table Three

| X | A |
|---|----|
| 1 | a1 |
| 1 | a2 |
| 2 | b1 |
| 2 | b2 |
| 4 | d |

Table Four

| X | B |
|---|----|
| 2 | x1 |
| 2 | x2 |
| 3 | y |
| 5 | v |

$$5 * 4 = 20$$

Partial Results Set

| X | A | X | B |
|---|----|---|----|
| 1 | a1 | 2 | x1 |
| 1 | a1 | 2 | x2 |
| 1 | a1 | 3 | y |
| 1 | a1 | 5 | v |
| 1 | a2 | 2 | x1 |
| 1 | a2 | 2 | x2 |
| 1 | a2 | 3 | y |
| 1 | a2 | 5 | v |
| 2 | b1 | 2 | x1 |
| 2 | b1 | 2 | x2 |

Inner Joins

Inner join syntax resembles Cartesian product syntax, but a WHERE clause restricts which rows are returned.

General form of an inner join:

```
SELECT column-1<, ...column-n>  
  FROM table-1/view-1<, ... table-n/view-n>  
  WHERE join-condition(s)  
        <AND other subsetting conditions>  
        <other clauses>;
```

Inner Joins

Inner join syntax resembles Cartesian product syntax, but a WHERE clause restricts which rows are returned.

General form of an inner join:

```
SELECT column-1<, ...column-n>
FROM table-1/view-1<, ... table-n/view-n>
WHERE join-condition(s)
        <AND other subsetting conditions>
        <other clauses>;
```

Significant syntax changes from earlier queries:

- The FROM clause references multiple tables.
- The WHERE clause includes join conditions in addition to other subsetting specifications.

Inner Joins

Conceptually, when processing an inner join, PROC SQL does the following:

1. builds the Cartesian product of all the tables listed
2. applies the WHERE clause to limit the rows returned

Inner Joins: Cartesian Product Built

Table One

| X | A |
|---|---|
| 1 | a |
| 4 | d |
| 2 | b |

Table Two

| X | B |
|---|---|
| 2 | x |
| 3 | y |
| 5 | v |

```
select *  
  from one, two
```

| X | A | X | B |
|---|---|---|---|
| 1 | a | 2 | x |
| 1 | a | 3 | y |
| 1 | a | 5 | v |
| 4 | d | 2 | x |
| 4 | d | 3 | y |
| 4 | d | 5 | v |
| 2 | b | 2 | x |
| 2 | b | 3 | y |
| 2 | b | 5 | v |

Inner Joins: WHERE Clause Restricts Rows

Table One

| X | A |
|---|---|
| 1 | a |
| 4 | d |
| 2 | b |

Table Two

| X | B |
|---|---|
| 2 | x |
| 3 | y |
| 5 | v |

```
select *  
  from one, two  
 where one.x=two.x;
```

| X | A | X | B |
|---|---|---|---|
| 1 | a | 2 | x |
| 1 | a | 3 | y |
| 1 | a | 5 | v |
| 4 | d | 2 | x |
| 4 | d | 3 | y |
| 4 | d | 5 | v |
| 2 | b | 2 | x |
| 2 | b | 3 | y |
| 2 | b | 5 | v |



Inner Joins: Results Are Returned

Table One

| X | A |
|---|---|
| 1 | a |
| 4 | d |
| 2 | b |

Table Two

| X | B |
|---|---|
| 2 | x |
| 3 | y |
| 5 | v |

```
select *  
  from one, two  
 where one.x=two.x;
```

| X | A | X | B |
|---|---|---|---|
| 2 | b | 2 | x |



Tables do not have to be sorted before they are joined.

Inner Joins

One method of displaying the X column only once is to use a table qualifier in the SELECT list.

Table One

| X | A |
|---|---|
| 1 | a |
| 4 | d |
| 2 | b |

Table Two

| X | B |
|---|---|
| 2 | x |
| 3 | y |
| 5 | v |

```
select one.x, a, b  
  from one, two  
 where one.x=two.x;
```

| X | A | B |
|---|---|---|
| 2 | b | x |

Inner Joins

Display all combinations of rows with matching keys, including duplicates.

Table Three

| X | A |
|---|----|
| 1 | a1 |
| 1 | a2 |
| 2 | b1 |
| 2 | b2 |
| 4 | d |

Table Four

| X | B |
|---|----|
| 2 | x1 |
| 2 | x2 |
| 3 | y |
| 5 | v |

```
proc sql;  
    select *  
        from three, four  
        where three.x=four.x;  
quit;
```

Inner Joins

Display all combinations of rows with matching keys, including duplicates.

Table Three

| X | A |
|---|----|
| 1 | a1 |
| 1 | a2 |
| 2 | b1 |
| 2 | b2 |
| 4 | d |

Table Four

| X | B |
|---|----|
| 2 | x1 |
| 2 | x2 |
| 3 | y |
| 5 | v |

Results Set

| X | A | X | B |
|---|----|---|----|
| 2 | b1 | 2 | x1 |
| 2 | b1 | 2 | x2 |
| 2 | b2 | 2 | x1 |
| 2 | b2 | 2 | x2 |

```
proc sql;  
    select *  
        from three, four  
        where three.x=four.x;  
quit;
```

Poll 

Quiz

Setup for the Poll

Run program **s105a02** and review the results to determine how many rows (observations) the DATA step MERGE statement produces in the output table.

Three

| X | A |
|---|----|
| 1 | a1 |
| 1 | a2 |
| 2 | b1 |
| 2 | b2 |
| 4 | d |

Four

| X | B |
|---|----|
| 2 | x1 |
| 2 | x2 |
| 3 | y |
| 5 | v |

```
data new;  
    merge three (in=InThree)  
           four (in=InFour);  
    by x;  
    if InThree and InFour;  
run;  
  
proc print data=new;  
run;
```

5.03 Multiple Choice Poll

How many rows (observations) result from the DATA step MERGE statement in program **s105a02**?

- a. 4
- b. 2
- c. 6
- d. 20
- e. None of the above

5.03 Multiple Choice Poll – Correct Answer

How many rows (observations) result from the DATA step MERGE statement in program **s105a02**?

- a. 4
- ☒ b. 2
- c. 6
- d. 20
- e. None of the above

Three

| X | A |
|---|----|
| 1 | a1 |
| 1 | a2 |
| 2 | b1 |
| 2 | b2 |
| 4 | d |

Four

| X | B |
|---|----|
| 2 | x1 |
| 2 | x2 |
| 3 | y |
| 5 | v |

New

| X | A | B |
|---|----|----|
| 2 | b1 | x1 |
| 2 | b2 | x2 |

Business Scenario

Display the name, city, and birth month of all Australian employees. Here is a sketch of the desired report:

| Australian Employees' Birth Months | | |
|------------------------------------|-----------|-------------|
| Name | City | Birth Month |
| Last, First | City Name | 1 |

Business Scenario

Considerations:

- **orion.Employee_Addresses** contains employee name, country, and city data.
- **orion.Payroll** contains employee birth dates.
- Both **orion.Employee_Addresses** and **orion.Payroll** contain **Employee_ID**.
- Names are stored in the **Employee_Name** column as Last, First.

Inner Joins

```
proc sql;  
title "Australian Employees' Birth Months";  
select Employee_Name as Name format=$25.,  
       City format=$25.,  
       month(Birth_Date) 'Birth Month' format=3.  
from   orion.Employee_Payroll,  
       orion.Employee_Addresses  
where  Employee_Payroll.Employee_ID=  
       Employee_Addresses.Employee_ID  
       and Country='AU'  
order by 3, City, Employee_Name;  
quit;
```

Inner Joins

Partial PROC SQL Output

Australian Employees Birthday Months

| Name | City | Birth Month |
|----------------------|-----------|-------------|
| Aisbitt, Sandy | Melbourne | 1 |
| Graham-Rowe, Jannene | Melbourne | 1 |
| Hieds, Merle | Melbourne | 1 |
| Sheedy, Sherie | Melbourne | 1 |
| Simms, Doungkamol | Melbourne | 1 |
| Tannous, Cos | Melbourne | 1 |
| Body, Meera | Sydney | 1 |
| Clarkson, Sharryn | Sydney | 1 |
| Dawes, Wilson | Sydney | 1 |
| Rusli, Skev | Sydney | 1 |
| Glattback, Ellis | Melbourne | 2 |
| Gromek, Gladys | Melbourne | 2 |

Inner Join Alternate Syntax

An inner join can also be accomplished using an alternate syntax, which limits the join to a maximum of two tables.

General form of an inner join:

```
SELECT column-1 <, ...column-n>  
FROM table-1  
INNER JOIN  
      table-2  
ON join-condition(s)  
    <other clauses>;
```



This syntax is common in SQL code produced by code generators such as SAS Enterprise Guide. The ON clause specifies the JOIN criteria; a WHERE clause **can** be added to subset the results.

Inner Join Alternate Syntax

```
proc sql;
title "Australian Employees' Birth Months";
select Employee_Name as Name format=$25.,
       City format=$25.,
       month(Birth_Date) 'Birth Month' format=3.
from   orion.Employee_Payroll
       inner join
       orion.Employee_Addresses
       on Employee_Payroll.Employee_ID=
          Employee_Addresses.Employee_ID
where  Country='AU'
order  by 3, City, Employee_Name;
quit;
```

Poll 

Quiz

5.04 Multiple Choice Poll

How many tables can be combined using a single inner join?

- a. 2
- b. 32
- c. 256
- d. 512
- e. Limited only by my computer's resources
- f. No limit

5.04 Multiple Choice Poll – Correct Answer

How many tables can be combined using a single inner join?

- a. 2
- b. 32
- ☒ c. 256
- d. 512
- e. Limited only by my computer's resources
- f. No limit



Question & Answer


Outer Joins

Inner joins returned only matching rows. When you join tables, you might want to include nonmatching rows as well as matching rows.


Outer Joins

You can retrieve both nonmatching and matching rows using an outer join.


Outer joins include left, full, and right outer joins. Outer joins can process only two tables at a time.



Left



Full



Right

Compare Inner Joins And Outer Joins

The following table is a comparison of inner and outer join syntax and limitations:

| Key Point | Inner Join | Outer Join |
|----------------|--|--|
| Table Limit | 256 | 2 |
| Join Behavior | Returns matching rows only | Returns matching and nonmatching rows |
| Join Options | Matching rows only | LEFT, FULL, RIGHT |
| Syntax changes | <ul style="list-style-type: none">■ Multiple tables in the FROM clause■ WHERE clause that specifies join criteria | ON clause that specifies join criteria |

Outer Joins

Outer join syntax is similar to the inner join alternate syntax.

General form of an outer join:

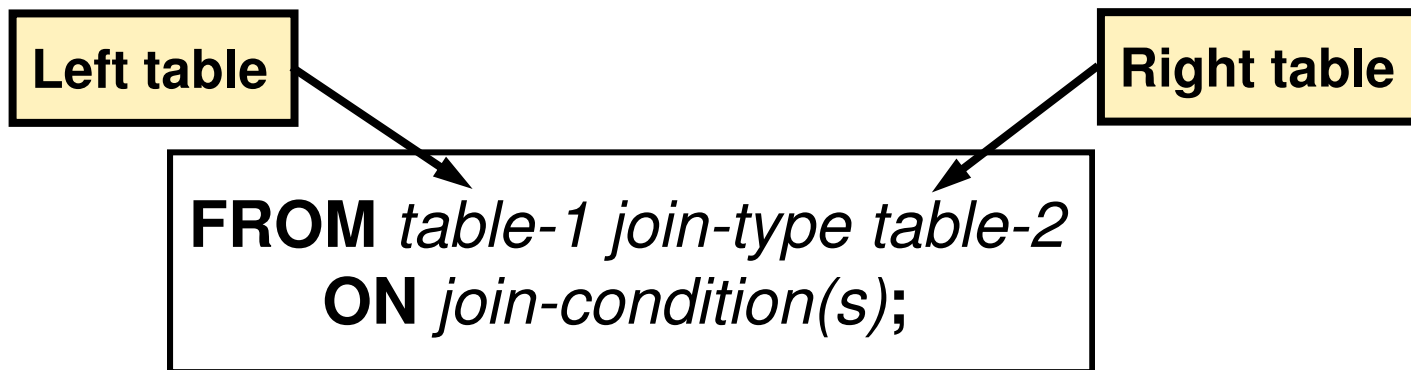
```
SELECT column-1 <, ...column-n>  
FROM table-1  
LEFT|RIGHT|FULL JOIN  
      table-2  
ON join-condition(s)  
    <other clauses>;
```

The ON clause specifies the join criteria in outer joins.

Determining Left and Right

Consider the position of the tables in the FROM clause.

- Left joins include all rows from the first (left) table, even if there are no matching rows in the second (right) table.
- Right joins include all rows from the second (right) table, even if there are no matching rows in the first (left) table.
- Full joins include all rows from both tables, even if there are no matching rows in either table.



Left Join

Table One

| X | A |
|---|---|
| 1 | a |
| 4 | d |
| 2 | b |

Table Two

| X | B |
|---|---|
| 2 | x |
| 3 | y |
| 5 | v |

```
select *  
  from one left join two  
        on one.x = two.x;
```

| X | A | X | B |
|---|---|---|---|
| 1 | a | . | |
| 2 | b | 2 | x |
| 4 | d | . | |

Right Join

Table Two

| X | B |
|---|---|
| 2 | x |
| 3 | y |
| 5 | v |

Table One

| X | A |
|---|---|
| 1 | a |
| 4 | d |
| 2 | b |

```
select *  
  from two right join one  
    on one.x = two.x;
```

| X | B | X | A |
|---|---|---|---|
| . | | 1 | a |
| 2 | x | 2 | b |
| . | | 4 | d |

Full Join

Table One

| X | A |
|---|---|
| 1 | a |
| 4 | d |
| 2 | b |

Table Two

| X | B |
|---|---|
| 2 | x |
| 3 | y |
| 5 | v |

```
select *  
  from one full join two  
    on one.x = two.x;
```

| X | A | X | B |
|---|---|---|---|
| 1 | a | . | |
| 2 | b | 2 | x |
| . | | 3 | y |
| 4 | d | . | |
| . | | 5 | v |

Business Scenario


List the employee ID and gender for all married employees. Include the names of any charities to which the employee donates via the company program.



Business Scenario

Considerations:


- The table `orion.Employee_Payroll` contains gender and marital status information.



Business Scenario

Considerations:


- The table **orion.Employee_Payroll** contains gender and marital status information.
- The table **orion.Employee_Donations** contains records only for those employees who donate to a charity via the company program.



Business Scenario

Considerations:


- The table **orion.Employee_Payroll** contains gender and marital status information.
- The table **orion.Employee_Donations** contains records only for those employees who donate to a charity via the company program.
- Less than half of all employees are married.



Business Scenario

Considerations:


- The table **orion.Employee_Payroll** contains gender and marital status information.
- The table **orion.Employee_Donations** contains records only for those employees who donate to a charity via the company program.
- Less than half of all employees are married.



Business Scenario

Considerations:

- The **orion.Employee_Payroll** table contains gender and marital status information.
- The **orion.Employee_Donations** table contains records only for those employees who donate to a charity via the company program.
- Less than half of all employees are married.



Poll 

Quiz

5.05 Multiple Choice Poll

For the report, you need the data for all married employees from **orion.Employee_Payroll**. You also want to include the charity names from the **orion.Employee_Donations** table if **Employee_ID** matches. What type of join should you use to combine the information from these two tables?

- a. Inner Join
- b. Left Join
- c. Full Join
- d. None of the above

5.05 Multiple Choice Poll – Correct Answer

For the report, you need the data for all married employees from **orion.Employee_Payroll**.

You also want to include the charity names from the **orion.Employee_Donations** table if

Employee_ID matches. What type of join should you use to combine the information from these two tables?

- a. Inner Join
- ☒ b. Left Join
- c. Full Join
- d. None of the above

Outer Joins

```
proc sql;  
    select Employee_payroll.Employee_ID,  
           Employee_Gender, Recipients  
    from orion.Employee_payroll  
    left join  
        orion.Employee_donations  
    on Employee_payroll.Employee_ID=  
        Employee_donations.Employee_ID  
    where Marital_Status="M"  
;  
quit;
```

Outer Joins

Partial PROC SQL Output (Rows 203-215)

| Employee_ID | Employee_ Gender | Recipients |
|-------------|---------------------|--|
| 121128 | F | Cancer Cures, Inc. |
| 121131 | M | Vox Victimas 40%, Conserve Nature, Inc. 60% |
| 121132 | M | EarthSalvors 50%, Vox Victimas 50% |
| 121133 | M | Disaster Assist, Inc. |
| 121138 | M | Cuidadores Ltd. |
| 121139 | F | |
| 121142 | M | AquaMissions International 10%, Child Survivors 90% |
| 121143 | M | Mitleid International 60%, Save the Baby Animals 40% |
| 121144 | F | |
| 121145 | M | Save the Baby Animals |
| 121146 | F | |
| 121147 | F | Cuidadores Ltd. 50%, Mitleid International 50% |
| 121148 | M | |



Remember that output order is not guaranteed unless you use an ORDER BY clause.

Using a Table Alias

An *alias* is a table nickname. You can assign an alias to a table by following the table name in the FROM clause with the AS keyword and a nickname for the table. Then use the alias in other clauses of the QUERY statement.

General form of the FROM clause:

```
SELECT alias-1.column-1<, ...alias-2.column-n>  
  FROM table-1 AS alias-1  
        join-type  
        table-2 AS alias-2  
  ON join-condition(s)  
    <other clauses>;
```

Using a Table Alias

```
proc sql;  
  select p.Employee_ID, Employee_Gender,  
         Recipients  
  from orion.Employee_payroll as p  
  left join  
    orion.Employee_donations as d  
  on p.Employee_ID=d.Employee_ID  
  where Marital_Status="M"  
;  
quit;
```

DATA Step Merge (Review)

A DATA step with MERGE and BY statements automatically overlays same-name columns.

Table One

| X | A |
|---|---|
| 1 | a |
| 4 | d |
| 2 | b |

Table Two

| X | B |
|---|---|
| 2 | x |
| 3 | y |
| 5 | v |

Table One must be sorted or indexed on column X before a merge can be performed.

```
data merged;  
    merge one two;  
    by x;  
run;  
proc print data=merged;  
run;
```

Output

| X | A | B |
|---|---|---|
| 1 | a | |
| 2 | b | x |
| 3 | | y |
| 4 | d | |
| 5 | | v |

SQL Join versus DATA Step Merge

SQL joins do not automatically overlay same-named columns.

Table One

| X | A |
|---|---|
| 1 | a |
| 4 | d |
| 2 | b |

Table Two

| X | B |
|---|---|
| 2 | x |
| 3 | y |
| 5 | v |

```
proc sql;  
    select one.x, a, b  
        from one full join two  
        on one.x=two.x  
;  
quit;
```

Output

| X | A | B |
|---|---|---|
| 1 | a | |
| 2 | b | x |
| ○ | | y |
| 4 | d | |
| ○ | | v |

The COALESCE Function

The COALESCE function returns the value of the first non-missing argument.

General form of the COALESCE function:

COALESCE(*argument-1*, *argument-2*<, ...*argument-n*)

argument can be a constant, expression, or variable name. When all arguments are missing, COALESCE returns a missing value.



All arguments must be of the same type (character or numeric).

SQL Join versus DATA Step Merge

You can use the COALESCE function to overlay columns.

Table One

| X | A |
|---|---|
| 1 | a |
| 4 | d |
| 2 | b |

Table Two

| X | B |
|---|---|
| 2 | x |
| 3 | y |
| 5 | v |

```
proc sql;  
  select coalesce(one.x,two.x)  
         as x,a,b  
        from one full join two  
        on one.x=two.x;  
quit;
```

Output

| X | A | B |
|---|---|---|
| 1 | a | |
| 2 | b | x |
| 3 | | y |
| 4 | d | |
| 5 | | v |

SQL Join versus DATA Step Merge

| Key Points | SQL Join | DATA Step Merge |
|--|--------------|-----------------|
| Explicit sorting of data before join/merge | Not required | Required |
| Same-named columns in join/merge expressions | Not required | Required |
| Equality in join or merge expressions | Not required | Required |



Question & Answer



Exercise

This exercise reinforces the concepts discussed previously.

Chapter 5: SQL Joins



5.1: Introduction to SQL Joins

5.2: Complex SQL Joins

Objectives

- Create and use in-line views.
- Use in-line views and subqueries to simplify coding a complex query.

In-Line Views

In-line views are often useful when you build complex SQL queries.

An *inline view* is

- a temporary “virtual table” that exists only during query execution
- created by placing a query expression in a FROM clause where a table name would normally be used.

In-Line Views

An in-line view is a query expression (SELECT statement) that resides in a FROM clause. It acts as a virtual table, used in place of a physical table in a query.

```
proc sql;  
    select *  
        from  
            (in-line view query expression)  
quit;
```

Business Scenario

List all active Sales employees having annual salaries significantly (more than 5%) lower than the average salary for everyone with the same job title.



Considerations

First, you must calculate the average salaries for active employees in the Sales department, grouped by job title.

Next, you must match each employee to a GROUP-BY job title.

Finally, you must compare the employee's salary to the group's average to determine if it is more than 5% below the group average.

In-Line Views

Build a query to produce the aggregate averages.

```
proc sql;
title   'Sales Department Average Salary';
title2  'By Job Title';
    select Job_Title,
           avg(Salary) as Job_Avg
           format=comma7.
    from  orion.Employee_payroll as p,
           orion.Employee_organization as o
    where p.Employee_ID=o.Employee_ID
           and not Employee_Term_Date
           and o.Department="Sales"
    group by Job_Title;
quit;
```


In-Line Views

PROC SQL Output


| Sales Department Average Salary by Job Title | |
|---|---------|
| Job_Title | Job_Avg |
| Sales Rep. I | 26,576 |
| Sales Rep. II | 27,348 |
| Sales Rep. III | 29,214 |
| Sales Rep. IV | 31,589 |

In-Line Views

If you create a table from the results of the query, you can join this table and the **orion.Employee_payroll** table and subset the appropriate rows to get the answer. This adds unnecessary I/O.

Would it be useful to use only the query itself in place of a table?

In SQL, you can with an in-line view!



In-Line Views

Using a query in the FROM clause in place of a table causes the query output to be used as an in-line view.

```
proc sql;
title  'Employees with salaries less than';
title2 '95% of the average for their job';
  select Employee_Name, emp.Job_Title,
         Salary format=comma7., Job_Avg format=comma7.
  from (select Job_Title,
              avg(Salary) as Job_Avg format=comma7.
        from orion.Employee_payroll as p,
             orion.Employee_organization as o
        where p.Employee_ID=o.Employee_ID
              and not Employee_Term_Date
              and o.Department="Sales"
        group by Job_Title) as job,
        orion.Salesstaff as emp
  where emp.Job_Title=job.Job_Title
        and Salary < Job_Avg*.95
  order by Job_Title, Employee_Name;
```

In-Line Views

PROC SQL Output

| Employees with salaries less than 95% of the average for their job | | | | |
|---|--------------------|------------------------------|---------|--|
| Employee_Name | Employee Job Title | Employee Annual Salary | Job_Avg | |
| Ould, Tulsidas | Sales Rep. I | 22,710 | 26,576 | |
| Polky, Asishana | Sales Rep. I | 25,110 | 26,576 | |
| Tilley, Kimiko | Sales Rep. I | 25,185 | 26,576 | |
| Voron, Tachaun | Sales Rep. I | 25,125 | 26,576 | |



Question & Answer

Business Scenario

In 2003, Top Sports launched a premium line of sleeping bags called Expedition Zero, which was sold through Orion Star.

The CEO of Top Sports wants to send a letter of thanks to the manager of each employee who sold Expedition Zero sleeping bags in 2003, with a \$50 reward certificate (in U.S. dollars) to be presented by the manager to the employee.

The Task:

Prepare a list of the managers' names and the cities in which they are located.

Planning the Complex Query

Step 1

Identify the employees who sold Expedition Zero merchandise in 2003.

Step 2

Find the employee identifier for the managers of these employees

Step 3

Obtain the managers' names and city information.

Complex Query: Step 1 Considerations

Step 1

Get employee IDs for employees who sold Expedition Zero merchandise in 2003.

Select the employee's identifier (**Employee_ID**) from the results of joining the **Order_Fact** and **Product_Dim** tables on **Product_ID**, where **Product_Name** contains Expedition Zero. Exclude Internet orders (**Employee_ID NE 99999999**).

Coding the Complex Query

Step 1

Write a query to obtain the employee ID of all employees who sold Expedition Zero merchandise in 2003.

```
select distinct Employee_ID
  from orion.Order_Fact as o,
       orion.Product_Dim as p
 where o.Product_ID=p.Product_ID
       and year(Order_Date)=2003
       and Product_Name contains
       'Expedition Zero'
       and Employee_ID ne 99999999;
```

Coding the Complex Query

Step 1 PROC SQL Output

| <u>Employee ID</u> |
|--------------------|
| 120145 |
| 120732 |

Complex Query: Step 2 Considerations

Step 2

Find the employee identifier for the managers of these employees.

Select the manager's identifier (**Manager_ID**) from the results of joining the **Employee_Organization** table with the first query's results on **Employee_ID**.

Poll 

Quiz

5.06 Multiple Choice Poll

To join the **Employee_Organization** table with the Step 1 query results, you use the query from Step 1 as which of the following?

- a. an in-line view
- b. a subquery

5.06 Multiple Choice Poll – Correct Answer

To join the **Employee_Organization** table with the Step 1 query results, you use the query from Step 1 as which of the following?

- ☒ a. an in-line view
- ☐ b. a subquery


A query used in place of a physical table in a SELECT statement FROM clause is called an in-line view.

Coding the Complex Query

Step 2

Write a query to obtain the manager ID of the employee's manager.

```
select Manager_ID
  from orion.Employee_Organization as o,
  (<Step 1 query results>) as ID
 where o.Employee_ID=ID.Employee_ID;
```



| Employee_ID |
|-------------|
| 120145 |
| 120732 |

Coding the Complex Query

Step 2

Write a query to obtain the manager ID of the employee's manager.

```
select Manager_ID
  from orion.Employee_Organization as o,
    (select distinct Employee_ID
      from orion.Order_Fact as o,
        orion.Product_Dim as p
     where o.Product_ID=p.Product_ID
           and year(Order_Date)=2003
           and Product_Name
             contains 'Expedition Zero'
           and Employee_ID ne 99999999) as ID
 where o.Employee_ID=ID.Employee_ID;
```


Coding the Complex Query

Step 2 PROC SQL Output

| Manager_ID |
|------------|
| 120103 |
| 120736 |

Complex Query: Step 3 Considerations

Step 3 Find the managers' names and cities.

Select the employee's name (**Employee_Name**) and **City** from the **Employee_Addresses** table, where **Employee_ID** matches **Manager_ID** in the results of the previous query.

Poll 

Quiz

5.07 Poll

Is it possible to use the entire query in Step 2 as a subquery?

- ☐ Yes
- ☐ No

5.07 Poll – Correct Answer

Is it possible to use the entire query in Step 2 as a subquery?

- ☒ Yes
- ☐ No


A subquery can return values for multiple rows, but must return values for only one column. When submitted on its own, the query in Step 2 returns two rows and only one column, so it can be used as a non-correlated subquery.

Coding the Complex Query

Step 3

Write a query to obtain the managers' names and city information.

```
proc sql;  
select Employee_Name format=$25. as Name, City  
  from orion.Employee_Addresses  
 where Employee_ID in  
    (<Step 2 query results>);
```



| Manager_ID |
|------------|
| 120145 |
| 120732 |

Coding the Complex Query

Step 3

```
proc sql;
select Employee_Name format=$25. as Name
      , City
      from orion.Employee_Addresses
      where Employee_ID in
          (select Manager_ID
           from orion.Employee_Organization as o,
           (select distinct Employee_ID
            from orion.Order_Fact as o,
              orion.Product_Dim as p
            where o.Product_ID=p.Product_ID
              and year(Order_Date)=2003
              and Product_Name contains
                  'Expedition Zero'
              and Employee_ID ne 99999999) as ID
           where o.Employee_ID=ID.Employee_ID);
```

Coding the Complex Query

Step 3 PROC SQL Output

| Name | City |
|---------------|------------|
| Dawes, Wilson | Sydney |
| Kiemle, Parie | Miami-Dade |

Coding the Complex Query

You can also solve this problem using a multiway join.

```
proc sql;
  select distinct Employee_Name format=$25. as Name, City
    from orion.Order_Fact as of,
         orion.Product_Dim as pd,
         orion.Employee_Organization as eo,
         orion.Employee_Addresses as ea
   where of.Product_ID=pd.Product_ID
         and of.Employee_ID=eo.Employee_ID
         and ea.Employee_ID=eo.Manager_ID
         and Product_Name contains 'Expedition Zero'
         and year(Order_Date)=2003
         and eo.Employee_ID ne 99999999
;
quit;
```

Chapter Review

1. How many rows are returned by the following query?

```
proc sql;  
    select *  
    from  
        table1, table2;  
quit;
```

Table1

| X | A |
|---|---|
| 1 | a |
| 3 | d |
| 2 | b |

Table2

| X | B |
|---|---|
| 2 | x |
| 1 | y |
| 3 | v |

Chapter Review Answers

1. How many rows are returned by the following query?

```
proc sql;  
    select *  
        from  
            table1, table2;  
quit;
```

Table1

| X | A |
|---|---|
| 1 | a |
| 3 | d |
| 2 | b |

Table2

| X | B |
|---|---|
| 2 | x |
| 1 | y |
| 3 | v |

**This query produces a Cartesian product.
Nine rows will be returned.**

Chapter Review

2. Which of the following statements describes an advantage of using a PROC SQL view?
 - a. Views often save space, because a view is usually quite small compared with the data that it accesses.
 - b. Views can provide users a simpler alternative to frequently retrieving and submitting query code to produce identical results.
 - c. Views hide complex query details from users.
 - d. All of the above

Chapter Review Answers

2. Which of the following statements describes an advantage of using a PROC SQL view?
- a. Views often save space, because a view is usually quite small compared with the data that it accesses.
 - b. Views can provide users a simpler alternative to frequently retrieving and submitting query code to produce identical results.
 - c. Views hide complex query details from users.
 - ☒ d. All of the above

Chapter Review

3. Outer and Inner Joins:
 - a. An **outer join** can operate on a maximum of ____ tables simultaneously.
 - b. An **inner join** can operate on a maximum of ____ tables simultaneously.

Chapter Review Answers

3. Outer and Inner Joins:

- a. An **outer join** can operate on a maximum of 2 tables simultaneously.
- b. An **inner join** can operate on a maximum of 256 tables simultaneously.

Chapter Review

4. True or False:

An in-line view can be used on a WHERE or HAVING clause and can return many rows of data, but must return only one column.

Chapter Review Answers

4. True or False:

An in-line view can be used on a WHERE or HAVING clause and can return many rows of data, but must return only one column.

False

An in-line view is a query used in the FROM clause in place of a table. An in-line view can return any number of rows or columns.